


# Catalyzing Accelerator Design and Programming for Machine Learning

Hongzheng Chen, Niansong Zhang, Shaojie Xiang, Zhiru Zhang

Cornell University


C4ML workshop @ CGO'25



Cornell University



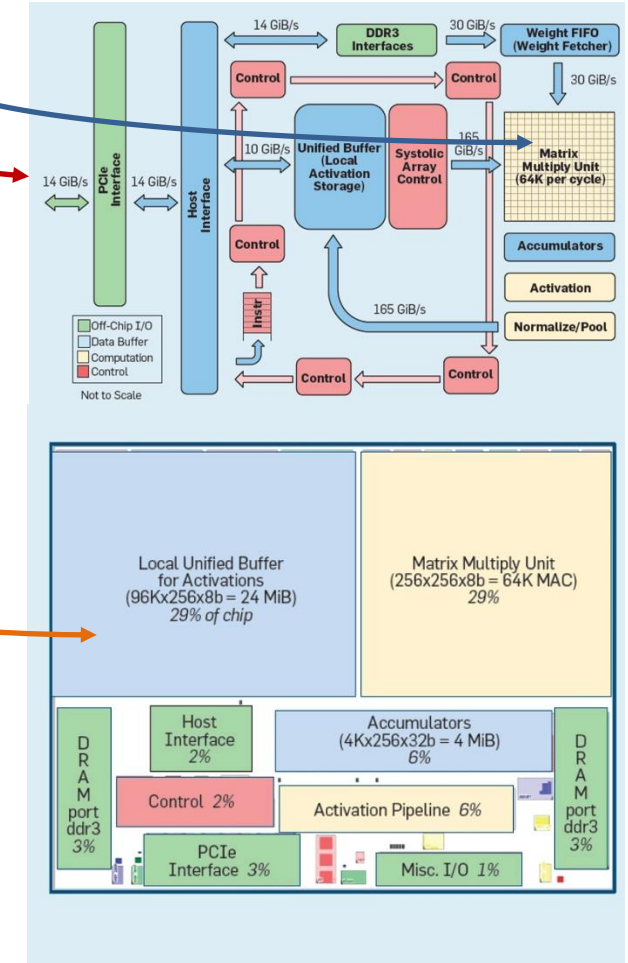

# Deep Learning Models Become Larger and Larger...



- ▶ The computational demands of large models have outgrown the capabilities of existing hardware
- ▶ Specialized accelerators are essential for speeding up model training and inference

# Complexity in Specialized Accelerator Design

- ▶ Accelerator design is different from programming on general processors
  - Custom processing engines (PEs)
  - Custom data communication
  - Custom memory hierarchy



An AI accelerator example

# Challenge: Balancing Manual Control & Compiler Optimization

```
void systolic_tile(int8_t A_tile[2][768],
int8_t B_tile[768][2],
int8_t C_tile[2][2]) {
#pragma dataflow
hls::stream<int8_t> A_fifo[2][3], B_fifo[2][3];
#pragma stream variable=A/B_fifo depth=3
#pragma partition variable=A/B/C_tile complete dim=1
for (int k4 = 0; k4 < 768; k4++) {
for (int m = 0; m < 2; m++) {
int8_t v105 = A_tile[m][k4];
A_fifo[m][0].write(v105);
// ... write B_fifo
}
for (int Ti = 0; Ti < 2; ++Ti) {
#pragma HLS unroll
for (int Tj = 0; Tj < 2; ++Tj) {
#pragma HLS unroll
// ... load A/B_fifo
PE_kernel(A_in, A_out, B_in, B_out, C, Ti, Tj);
}
}
}
}
```

Compute  
Customization

```
void matmul(int8_t A[512][768], int8_t B[768][768],
int8_t C[512][768]) {
int8_t local_A[2][768], local_B[768][2], local_C[2][2];
for (int mi = 0; mi < 256; mi++) {
for (int ni = 0; ni < 384; ni++) {
// ... load A, B
systolic_tile(local_A, local_B, local_C);
}
}
}
```



Vanilla Matmul (1% theoretical peak perf.)

+ Compute customization (~30% peak)

- + Loop tiling
- + Loop unrolling
- + Loop & function pipelining

Increasing programming effort

Memory



Accelerator

# Challenge: Balancing Manual Control & Compiler Optimization

```
void systolic_tile(int8_t A_tile[2][768],
int8_t B_tile[768][2],
int8_t C_tile[2][2]) {
#pragma dataflow
hls::stream<int8_t> A_fifo[2][3], B_fifo[2][3];
#pragma stream variable=A/B_fifo depth=3
#pragma partition variable=A/B/C_tile complete dim=1
for (int k4 = 0; k4 < 768; k4++) {
for (int m = 0; m < 2; m++) {
int8_t v105 = A_tile[m][k4];
A_fifo[m][0].write(v105);
// ... write B_fifo
}}
for (int Ti = 0; Ti < 2; ++Ti) {
#pragma HLS unroll
for (int Tj = 0; Tj < 2; ++Tj) {
#pragma HLS unroll
// ... load A/B_fifo
PE_kernel(A_in, A_out, B_in, B_out, C, Ti, Tj);
}}}
}
```

**Memory  
Customization**

**Compute  
Customization**

```
void matmul(int8_t A[512][768], int8_t B[768][768],
int8_t C[512][768]) {
int8_t local_A[2][768], local_B[768][2], local_C[2][2];
for (int mi = 0; mi < 256; mi++) {
for (int ni = 0; ni < 384; ni++) {
// ... load A, B
systolic_tile(local_A, local_B, local_C);
}}}
}
```

**Vanilla Matmul (1% theoretical peak perf.)**

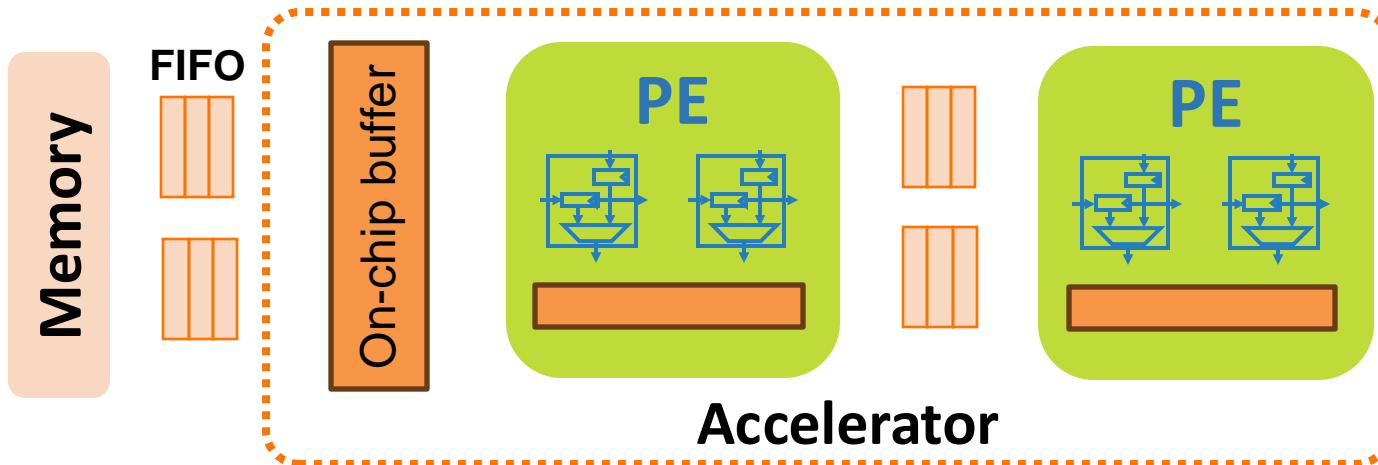
+ Compute customization (~30% peak)

- + Loop tiling
- + Loop unrolling
- + Loop & function pipelining

+ Custom memory hierarchy (~50% peak)

- + Tiling & data reuse buffers
- + Memory banking/partitioning

Increasing programming effort



# Challenge: Balancing Manual Control & Compiler Optimization

```
void systolic_tile(int8_t A_tile[2][768],
int8_t B_tile[768][2],
int8_t C_tile[2][2]) {
#pragma dataflow
hls::stream<int8_t> A_fifo[2][3], B_fifo[2][3];
#pragma stream variable=A/B_fifo depth=3
#pragma partition variable=A/B/C_tile complete dim=1
for (int k4 = 0; k4 < 768; k4++) {
for (int m = 0; m < 2; m++) {
int8_t v105 = A_tile[m][k4];
A_fifo[m][0].write(v105);
// ... write B_fifo
}
}
for (int Ti = 0; Ti < 2; ++Ti) {
#pragma HLS unroll
for (int Tj = 0; Tj < 2; ++Tj) {
#pragma HLS unroll
// ... load A/B_fifo
PE_kernel(A_in, A_out, B_in, B_out, C, Ti, Tj);
}
}
}
```

**Communication  
Customization**

**Memory  
Customization**

**Compute  
Customization**

```
void matmul(int8_t A[512][768], int8_t B[768][768],
int8_t C[512][768]) {
int8_t local_A[2][768], local_B[768][2], local_C[2][2];
for (int mi = 0; mi < 256; mi++) {
for (int ni = 0; ni < 384; ni++) {
// ... load A, B
systolic_tile(local_A, local_B, local_C);
}
}
}
```

**Vanilla Matmul (1% theoretical peak perf.)**

+ Compute customization (~30% peak)

- + Loop tiling
- + Loop unrolling
- + Loop & function pipelining


+ Custom memory hierarchy (~50% peak)

- + Tiling & data reuse buffers
- + Memory banking/partitioning

+ Data movement optimization (~95% peak)

- + Data streaming
- + Data packing (vectorization)
- + Memory coalescing
- + Systolic communication

Increasing programming effort



# Challenge: Balancing Manual Control & Compiler Optimization

```
void systolic_tile(int8_t A_tile[2][768],
int8_t B_tile[768][2],
int8_t C_tile[2][2]) {
#pragma adataflow
hls::stream<int8_t> A_fifo[2][3], B_fifo[2][3];
#pragma stream variable=A/B_fifo depth=3
#pragma partition variable=A/B/C_tile complete dim=1
for (int k4 = 0; k4 < 768; k4++) {
for (int m = 0; m < 2; m++) {
int8_t v105 = A_tile[m][k4];
A_fifo[m][0].write(v105);
// ... write B_fifo
}}
for (int Ti = 0; Ti < 2; ++Ti) {
#pragma HLS unroll
for (int Tj = 0; Tj < 2; ++Tj) {
#pragma HLS unroll
// ... load A/B_fifo
PE_kernel(A_in, A_out, B_in, B_out, C, Ti, Tj);
}}}

void matmul(int8_t A[512][768], int8_t B[768][768],
int8_t C[512][768])
{
int8_t local_A[2][768], local_B[768][2], local_C[2][2];
for (int mi = 0; mi < 256; mi++) {
for (int ni = 0; ni < 384; ni++) {
// ... load A, B
systolic_tile(local_A, local_B, local_C);
}}}

```

**Communication  
Customization**

**Memory  
Customization**

**Compute  
Customization**

**Vanilla Matmul (1% theoretical peak perf.)**

- + Compute customization (~30% peak)
  - + Loop tiling
  - + Loop unrolling
  - + Loop & function pipelining


**Existing HLS compiler**

**e.g., ScaleHLS [HPCA'22]**

- + Custom memory hierarchy (~50% peak)
  - + Tiling & data reuse buffers
  - + Memory banking/partitioning

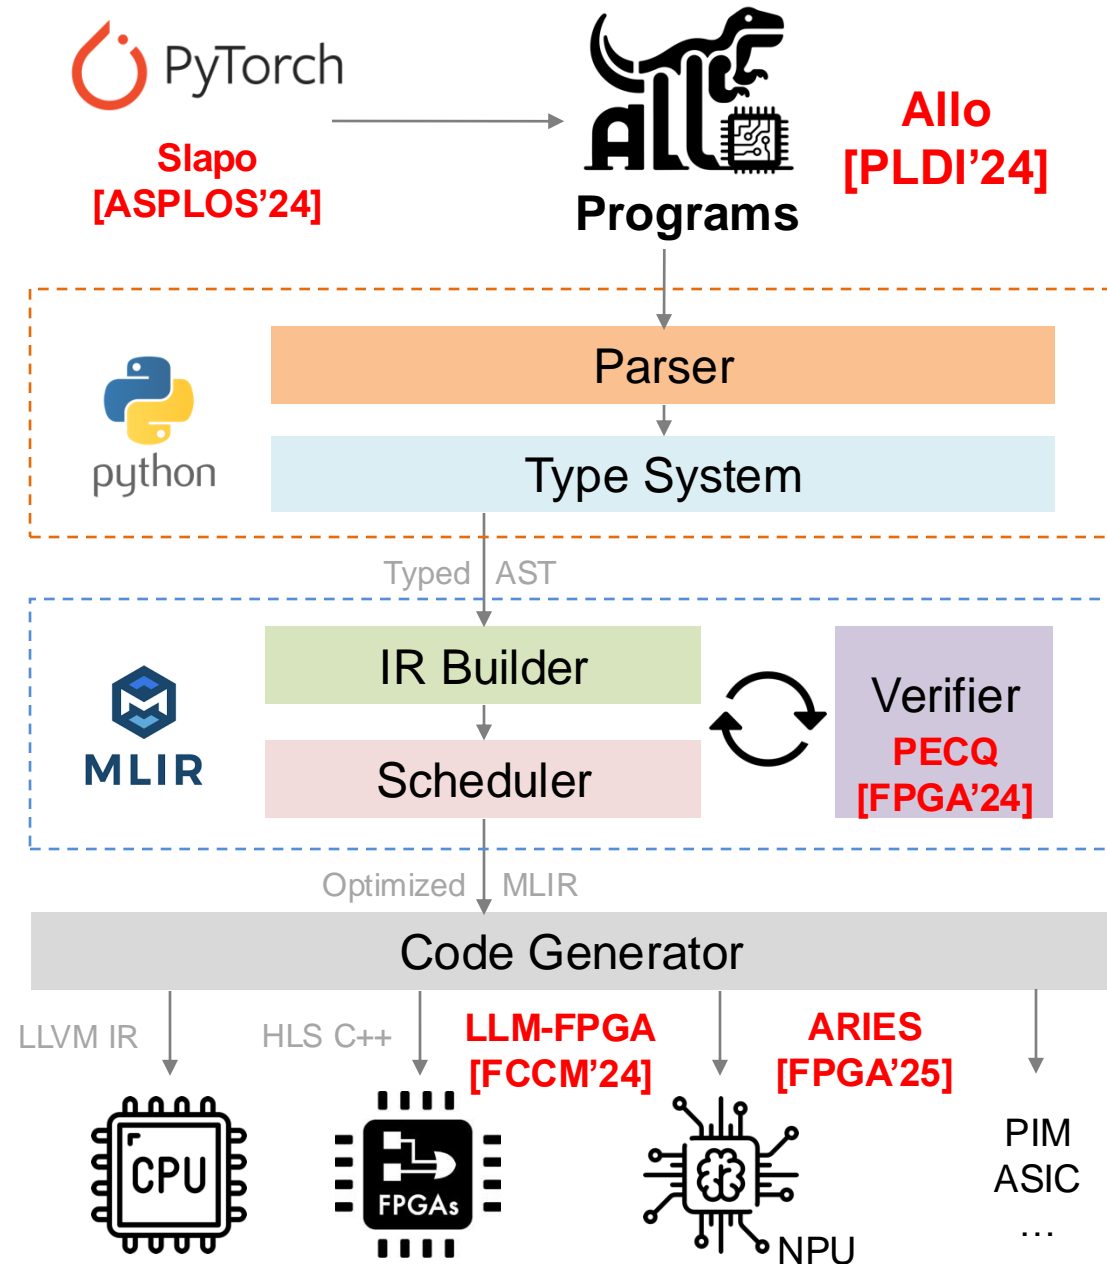
- + Data movement optimization (~95% peak)
  - + Data streaming
  - + Data packing (vectorization)
  - + Memory coalescing
  - + Systolic communication

Increasing programming effort




~500 lines of HLS code for a GEMM systolic array  
**vendor-specific, hard to maintain & reuse**

# Allo Accelerator Design Language (ADL) and Compiler





# Allo Accelerator Design Language (ADL) and Compiler



Programs

Parser

Type System

Typed AST

IR Builder

Scheduler



Verifier

Optimized MLIR

Code Generator

LLVM IR

HLS C++



Pythonic: No need to learn a new DSL!

Algorithm  
Specification

```
def matmul(A: float32[M, K],  
           B: float32[K, N])  
    -> float32[M, N]:  
    C: float32[M, N] = 0.0  
    for i, j, k in allo.grid(M, N, K):  
        C[i, j] += A[i, k] * B[k, j]  
    return C
```

- Free-form imperative programming
- Strong type system

# Allo Accelerator Design Language (ADL) and Compiler



Programs

Parser

Type System

Typed AST

IR Builder

Scheduler



Verifier

Optimized MLIR

Code Generator

LLVM IR

HLS C++



Pythonic: No need to learn a new DSL!

Algorithm  
Specification

```
def matmul(A: float32[M, K],  
           B: float32[K, N])  
    -> float32[M, N]:  
    C: float32[M, N] = 0.0  
    for i, j, k in allo.grid(M, N, K):  
        C[i, j] += A[i, k] * B[k, j]  
    return C
```

Decoupled customization

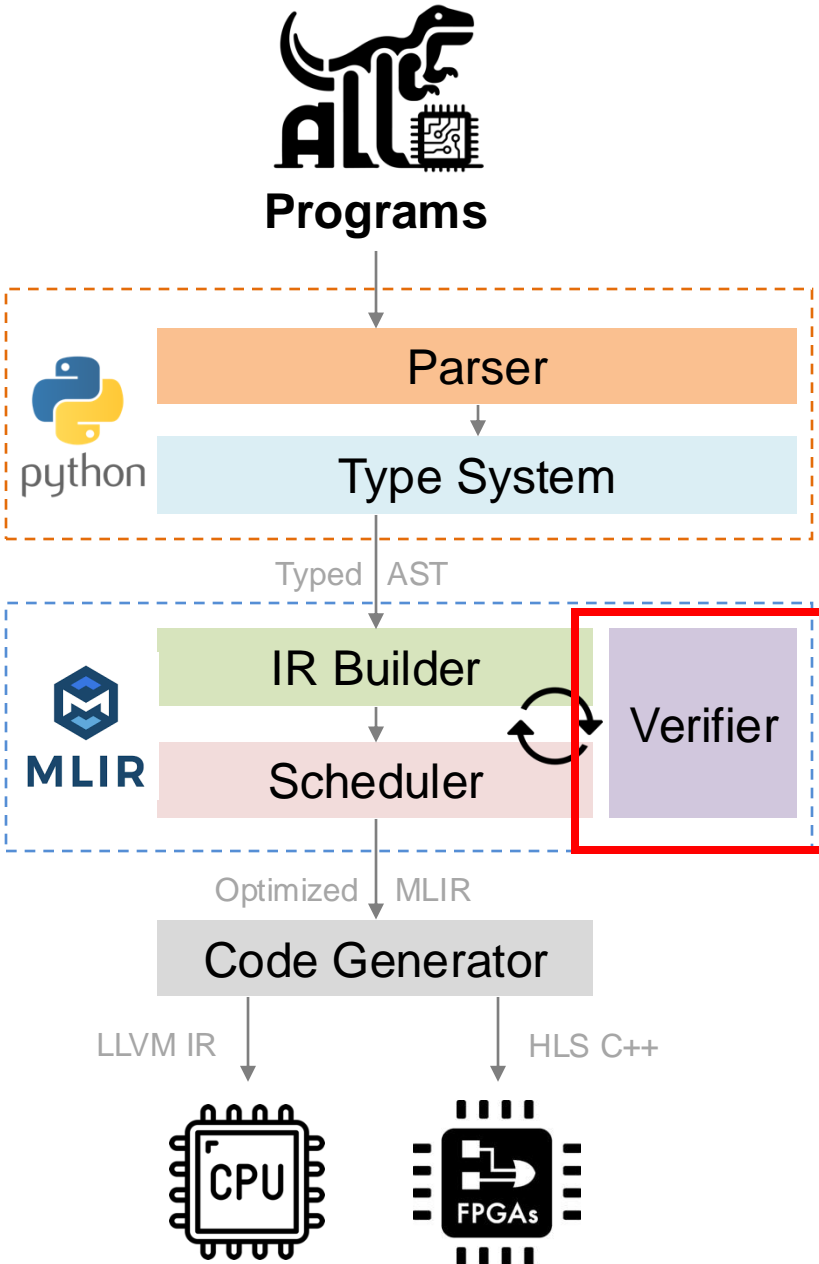
Compute Cust.

Memory Cust.

Comm. Cust.

```
s = allo.customize(matmul)  
s.reorder("k", "j")  
s.buffer_at(s.C, axis="i")  
s.pipeline("j")
```

# Allo Accelerator Design Language (ADL) and Compiler



**Pythonic: No need to learn a new DSL!**

Algorithm  
Specification

```
def matmul(A: float32[M, K],
           B: float32[K, N])
    -> float32[M, N]:
    C: float32[M, N] = 0.0
    for i, j, k in allo.grid(M, N, K):
        C[i, j] += A[i, k] * B[k, j]
    return C
```

**Stepwise verifiable rewrites**

```
s = allo.customize(gemm)
s.reorder("k", "j")
print(s.module)
```

```
s.buffer_at(s.C, axis="i")
print(s.module)
```

```
s.pipeline("j")
print(s.module)
```

```
module {
func.func @gemm(%arg0: memref<1024x1024xf32>, %arg1: memref<1024x1024xf32>) -> memref<1024x1024xf32> {
  %alloc = memref.alloc() {name = "C"} : memref<1024x1024xf32>
  affine.for %arg2 = 0 to 1024 {
    affine.for %arg3 = 0 to 1024 {
      affine.for %arg4 = 0 to 1024 {
        ...
      } {loop_name = "j", pipeline}
    } {loop_name = "k", op_name = "S_k_0", reduction}
    ...
  } {loop_name = "i", op_name = "S_i_j_0"}
  return %alloc : memref<1024x1024xf32>
}
```

```
module {
func.func @gemm(%arg0: memref<1024x1024xf32>, %arg1: memref<1024x1024xf32>) -> memref<1024x1024xf32> {
  %alloc = memref.alloc() {name = "C"} : memref<1024x1024xf32>
  affine.for %arg2 = 0 to 1024 {
    affine.for %arg3 = 0 to 1024 {
      affine.for %arg4 = 0 to 1024 {
        ...
      } {loop_name = "k"}
    } {loop_name = "j"}
  } {loop_name = "i", op_name = "S_i_j_0"}
  return %alloc : memref<1024x1024xf32>
}
```

MLIR

```
module {
func.func @gemm(%arg0: memref<1024x1024xf32>, %arg1: memref<1024x1024xf32>) -> memref<1024x1024xf32> {
  %alloc = memref.alloc() {name = "C"} : memref<1024x1024xf32>
  %alloc_0 = memref.alloc() : memref<1024xf32>
  affine.for %arg3 = 0 to 1024 {
    ...
  } {buffer, loop_name = "j_init", pipeline_ii = 1 : i32}
  affine.for %arg3 = 0 to 1024 {
    affine.for %arg4 = 0 to 1024 {
      ...
    } {loop_name = "j"}
  } {loop_name = "k", op_name = "S_k_0", reduction}
  affine.for %arg3 = 0 to 1024 {
    ...
  } {buffer, loop_name = "j_back", pipeline_ii = 1 : i32}
  } {loop_name = "i", op_name = "S_i_j_0"}
  return %alloc : memref<1024x1024xf32>
}
```

# Transforming GEMM to Systolic Array

## → Algorithm specification


```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N]):  
    for i, j in allo.grid(M, N, "PE"):  
        for k in range(K):  
            C[i, j] += A[i, k] * B[k, j]
```



## → Schedule construction

```
s = allo.customize(matmul)
```

## → Architectural Diagram



\* Schedule: A sequence of customization primitives

# Transforming GEMM to Systolic Array

## → Algorithm specification

```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N]):  
    for i, j in allo.grid(M, N, "PE"):  
        for k in range(K):  
            C[i, j] += A[i, k] * B[k, j]
```

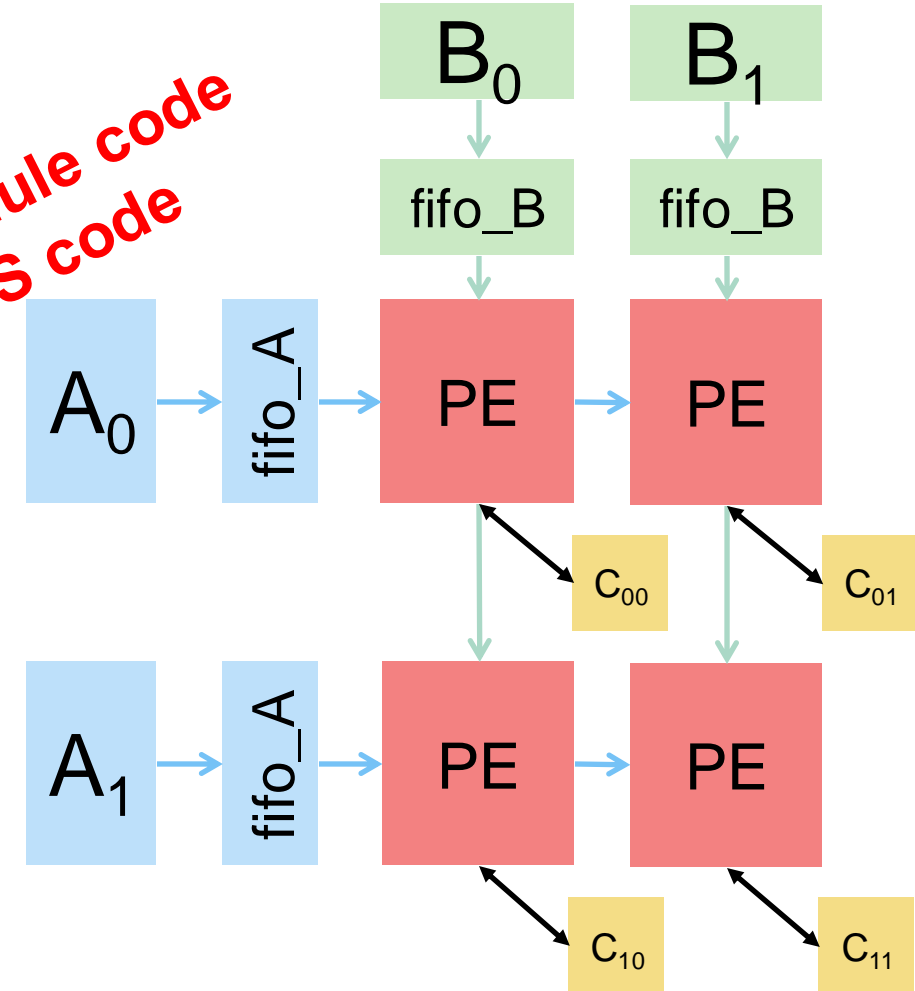


Only 8 lines of schedule code  
vs ~500 lines HLS code

## → Schedule construction

```
s = allo.customize(matmul)  
buf_A = s.buffer_at(s.A, "j")  
buf_B = s.buffer_at(s.B, "j")  
pe = s.unfold("PE", axis=[0, 1],  
              factor=[M, N])  
s.partition(s.A, dim=0)  
s.partition(s.B, dim=1)  
s.partition(s.C, dim=[0, 1])  
s.to(buf_A, pe, axis=0, depth=M + 1)  
s.to(buf_B, pe, axis=1, depth=N + 1)
```

## → Architectural Diagram



# Transforming GEMM to Systolic Array

## → Algorithm specification

```
def matmul(A: int8[M, K],
           B: int8[K, N],
           C: int16[M, N]):
    for i, j in allo.grid(M, N, "PE"):
        for k in range(K):
            C[i, j] += A[i, k] * B[k, j]
```




## → Schedule construction

```
s = allo.customize(matmul)
buf_A = s.buffer_at(s.A, "j")
buf_B = s.buffer_at(s.B, "j")
pe = s.unfold("PE", axis=[0, 1],
              factor=[M, N])
s.partition(s.A, dim=0)
s.partition(s.B, dim=1)
s.partition(s.C, dim=[0, 1])
s.to(buf_A, pe, axis=0, depth=M + 1)
s.to(buf_B, pe, axis=1, depth=N + 1)
```

## → CPU Simulation

```
mod = s.build(target="llvm")


A = np.random.randint(-8, 8, size=(M, K)) \
    .astype(np.int8)
B = np.random.randint(-8, 8, size=(K, N)) \
    .astype(np.int8)
C = np.zeros((M, N), dtype=np.int16)
mod(A, B, C)
np.testing.assert_allclose(C, A @ B, atol=1e-3)
```



## → Bitstream generation

```
mod = s.build(target="vitis_hls",
               mode="hw",
               project="systolic.prj")
```

*# Automatically invoke the HLS toolchain*  
mod(A, B, C)



# Transforming GEMM to Systolic Array

## → Algorithm specification

```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N]):  
    for i, j in allo.grid(M, N, "PE"):  
        for k in range(K):  
            C[i, j] += A[i, k] * B[k, j]
```





## → Schedule construction

```
s = allo.customize(matmul)  
buf_A = s.buffer_at(s.A, "j")  
buf_B = s.buffer_at(s.B, "j")  
pe = s.unfold("PE", axis=[0, 1],  
             factor=[M, N])  
s.partition(s.A, dim=0)  
s.partition(s.B, dim=1)  
s.partition(s.C, dim=[0, 1])  
s.to(buf_A, pe, axis=0, depth=M + 1)  
s.to(buf_B, pe, axis=1, depth=N + 1)
```

## → Formal Verification [FPGA'24 Best Paper\*]

```
mod = s.verify()
```

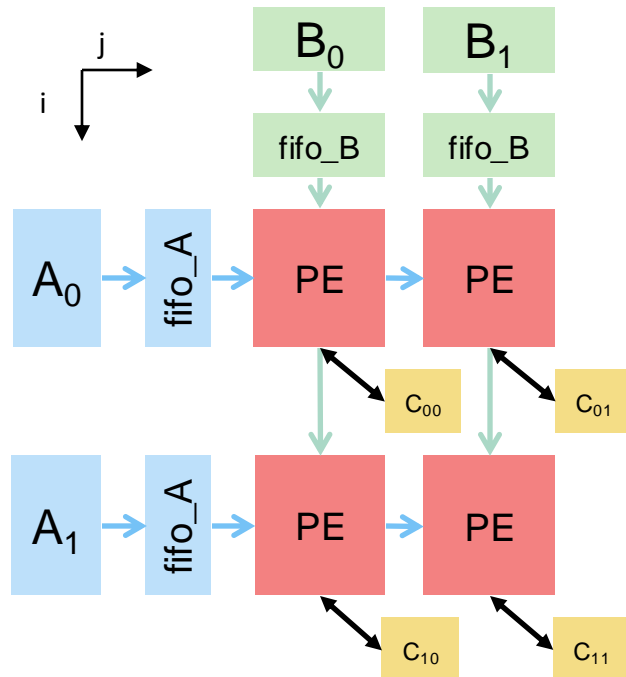


- A **formal equivalence checker** of source-to-source HLS transformations via symbolic execution
  - Support statically interpretable control-flow (SICF)
- Verification in time/space linear w.r.t. OPs executed
  - ~500K Ops/sec in verification throughput
  - A complex **64x64 systolic array verified in 16 minutes**

# Transforming GEMM to Systolic Array


## → Spatial loop: i, j

```
s = allo.customize(matmul)
buf_A = s.buffer_at(s.A, "j")
buf_B = s.buffer_at(s.B, "j")
pe = s.unfold("PE", axis=[0, 1],
             factor=[M, N])
s.partition(s.A, dim=0)
s.partition(s.B, dim=1)
s.partition(s.C, dim=[0, 1])
s.to(buf_A, pe, axis=0, depth=M + 1)
s.to(buf_B, pe, axis=1, depth=N + 1)
```




## → Spatial loop: k, j

```
s = allo.customize(matmul)
s.reorder("k", "j", "i")
buf_A = s.buffer_at(s.A, "j")
buf_C = s.buffer_at(s.C, "j")
pe = s.unfold("PE", axis=[0, 1],
             factor=[K, N])
s.partition(s.A, dim=0)
s.partition(s.B, dim=[0, 1])
s.partition(s.C, dim=1)
s.to(buf_A, pe, axis=0)
s.to(buf_C, pe, axis=1)
```



## → Spatial loop: j


```
s = allo.customize(matmul)
buf_A = s.buffer_at(s.A, "j")
buf_B = s.buffer_at(s.B, "j")
pe = s.unfold("PE", axis=1,
             factor=M)
s.partition(s.B, dim=1)
s.partition(s.C, dim=1)
s.to(buf_A, pe, axis=0)
s.to(buf_B, pe, axis=1)
```




**Realize different dataflows  
with minimal schedule code**



# How to Compose Optimized Kernels into Complete Accelerator?



**Optimized  
Kernels**



**High-performance  
Accelerator**

# Composable Schedules

→ Given different optimized kernel implementations

```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N])
```



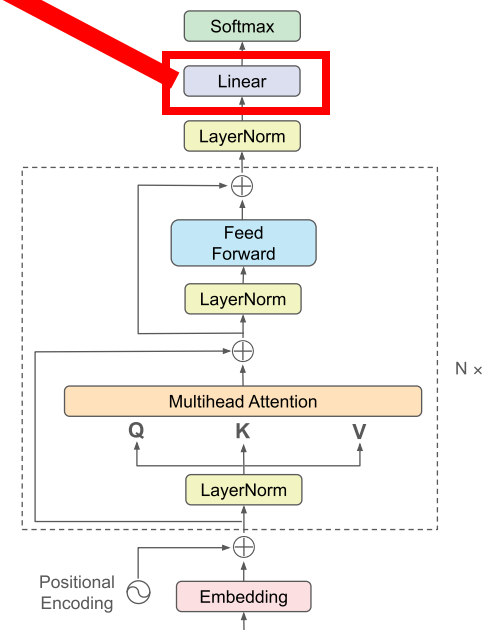
→ Algorithm specification (Hierarchical)

```
def top(X: int8[M, K],  
       W_A: int8[K, N],  
       W_B: int8[N, K]):  
    Y: int8[M, N] = 0  
    Z: int8[M, K] = 0  
    matmul(X, W_A, Y)  
    matmul(Y, W_B, Z)  
    return Z
```


→ Schedule composition

```
# Previous customizations for matmul  
s_matmul = allo.customize(matmul)  
# ...
```


```
s_top = allo.customize(top)  
s_top.compose(s_matmul)
```




# Interconnection between Kernels




Temporal Composition



```
def top(
  X: int8[M, K],
  W_A: int8[K, N],
  W_B: int8[N, K]):
  Y: ?
  Z: int8[M, K] = 0
  matmul(X, W_A, Y)
  matmul(Y, W_B, Z)
  return Z
```



Spatial Composition



# Temporal Composition

→ Given optimized kernel implementations

```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N])
```



Mismatched  
interface

`s2.partition(s.A, dim=1)`



`s1.partition(s.C, dim=0)`



Matched  
interface!




`s2.partition(s.A, dim=0)`




→ Algorithm specification (Hierarchical)

```
def top(X: int8[M, K], W_A: int8[K, N],  
       W_B: int8[N, K]):  
  Y: int8[M, N] = 0  
  Z: int8[M, K] = 0  
  matmul(X, W_A, Y)  
  matmul(Y, W_B, Z)  
  return Z
```



Need to ensure  
interface consistency

# Temporal Composition



Hierarchical use-def graph

## → Algorithm specification (Hierarchical)

```
def top(X: int8[M, K], W_A: int8[K, N],  
        W_B: int8[N, K], Y: int8[M, K]):  
    Y: int8[M, N] = 0  
    Z: int8[M, K] = 0  
    matmul(X, W_A, Y) ← Caller definition  
    matmul(Y, W_B, Z)  
    return Z
```


## → Callee definition

```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N])
```

Need to ensure interface consistency




# Temporal Composition



Hierarchical use-def graph

- **Key idea: Model memory banking as a type** to ensure kernel interfaces are consistent

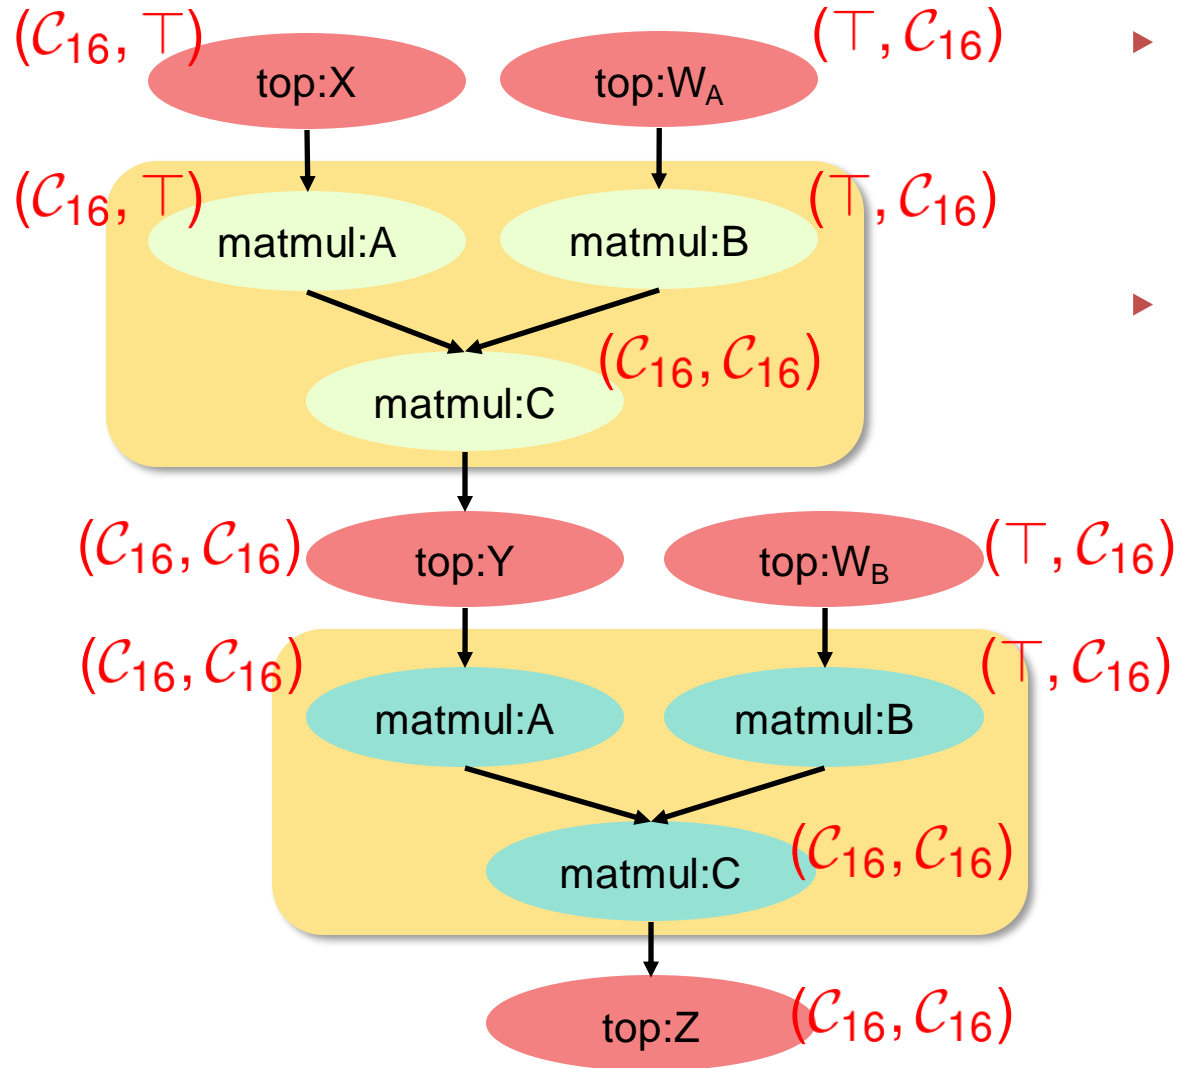
Subtyping relation forms a **lattice**!



**Intuition:**

We can supply more read/write parallelism, but not less!

# Temporal Composition

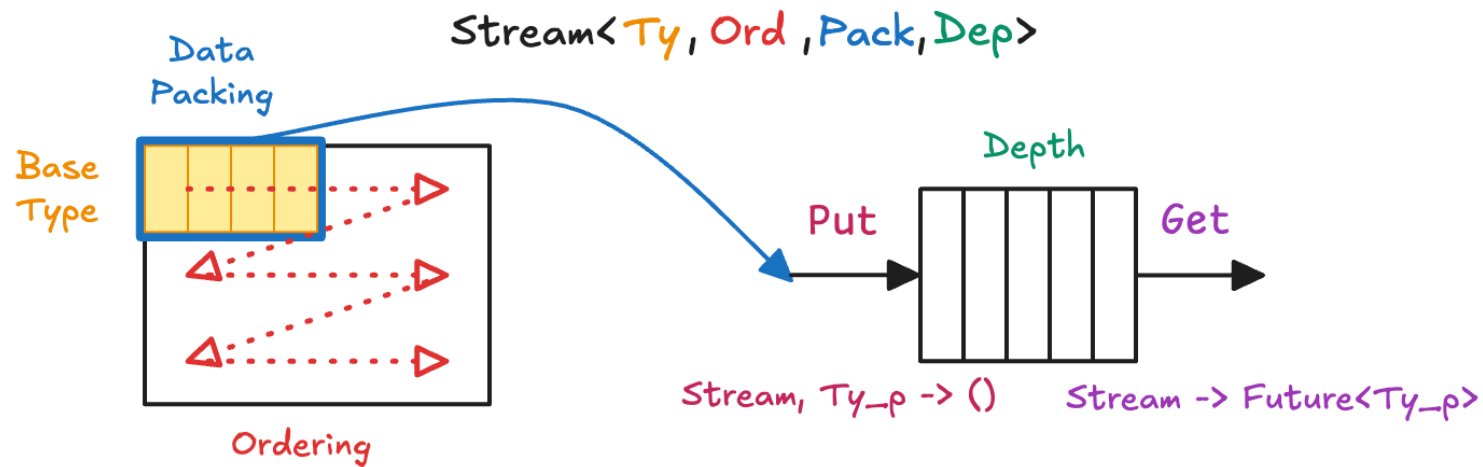


Hierarchical use-def graph

- **Key idea: Model data layout as a type** to ensure the kernel interfaces are consistent
- Lattice-based subtyping relation permits **linear-time layout inference** through Worklist algorithm

# Spatial Composition

## ► Model data streaming as a type



```
@df.region
def top():
    pipe: Stream[Ty[M, N],
                "R", 1, 4]
```

```
@df.kernel(mapping=[1])
def producer(A: Ty[M, N]):
    for i, j in allo.grid(M, N):
        # send data
        pipe.put(A[i, j])
```

```
@df.kernel(mapping=[1])
def consumer(B: Ty[M, N]):
    for i, j in allo.grid(M, N):
        # receive data
        B[i, j] = pipe.get() + 1
```



# Spatial Composition

```
@df.region()
```

```
def top():
```

```
    fifo_A = df.array(df.pipe(dtype=float32, depth=4), shape=(P0, P1))
```

```
    fifo_B = df.array(df.pipe(dtype=float32, depth=4), shape=(P0, P1))
```

```
@df.kernel(mapping=[P0, P1])
```

```
def gemm(A: float32[M, K],
```

```
        B: float32[K, N],
```

```
        C: float32[M, N]):
```

```
    i, j = df.get_pid()
```

```
    # peripheral kernels
```

```
    with allo.meta_if(i in {0, M + 1}
                      and j in {0, N + 1}):
```

```
        pass
```

```
    with allo.meta_elif(j == 0):
```

```
        # i > 0
```

```
        for k in range(K):
```


```
            fifo_A[i, j + 1].put(A[i - 1, k])
```

```
    with allo.meta_elif(i == 0):
```

```
        # j > 0
```

```
        for k in range(K):
```

```
            fifo_B[i + 1, j].put(B[k, j - 1])
```



# Spatial Composition

```
@df.region()
```

```
def top():
```

```
    fifo_A = df.array(df.pipe(dtype=float32, depth=4), shape=(P0, P1))
```

```
    fifo_B = df.array(df.pipe(dtype=float32, depth=4), shape=(P0, P1))
```

```
# drain
```

```
with allo.meta_elif(i == M + 1 and j > 0):
```

```
    for k in range(K):
```

```
        b: float32 = fifo_B[i, j].get()
```

```
with allo.meta_elif(j == N + 1 and i > 0):
```

```
    for k in range(K):
```

```
        a: float32 = fifo_A[i, j].get()
```

```
# main body
```

```
with allo.meta_else():
```

```
    c: float32 = 0
```

```
    for k in range(K):
```

```
        a: float32 = fifo_A[i, j].get()
```


```
        b: float32 = fifo_B[i, j].get()
```

```
        c += a * b
```

```
        fifo_A[i, j + 1].put(a)
```


```
        fifo_B[i + 1, j].put(b)
```

```
        C[i - 1, j - 1] = c
```



# Single-Kernel Evaluation

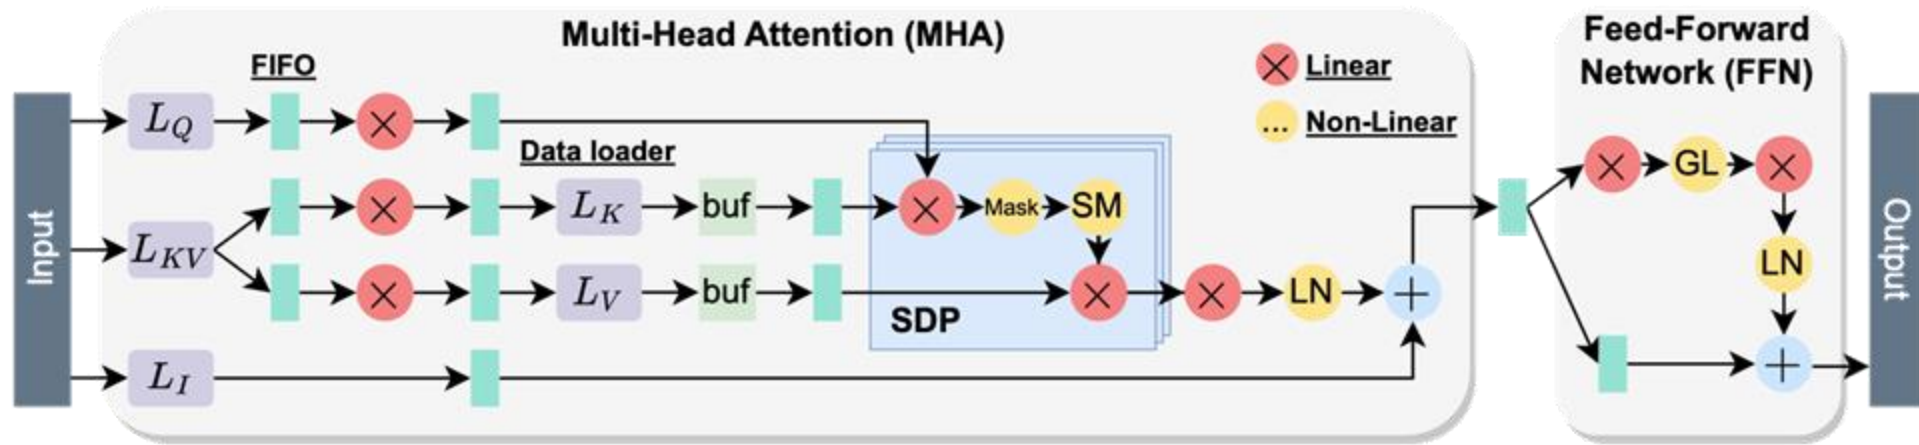
- ▶ Benchmarks from PolyBench; Target hardware: AMD U280 FPGA
- ▶ Normalized against AMD VitisHLS-auto (pragmas automatically inserted)
- ▶ Other baselines
  - ADLs: HeteroCL [FPGA'19], Dahlia [PLDI'20], PyLog [TC'21]
  - Automated DSE for HLS: ScaleHLS [HPCA'22], Merlin [TRETS'22]



Allo achieves (much) higher performance by optimizing data placement with a custom memory hierarchy

# Multi-Kernel Evaluation: A Complete LLM Accelerator

- ▶ GPT2 model (the only open-source LLM in the GPT family)
  - 355M parameters, 24 hidden layers, 16 heads
  - W4A8 quantization




```
def GPT_layer(inp: float32[L, D], ...)
    -> float32[L, D]:
    # 1. Multi-Head Attention (MHA)
    Q = Linear_layer_qkv(inp, Wq, Bq)
    K = Linear_layer_qkv(inp, Wk, Bk)
    V = Linear_layer_qkv(inp, Wv, Bv)
    attn_sf_outp = Self_attention(Q, K, V)
    # ...
    # 2. Feed Forward Network (FFN)
    # ...
    return ff_n_res_outp
```

Compose all the schedules together

```
s = allo.customize(GPT_layer)
s.compose([s_qkv, ..., s_gelu])
```

# LLM Accelerator Evaluation

- ▶ GPT2: single-batch, low-latency, generative inference settings
  - AMD U280 FPGA (16nm), 250MHz
  - **2.2x speedup in prefill stage** compared to DFX [MICRO'22] (an FPGA-based overlay)
  - **1.9x speedup for long output** sequences and **5.7x more energy-efficient** vs. NVIDIA A100 GPU (7nm)
  - Fewer than 50 lines of schedule code in Allo



	Allo	DFX
Device	U280	U280
Freq.	250MHz	200MHz
Quant.	W4A8	fp16
BRAM	384 (19.0%)	1192 (59.1%)
DSP	1780 (19.73%)	3533 (39.2%)
FF	652K (25.0%)	1107K (42.5%)
LUT	508K (39.0%)	520K (39.9%)

# High-Level PyTorch Frontend

- ▶ Predefined schedules for commonly used NN operators
- ▶ Can directly import model from PyTorch and build optimized xcel design
  - Through TorchDynamo and torch.fx

```
import torch
import allo
import numpy as np
from transformers import AutoConfig
from transformers.models.llama.modeling_llama import LlamaModel

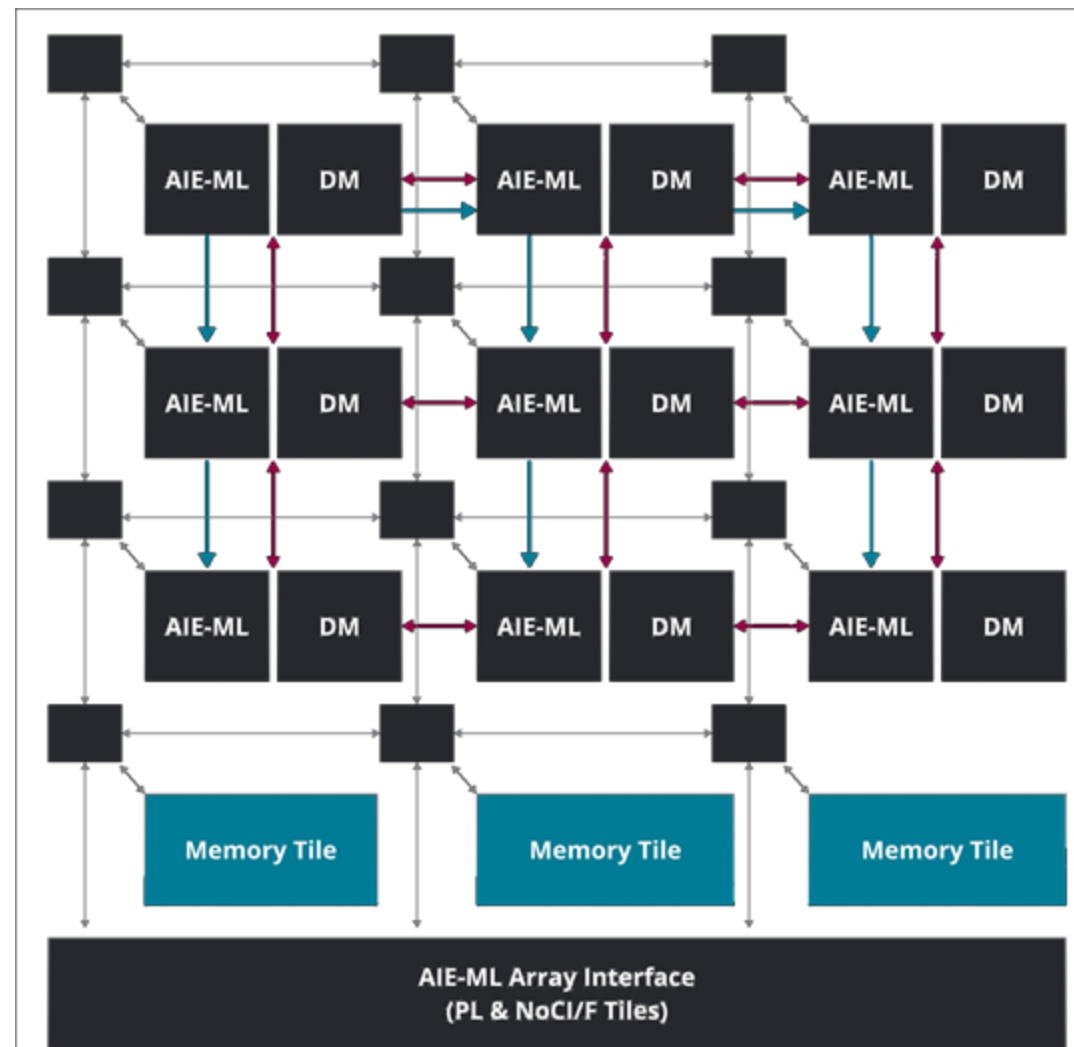

bs, seq, hs = 1, 512, 1024

example_inputs = [torch.rand(bs, seq, hs)]
config = AutoConfig.from_pretrained("meta-llama/Meta-Llama-3-8B")
module = GPT2Model(config).eval()
mlir_mod = allo.frontend.from_pytorch(
    module,
    example_inputs=example_inputs,
)
```

# Other Backends: AMD AI Engine Evaluation


- ▶ GEMM on AMD Ryzen 7940 SoC (CPU + AIE NPU)
  - Orders of magnitude improvement using the streaming interface

```
mod = df.build(target="aie")
```



## Other Backends: AMD AI Engine Evaluation

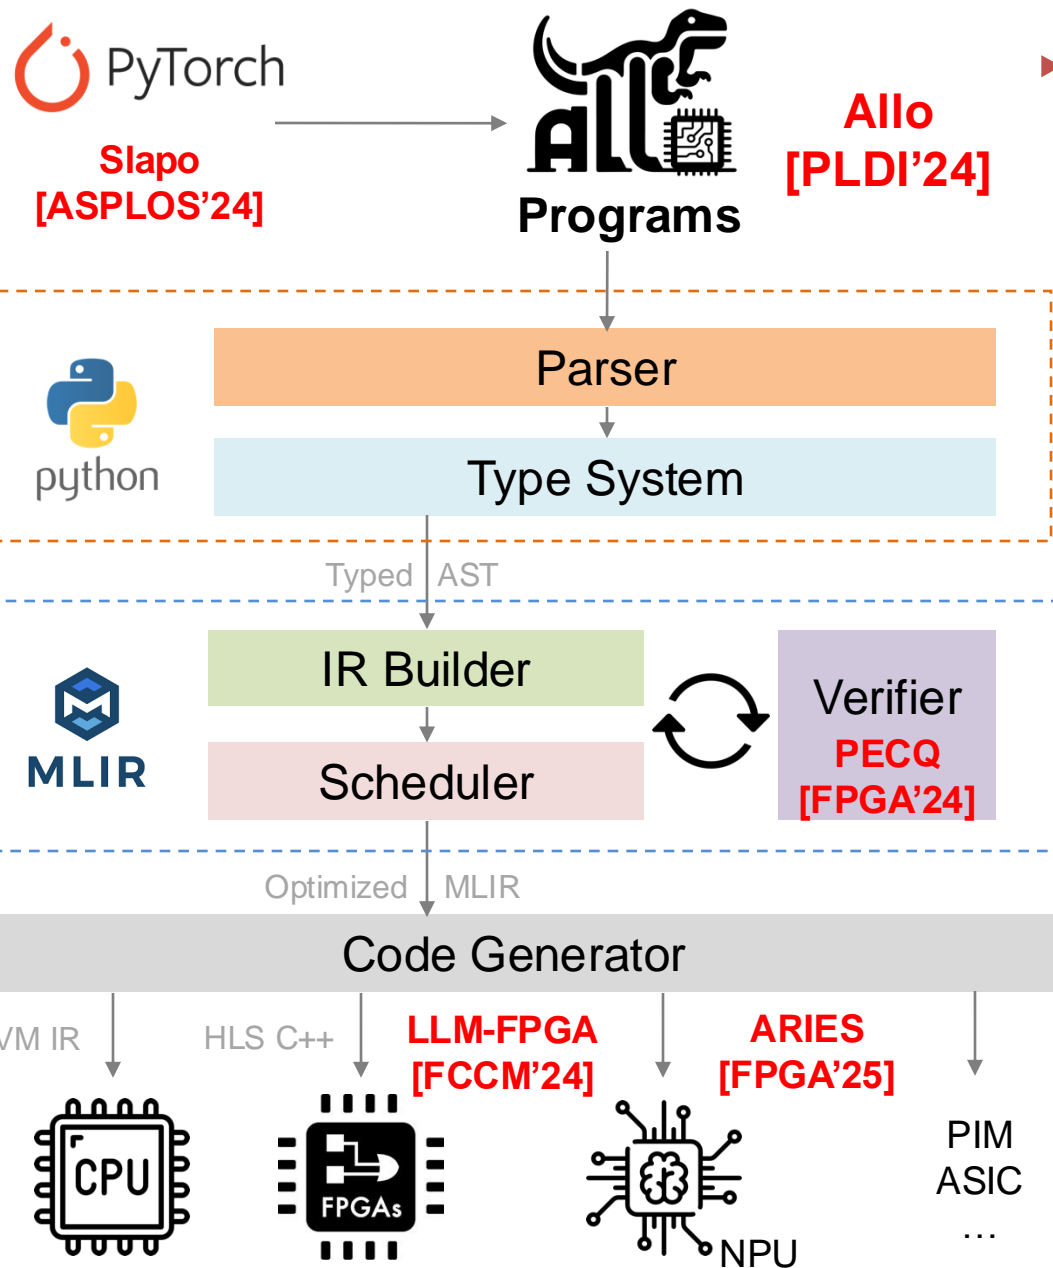
- ▶ ResNet on on AMD Ryzen 7940 SoC (CPU + AIE NPU)
  - 23x speedup compared to AMD Riallto
- ▶ LLaMA models being evaluated



Designs	Tile	SL	IP	Util (%)	RT (ms)	Speedup
D1 scalar	[16, 16]	1	No	15	57.40	1.24
D2 +vectorized	[16, 16]	4	No	15	16.63	4.29
D3 +vectorized	[16, 16]	8	No	15	9.82	7.27
D4 +conv3×3-2core	[16, 16]	8	No	20	9.18	7.77
D5 +inst-pipeline	[16, 16]	8	Yes	20	8.44	8.45
D6 +opt-tile-size	[32, 16]	8	Yes	20	5.72	13.70
D7 +opt-tile-size	[32, 32]	8	Yes	20	5.21	12.48
D8 +opt-tile-size	[32, 64]	8	Yes	20	4.99	14.30
D9 +more-cores	[32, 64]	8	Yes	40	3.16	22.58
Riallto / RAI-SW[20]	-	-	-	20	71.36	1x



# Summary



## ► Features of Allo ADL

- Pythonic
- Decoupled & verifiable customizations
- Composability



<https://github.com/cornell-zhang/allo>



Cornell University



UNIVERSITY OF ILLINOIS  
URBANA-CHAMPAIGN



BROWN



UNIVERSITY OF TORONTO



UNIVERSITY OF VIRGINIA



UNIVERSITY OF ILLINOIS CHICAGO



上海交通大学  
SHANGHAI JIAO TONG UNIVERSITY