


# Python-Based Accelerator Design

Hongzheng Chen, Niansong Zhang, Shaojie Xiang, Zhiru Zhang

Cornell University

FPGA'24 Tutorial



Cornell University

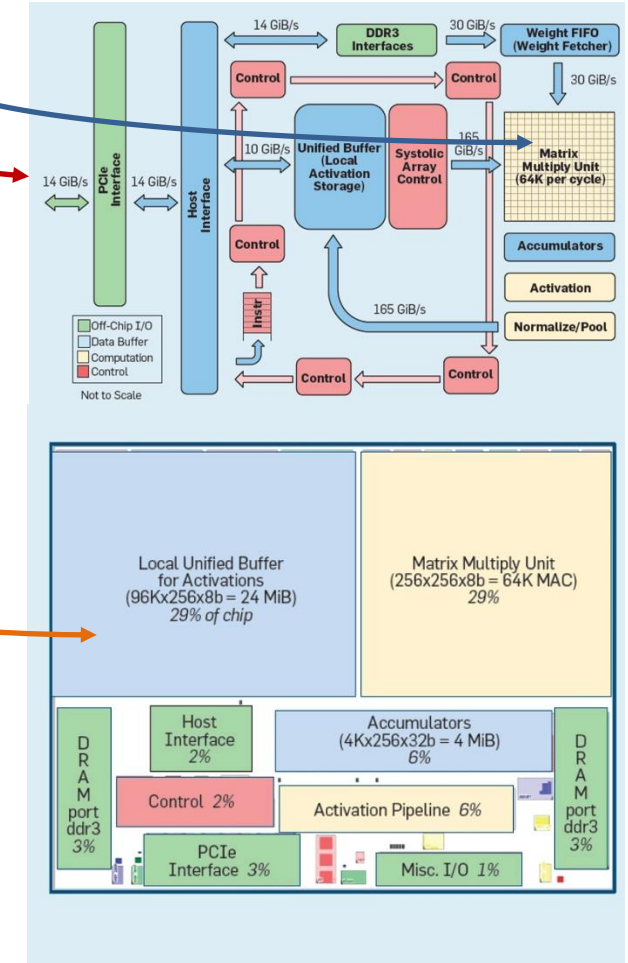



# Outline

- ▶ Part I: Introduction (40min)
  - Background
  - Decoupled customizations
  - Large-scale composability
- ▶ Part II: Hands-on tutorial (90min)
  - Basic syntax
  - Optimization guide
  - Verification
  - Multi-kernel composition
  - Different targets (CPU/FPGA/AIE)
- ▶ Part III: Free exploration and Q&A

# Complexity in Specialized Accelerator Design

- ▶ Accelerator design is different from programming on general processors
  - Custom processing engines (PEs)
  - Custom data communication
  - Custom memory hierarchy



An AI accelerator example

# Challenge: Balancing Manual Control & Compiler Optimization

```
void systolic_tile(int8_t A_tile[2][768],
int8_t B_tile[768][2],
int8_t C_tile[2][2]) {
#pragma dataflow
hls::stream<int8_t> A_fifo[2][3], B_fifo[2][3];
#pragma stream variable=A/B_fifo depth=3
#pragma partition variable=A/B/C_tile complete dim=1
for (int k4 = 0; k4 < 768; k4++) {
for (int m = 0; m < 2; m++) {
int8_t v105 = A_tile[m][k4];
A_fifo[m][0].write(v105);
// ... write B_fifo
}
for (int Ti = 0; Ti < 2; ++Ti) {
#pragma HLS unroll
for (int Tj = 0; Tj < 2; ++Tj) {
#pragma HLS unroll
// ... load A/B_fifo
PE_kernel(A_in, A_out, B_in, B_out, C, Ti, Tj);
}
}
}
}
```

Compute  
Customization

```
void matmul(int8_t A[512][768], int8_t B[768][768],
int8_t C[512][768]) {
int8_t local_A[2][768], local_B[768][2], local_C[2][2];
for (int mi = 0; mi < 256; mi++) {
for (int ni = 0; ni < 384; ni++) {
// ... load A, B
systolic_tile(local_A, local_B, local_C);
}
}
}
```



Vanilla Matmul (1% theoretical peak perf.)

+ Compute customization (~30% peak)

- + Loop tiling
- + Loop unrolling
- + Loop & function pipelining

Increasing programming effort

Memory



Accelerator

# Challenge: Balancing Manual Control & Compiler Optimization

```
void systolic_tile(int8_t A_tile[2][768],
int8_t B_tile[768][2],
int8_t C_tile[2][2]) {
#pragma dataflow
hls::stream<int8_t> A_fifo[2][3], B_fifo[2][3];
#pragma stream variable=A/B_fifo depth=3
#pragma partition variable=A/B/C_tile complete dim=1
for (int k4 = 0; k4 < 768; k4++) {
for (int m = 0; m < 2; m++) {
int8_t v105 = A_tile[m][k4];
A_fifo[m][0].write(v105);
// ... write B_fifo
}}
for (int Ti = 0; Ti < 2; ++Ti) {
#pragma HLS unroll
for (int Tj = 0; Tj < 2; ++Tj) {
#pragma HLS unroll
// ... load A/B_fifo
PE_kernel(A_in, A_out, B_in, B_out, C, Ti, Tj);
}}}
}
```

**Memory  
Customization**

**Compute  
Customization**

```
void matmul(int8_t A[512][768], int8_t B[768][768],
int8_t C[512][768]) {
int8_t local_A[2][768], local_B[768][2], local_C[2][2];
for (int mi = 0; mi < 256; mi++) {
for (int ni = 0; ni < 384; ni++) {
// ... load A, B
systolic_tile(local_A, local_B, local_C);
}}}
}
```

**Vanilla Matmul (1% theoretical peak perf.)**

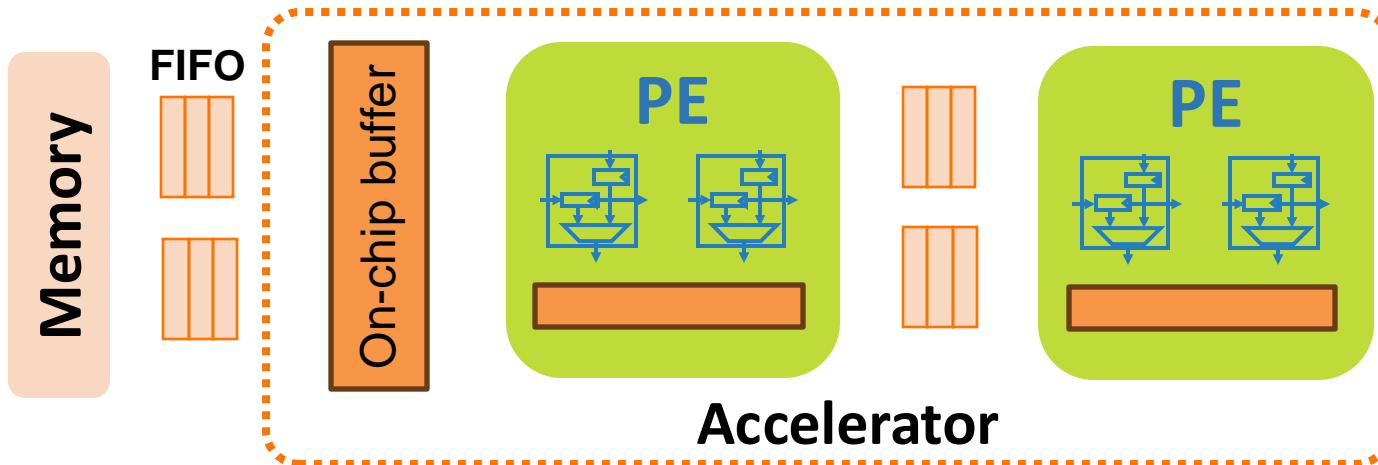
+ Compute customization (~30% peak)

- + Loop tiling
- + Loop unrolling
- + Loop & function pipelining

+ Custom memory hierarchy (~50% peak)

- + Tiling & data reuse buffers
- + Memory banking/partitioning

Increasing programming effort



# Challenge: Balancing Manual Control & Compiler Optimization

```
void systolic_tile(int8_t A_tile[2][768],
int8_t B_tile[768][2],
int8_t C_tile[2][2]) {
#pragma dataflow
hls::stream<int8_t> A_fifo[2][3], B_fifo[2][3];
#pragma stream variable=A/B_fifo depth=3
#pragma partition variable=A/B/C_tile complete dim=1
for (int k4 = 0; k4 < 768; k4++) {
for (int m = 0; m < 2; m++) {
int8_t v105 = A_tile[m][k4];
A_fifo[m][0].write(v105);
// ... write B_fifo
}
}
for (int Ti = 0; Ti < 2; ++Ti) {
#pragma HLS unroll
for (int Tj = 0; Tj < 2; ++Tj) {
#pragma HLS unroll
// ... load A/B_fifo
PE_kernel(A_in, A_out, B_in, B_out, C, Ti, Tj);
}
}
}
```

**Communication  
Customization**

**Memory  
Customization**

**Compute  
Customization**

```
void matmul(int8_t A[512][768], int8_t B[768][768],
int8_t C[512][768]) {
int8_t local_A[2][768], local_B[768][2], local_C[2][2];
for (int mi = 0; mi < 256; mi++) {
for (int ni = 0; ni < 384; ni++) {
// ... load A, B
systolic_tile(local_A, local_B, local_C);
}
}
}
```

**Vanilla Matmul (1% theoretical peak perf.)**

+ Compute customization (~30% peak)

- + Loop tiling
- + Loop unrolling
- + Loop & function pipelining


+ Custom memory hierarchy (~50% peak)

- + Tiling & data reuse buffers
- + Memory banking/partitioning

+ Data movement optimization (~95% peak)

- + Data streaming
- + Data packing (vectorization)
- + Memory coalescing
- + Systolic communication

Increasing programming effort



# Challenge: Balancing Manual Control & Compiler Optimization


```
void systolic_tile(int8_t A_tile[2][768],
int8_t B_tile[768][2],
int8_t C_tile[2][2]) {
#pragma a dataflow
hls::stream<int8_t> A_fifo[2][3], B_fifo[2][3];
#pragma stream variable=A/B_fifo depth=3
#pragma partition variable=A/B/C_tile complete dim=1
for (int k4 = 0; k4 < 768; k4++) {
for (int m = 0; m < 2; m++) {
int8_t v105 = A_tile[m][k4];
A_fifo[m][0].write(v105);
// ... write B_fifo
}}
for (int Ti = 0; Ti < 2; ++Ti) {
#pragma HLS unroll
for (int Tj = 0; Tj < 2; ++Tj) {
#pragma HLS unroll
// ... load A/B_fifo
PE_kernel(A_in, A_out, B_in, B_out, C, Ti, Tj);
}}
}
```

**Communication  
Customization**

**Memory  
Customization**

**Compute  
Customization**

```
void matmul(int8_t A[512][768], int8_t B[768][768],
int8_t C[512][768]) {
int8_t local_A[2][768], local_B[768][2], local_C[2][2];
for (int mi = 0; mi < 256; mi++) {
for (int ni = 0; ni < 384; ni++) {
// ... load A, B
systolic_tile(local_A, local_B, local_C);
}}
}
```



**Vanilla Matmul (1% theoretical peak perf.)**

+ Compute customization (~30% peak)

- + Loop tiling **Existing HLS compiler**
- + Loop unrolling **e.g., ScaleHLS [HPCA'22]**
- + Loop & function pipelining

+ Custom memory hierarchy (~50% peak)

- + Tiling & data reuse buffers
- + Memory banking/partitioning

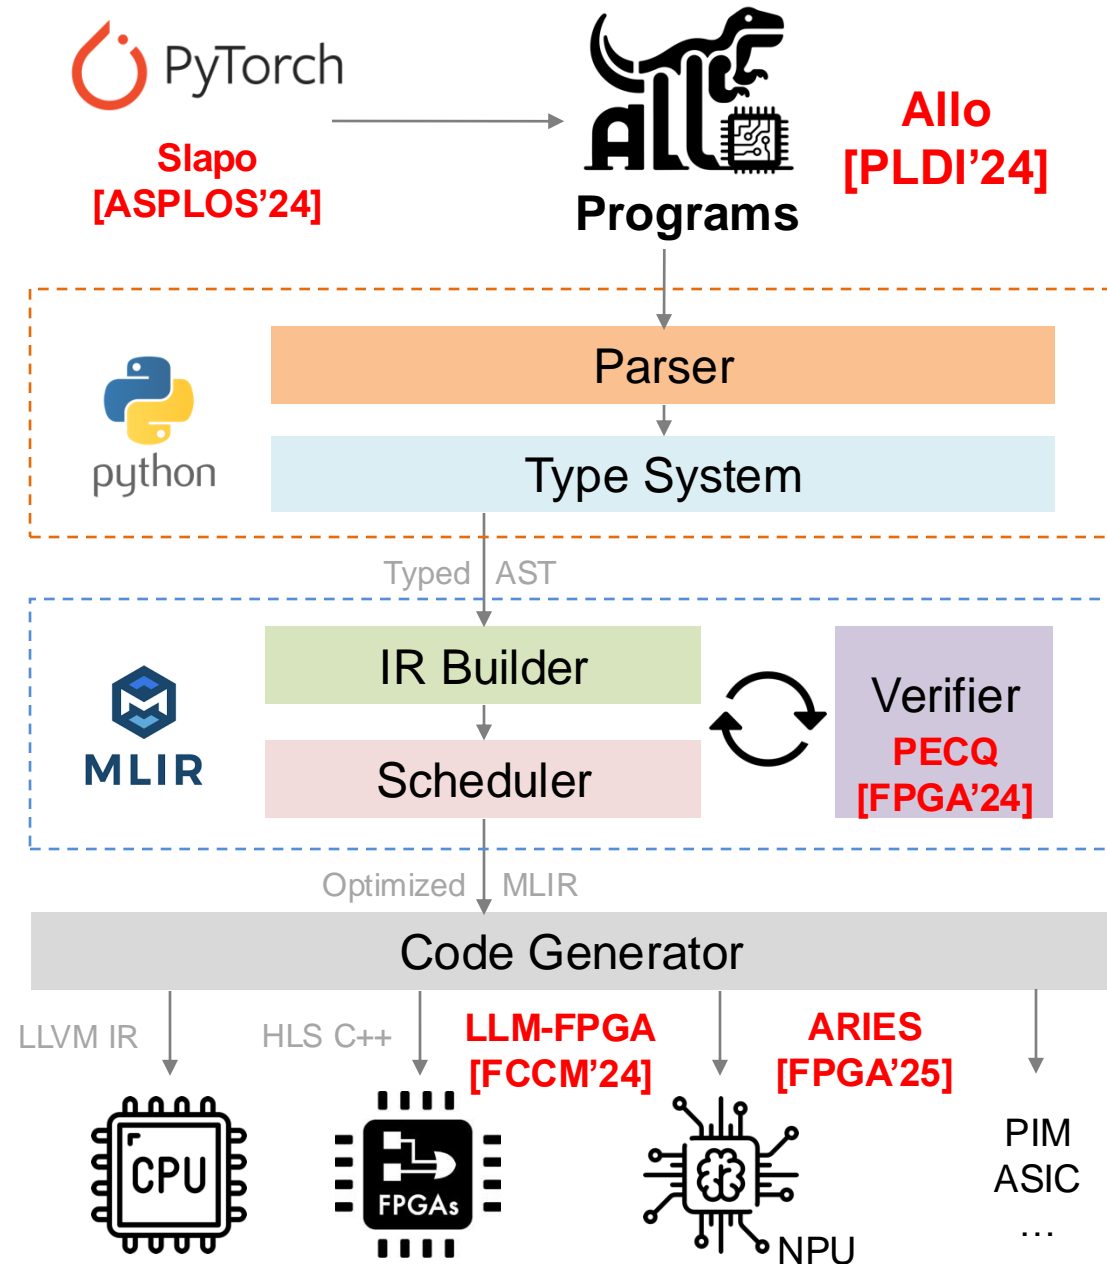
+ Data movement optimization (~95% peak)

- + Data streaming
- + Data packing (vectorization)
- + Memory coalescing
- + Systolic communication

~500 lines of HLS code for a small systolic array  
**vendor-specific, hard to maintain & reuse**

Increasing programming effort

# Allo Accelerator Design Language (ADL) and Compiler





# Allo Accelerator Design Language (ADL) and Compiler



Programs

Parser

Type System

Typed AST

IR Builder

Scheduler



Verifier

Optimized MLIR

Code Generator

LLVM IR

HLS C++



Pythonic: No need to learn a new DSL!

Algorithm  
Specification

```
def matmul(A: float32[M, K],  
           B: float32[K, N])  
    -> float32[M, N]:  
    C: float32[M, N] = 0.0  
    for i, j, k in allo.grid(M, N, K):  
        C[i, j] += A[i, k] * B[k, j]  
    return C
```

Framework	Input	Paradigm
Spatial [PLDI'18]	Scala	Imperative
TVM TE [OSDI'18]	Python	Declarative
<b>Allo</b>	<b>Python</b>	<b>Imperative</b>

# Allo Accelerator Design Language (ADL) and Compiler



Programs

Parser

Type System

Typed AST

IR Builder

Scheduler



Verifier

Optimized MLIR

Code Generator

LLVM IR

HLS C++



Pythonic: No need to learn a new DSL!

Algorithm Specification

```
def matmul(A: float32[M, K],  
           B: float32[K, N])  
    -> float32[M, N]:  
    C: float32[M, N] = 0.0  
    for i, j, k in allo.grid(M, N, K):  
        C[i, j] += A[i, k] * B[k, j]  
    return C
```

Decoupled customization

Compute Cust.

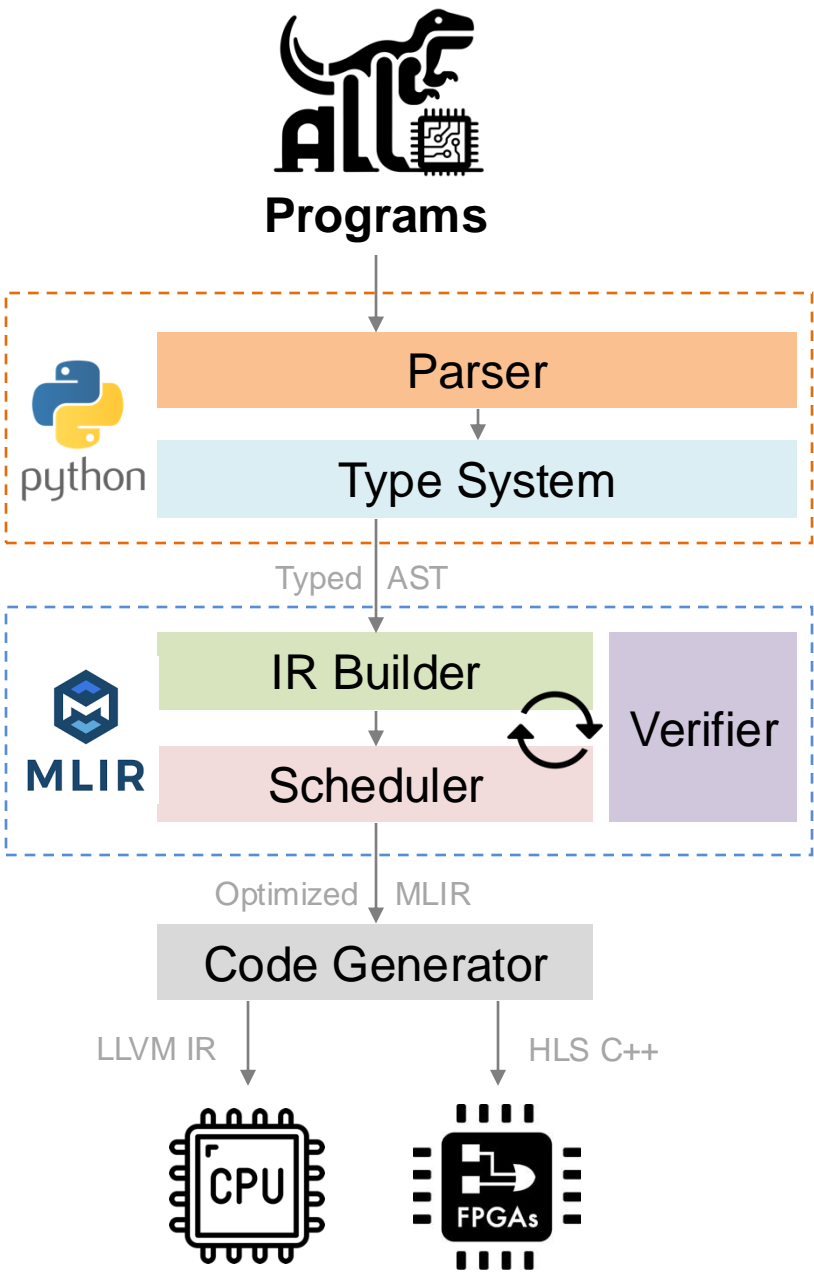
Memory Cust.

Comm. Cust.

```
s = allo.customize(matmul)  
s.reorder("k", "j")  
s.buffer_at(s.C, axis="i")  
s.pipeline("j")
```

Framework	Input	Paradigm	Decoupled Cust.
Spatial [PLDI'18]	Scala	Imperative	X
TVM TE [OSDI'18]	Python	Declarative	✓
<b>Allo</b>	<b>Python</b>	<b>Imperative</b>	<b>✓</b>

# Allo Accelerator Design Language (ADL) and Compiler



## Stepwise verifiable rewrites

```
s = allo.customize(gemm)
s.reorder("k", "j")
print(s.module)
```

```
s.buffer_at(s.C, axis="i")
print(s.module)
```

```
s.pipeline("j")
print(s.module)
```

```
module {
func.func @gemm(%arg0: memref<1024x1024xf32>, %arg1: memref<1024x1024xf32>) ->
memref<1024x1024xf32> {
  %alloc = memref.alloc() {name = "C"} : memref<1024x1024xf32>
  affine.for %arg2 = 0 to 1024 {
    affine.for %arg3 = 0 to 1024 {
      affine.for %arg4 = 0 to 1024 {
        ...
      }
    }
  }
  {loop_name = "j", pipeline}
  {loop_name = "k", op_name = "S_k_0", reduction}
  ...
  {loop_name = "i", op_name = "S_i_j_0"}
  return %alloc : memref<1024x1024xf32>
}}
```



```
module {
func.func @gemm(%arg0: memref<1024x1024xf32>, %arg1: memref<1024x1024xf32>) ->
memref<1024x1024xf32> {
  %alloc = memref.alloc() {name = "C"} : memref<1024x1024xf32>
  affine.for %arg2 = 0 to 1024 {
    affine.for %arg3 = 0 to 1024 {
      affine.for %arg4 = 0 to 1024 {
        ...
      }
    }
  }
  {loop_name = "k"}
  {loop_name = "j"}
  {loop_name = "i", op_name = "S_i_j_0"}
  return %alloc : memref<1024x1024xf32>
}}
```

```
module {
func.func @gemm(%arg0: memref<1024x1024xf32>, %arg1: memref<1024x1024xf32>) ->
memref<1024x1024xf32> {
  %alloc = memref.alloc() {name = "C"} : memref<1024x1024xf32>
  %alloc_0 = memref.alloc() : memref<1024xf32>
  affine.for %arg3 = 0 to 1024 {
    ...
  }
  {buffer, loop_name = "j_init", pipeline_ii = 1 : i32}
  affine.for %arg3 = 0 to 1024 {
    affine.for %arg4 = 0 to 1024 {
      ...
    }
  }
  {loop_name = "j"}
  {loop_name = "k", op_name = "S_k_0", reduction}
  affine.for %arg3 = 0 to 1024 {
    ...
  }
  {buffer, loop_name = "j_back", pipeline_ii = 1 : i32}
  {loop_name = "i", op_name = "S_i_j_0"}
  return %alloc : memref<1024x1024xf32>
}}
```

Framework	Input	Paradigm	Decoupled Cust.	Verified Rewrite
Spatial [PLDI'18]	Scala	Imperative	✗	✗
TVM TE [OSDI'18]	Python	Declarative	✓	✗
<b>Allo</b>	<b>Python</b>	<b>Imperative</b>	<b>✓</b>	<b>✓</b>

# Transforming GEMM to Systolic Array

## → Algorithm specification


```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N]):  
    for i, j in allo.grid(M, N, "PE"):  
        for k in range(K):  
            C[i, j] += A[i, k] * B[k, j]
```



## → Schedule construction

```
s = allo.customize(matmul)
```

## → Architectural Diagram



\* Schedule: A sequence of customization primitives

# Transforming GEMM to Systolic Array

## → Algorithm specification

```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N]):  
    for i, j in allo.grid(M, N, "PE"):  
        for k in range(K):  
            C[i, j] += A[i, k] * B[k, j]
```

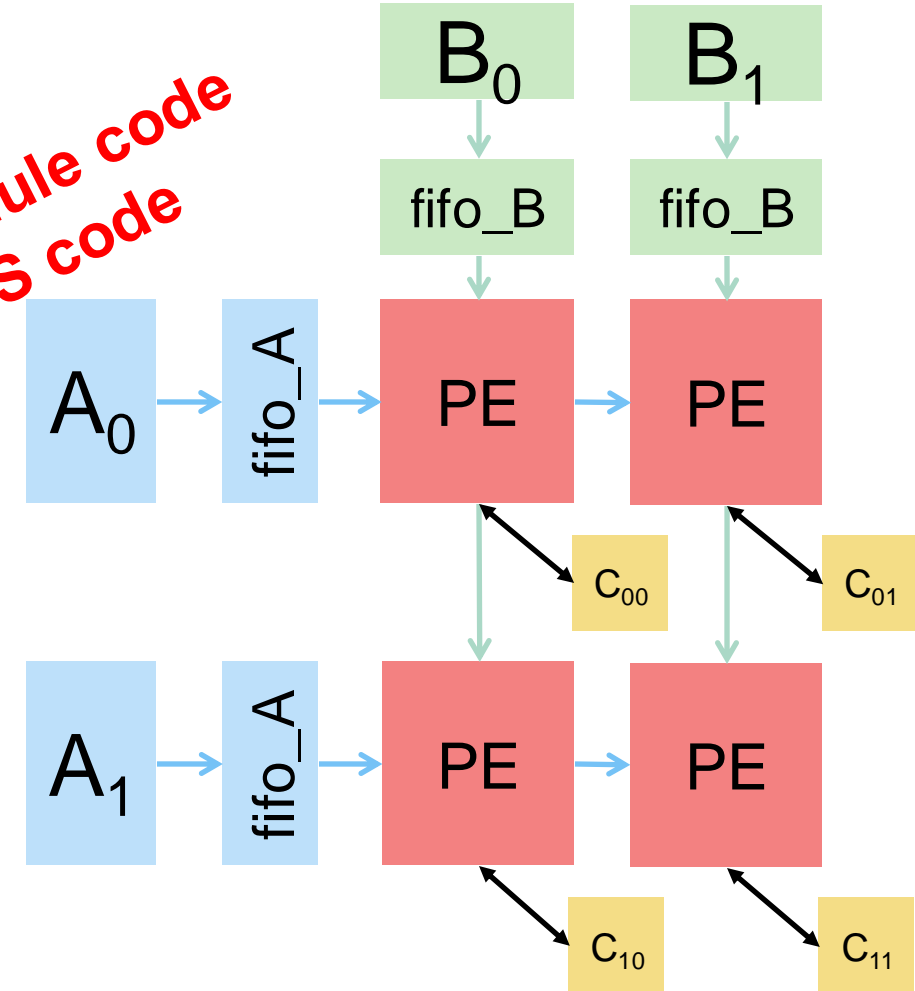


Only 8 lines of schedule code  
vs ~500 lines HLS code

## → Schedule construction

```
s = allo.customize(matmul)  
buf_A = s.buffer_at(s.A, "j")  
buf_B = s.buffer_at(s.B, "j")  
pe = s.unfold("PE", axis=[0, 1],  
              factor=[M, N])  
s.partition(s.A, dim=0)  
s.partition(s.B, dim=1)  
s.partition(s.C, dim=[0, 1])  
s.relay(buf_A, pe, axis=0, depth=M + 1)  
s.relay(buf_B, pe, axis=1, depth=N + 1)
```

## → Architectural Diagram



\* Schedule: A sequence of customization primitives

# Transforming GEMM to Systolic Array

## → Algorithm specification

```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N]):  
    for i, j in allo.grid(M, N, "PE"):  
        for k in range(K):  
            C[i, j] += A[i, k] * B[k, j]
```




## → Schedule construction

```
s = allo.customize(matmul)  
buf_A = s.buffer_at(s.A, "j")  
buf_B = s.buffer_at(s.B, "j")  
pe = s.unfold("PE", axis=[0, 1],  
             factor=[M, N])  
s.partition(s.A, dim=0)  
s.partition(s.B, dim=1)  
s.partition(s.C, dim=[0, 1])  
s.relay(buf_A, pe, axis=0, depth=M + 1)  
s.relay(buf_B, pe, axis=1, depth=N + 1)
```

## → CPU Simulation

```
mod = s.build(target="llvm")  
  
A = np.random.randint(-8, 8, size=(M, K)) \  
    .astype(np.int8)  
B = np.random.randint(-8, 8, size=(K, N)) \  
    .astype(np.int8)  
C = np.zeros((M, N), dtype=np.int16)  
mod(A, B, C)  
np.testing.assert_allclose(C, A @ B, atol=1e-3)
```



# Transforming GEMM to Systolic Array

## → Algorithm specification

```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N]):  
    for i, j in allo.grid(M, N, "PE"):  
        for k in range(K):  
            C[i, j] += A[i, k] * B[k, j]
```





## → Schedule construction

```
s = allo.customize(matmul)  
buf_A = s.buffer_at(s.A, "j")  
buf_B = s.buffer_at(s.B, "j")  
pe = s.unfold("PE", axis=[0, 1],  
             factor=[M, N])  
s.partition(s.A, dim=0)  
s.partition(s.B, dim=1)  
s.partition(s.C, dim=[0, 1])  
s.relay(buf_A, pe, axis=0, depth=M + 1)  
s.relay(buf_B, pe, axis=1, depth=N + 1)
```

## → Formal Verification [FPGA'24 Best Paper\*]

```
mod = s.verify()
```



- A **formal equivalence checker** of source-to-source HLS transformations via symbolic execution
  - Support statically interpretable control-flow (SICF)
- Verification in time/space linear w.r.t. OPs executed
  - ~500K Ops/sec in verification throughput
  - A complex **64x64 systolic array verified in 16 minutes**

# Transforming GEMM to Systolic Array


## → Algorithm specification

```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N]):  
    for i, j in allo.grid(M, N, "PE"):  
        for k in range(K):  
            C[i, j] += A[i, k] * B[k, j]
```



## → Bitstream generation

```
mod = s.build(target="vitis_hls",  
             mode="hw",  
             project="systolic.prj")  
  
# Automatically invoke the HLS toolchain  
mod(A, B, C)
```



First FFN layer in BERT-base (512, 768)x(768, 3072) w/ 16x16 SA

	Latency (ms)	BRAM	DSP	FF	LUT
Allo	<b>15.73</b>	0 (0%)	<b>128 (1%)</b>	79969 (3%)	244439 (18%)
AutoSA	15.71	514 (12%)	256 (2%)	100138 (3%)	244032 (18%)


Same level of performance but much lower resource usage



# Transforming GEMM to Systolic Array

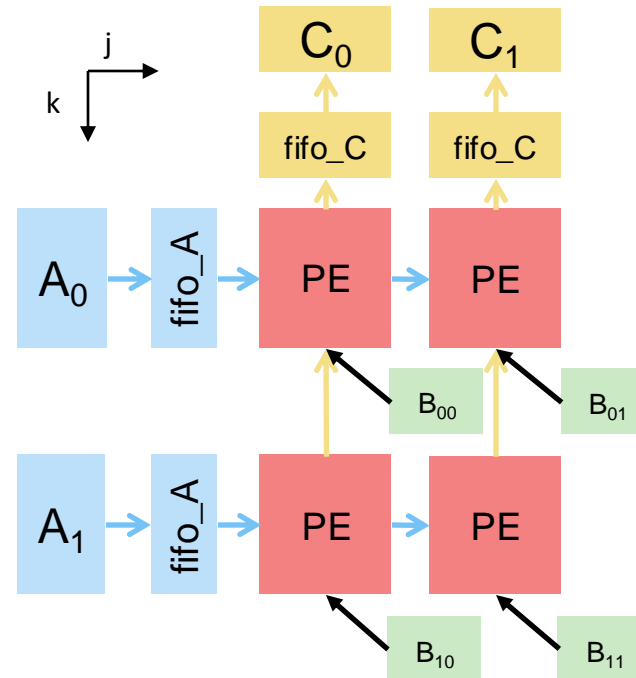
## → Spatial loop: i, j

```
s = allo.customize(matmul)
buf_A = s.buffer_at(s.A, "j")
buf_B = s.buffer_at(s.B, "j")
pe = s.unfold("PE", axis=[0, 1],
             factor=[M, N])
s.partition(s.A, dim=0)
s.partition(s.B, dim=1)
s.partition(s.C, dim=[0, 1])
s.relay(buf_A, pe, axis=0, depth=M + 1)
s.relay(buf_B, pe, axis=1, depth=N + 1)
```




## → Spatial loop: k, j

```
s = allo.customize(matmul)
s.reorder("k", "j", "i")
buf_A = s.buffer_at(s.A, "j")
buf_C = s.buffer_at(s.C, "j")
pe = s.unfold("PE", axis=[0, 1],
             factor=[K, N])
s.partition(s.A, dim=0)
s.partition(s.B, dim=[0, 1])
s.partition(s.C, dim=1)
s.relay(buf_A, pe, axis=0)
s.relay(buf_C, pe, axis=1)
```




## → Spatial loop: j

```
s = allo.customize(matmul)
buf_A = s.buffer_at(s.A, "j")
buf_B = s.buffer_at(s.B, "j")
pe = s.unfold("PE", axis=1,
             factor=M)
s.partition(s.B, dim=1)
s.partition(s.C, dim=1)
s.relay(buf_A, pe, axis=0)
s.relay(buf_B, pe, axis=1)
```




**Realize different dataflows  
with minimal schedule code**

# How to Compose Optimized Kernels into Complete Accelerator?



**Optimized  
Kernels**



**High-performance  
Accelerator**

# Composable Schedules

## → Given different optimized kernel implementations

```
def matmul(A: int8[M, K],  
           B: int8[K, N],  
           C: int16[M, N])
```



## → Algorithm specification (Hierarchical)


```
def top(X: int8[M, K],  
        W_A: int8[K, N],  
        W_B: int8[N, K]):  
    Y: int8[M, N] = 0  
    Z: int8[M, K] = 0  
    matmul(X, W_A, Y)  
    matmul(Y, W_B, Z)  
    return Z
```

## → Schedule composition


```
# Previous customizations for matmul  
s_matmul = allo.customize(matmul)  
# ...
```

```
s_top = allo.customize(top)  
s_top.compose(s_matmul)
```


# Interconnection between Kernels




Temporal Composition



```
def top(
  X: int8[M, K],
  W_A: int8[K, N],
  W_B: int8[N, K]):
  Y: ?
  Z: int8[M, K] = 0
  matmul(X, W_A, Y)
  matmul(Y, W_B, Z)
  return Z
```



Spatial Composition



# Temporal Composition

## → Given optimized kernel implementations

```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N])
```



Mismatched  
interface

`s2.partition(s.A, dim=1)`



`s1.partition(s.C, dim=0)`




Matched  
interface!



`s2.partition(s.A, dim=0)`

## → Algorithm specification (Hierarchical)

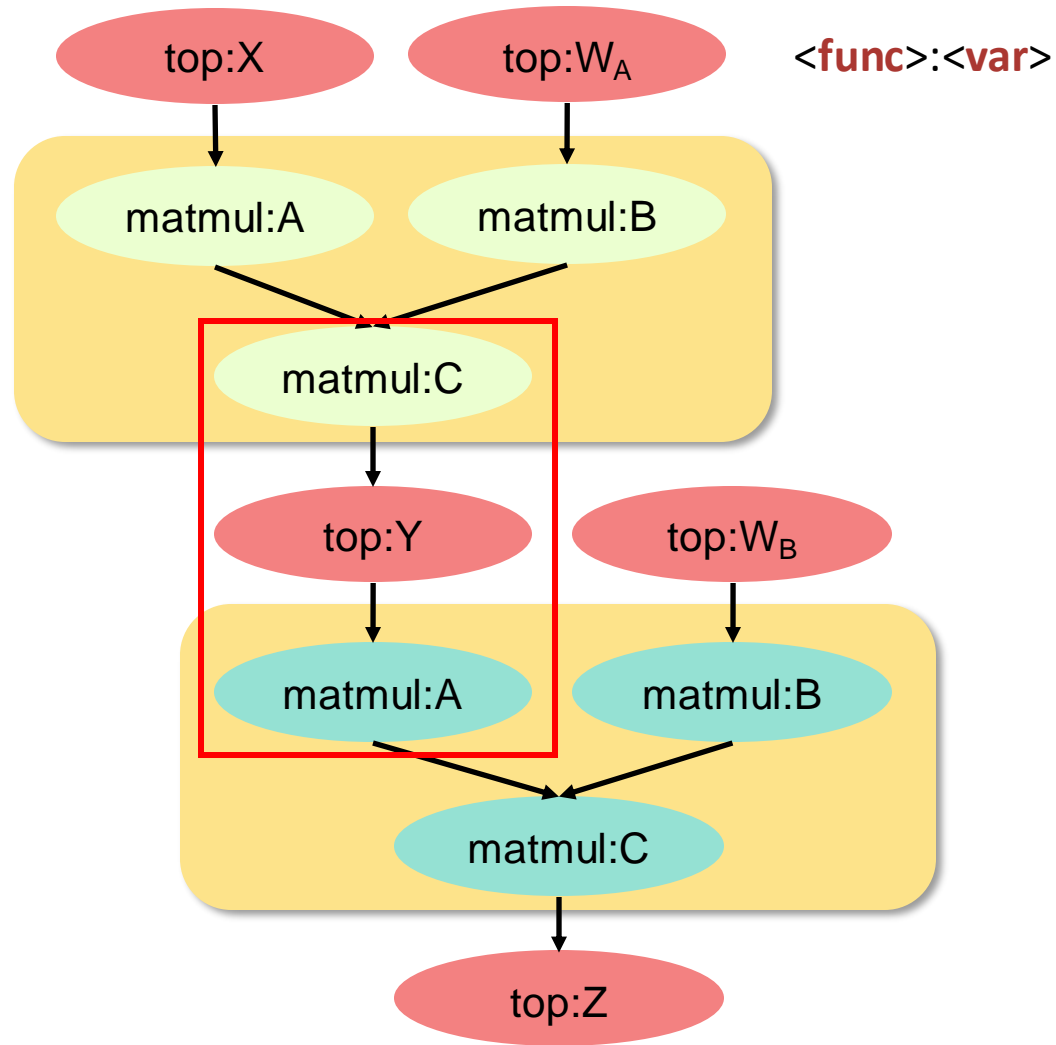
```
def top(X: int8[M, K], W_A: int8[K, N],  
       W_B: int8[N, K]):  
    Y: int8[M, N] = 0  
    Z: int8[M, K] = 0  
    matmul(X, W_A, Y)  
    matmul(Y, W_B, Z)  
    return Z
```



Need to ensure  
interface consistency



# Temporal Composition



Hierarchical use-def graph

## → Algorithm specification (Hierarchical)

```
def top(X: int8[M, K], W_A: int8[K, N],  
        W_B: int8[N, K], Y: int8[M, K]):  
    Y: int8[M, N] = 0  
    Z: int8[M, K] = 0  
    matmul(X, W_A, Y) ← Caller definition  
    matmul(Y, W_B, Z)  
    return Z
```

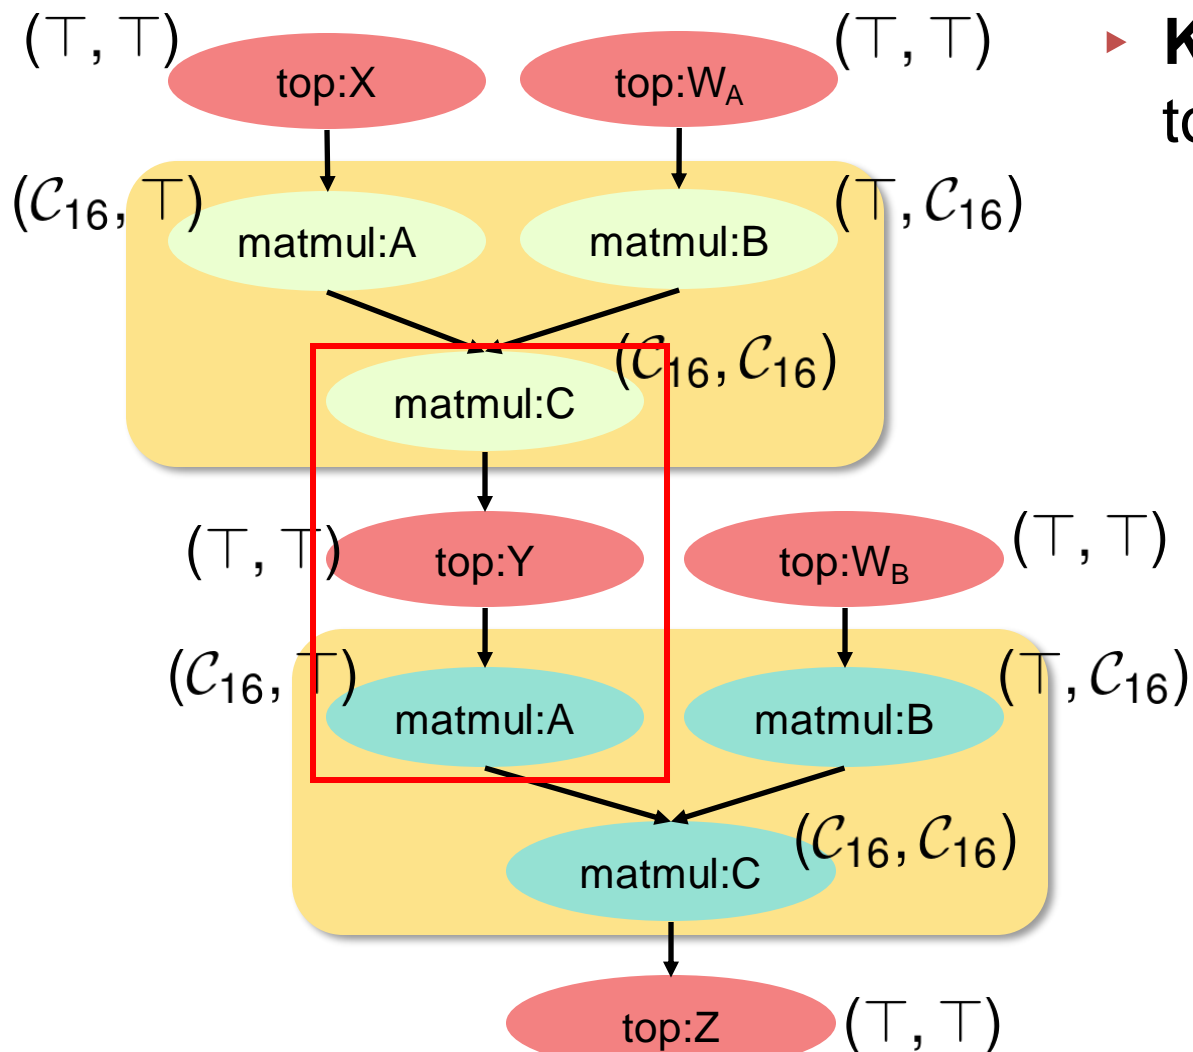
← Caller definition

## → Callee definition

```
def matmul(A: int8[M, K],  
           B: int8[K, N],  
           C: int16[M, N])
```

**Need to ensure interface consistency**


# Temporal Composition



Hierarchical use-def graph

- **Key idea: Model memory banking as a type** to ensure kernel interfaces are consistent


Subtyping relation forms a **lattice**!



**Intuition:**

We can supply more read/write parallelism, but not less!

# Temporal Composition





Hierarchical use-def graph

- **Key idea: Model data layout as a type** to ensure the kernel interfaces are consistent
- Lattice-based subtyping relation permits **linear-time layout inference** through Worklist algorithm



# Spatial Composition

- Model data streaming as a type



```
import allo.dataflow as df
```

```
@df.region
def top():
    pipe: Stream[Ty] =
        df.pipe(dtype=Ty, depth=4)
```

```
@df.kernel(mapping=[1])
def producer(A: Ty[M, N]):
    for i, j in allo.grid(M, N):
        # send data
        pipe.put(A[i, j])
```

```
@df.kernel(mapping=[1])
def consumer(B: Ty[M, N]):
    for i, j in allo.grid(M, N):
        # receive data
        B[i, j] = pipe.get() + 1
```

# Spatial Composition

```
@df.region()
```

```
def top():
```

```
    fifo_A = df.array(df.pipe(dtype=float32, depth=4), shape=(P0, P1))
```

```
    fifo_B = df.array(df.pipe(dtype=float32, depth=4), shape=(P0, P1))
```

```
@df.kernel(mapping=[P0, P1])
```

```
def gemm(A: float32[M, K],
```

```
        B: float32[K, N],
```

```
        C: float32[M, N]):
```

```
    i, j = df.get_pid()
```

```
    # peripheral kernels
```

```
    with allo.meta_if(i in {0, M + 1}
                      and j in {0, N + 1}):
```

```
        pass
```

```
    with allo.meta_elif(j == 0):
```

```
        # i > 0
```

```
        for k in range(K):
```


```
            fifo_A[i, j + 1].put(A[i - 1, k])
```

```
    with allo.meta_elif(i == 0):
```

```
        # j > 0
```

```
        for k in range(K):
```

```
            fifo_B[i + 1, j].put(B[k, j - 1])
```



# Spatial Composition

```
@df.region()
```

```
def top():
```

```
    fifo_A = df.array(df.pipe(dtype=float32, depth=4), shape=(P0, P1))
```

```
    fifo_B = df.array(df.pipe(dtype=float32, depth=4), shape=(P0, P1))
```

```
# drain
```

```
with allo.meta_elif(i == M + 1 and j > 0):
```

```
    for k in range(K):
```

```
        b: float32 = fifo_B[i, j].get()
```

```
with allo.meta_elif(j == N + 1 and i > 0):
```

```
    for k in range(K):
```

```
        a: float32 = fifo_A[i, j].get()
```

```
# main body
```

```
with allo.meta_else():
```

```
    c: float32 = 0
```

```
    for k in range(K):
```

```
        a: float32 = fifo_A[i, j].get()
```


```
        b: float32 = fifo_B[i, j].get()
```

```
        c += a * b
```

```
        fifo_A[i, j + 1].put(a)
```


```
        fifo_B[i + 1, j].put(b)
```

```
        C[i - 1, j - 1] = c
```



# Single-Kernel Evaluation

- ▶ Benchmarks from PolyBench; Target hardware: AMD U280 FPGA
- ▶ Normalized against AMD VitisHLS-auto (pragmas automatically inserted)
- ▶ Other baselines
  - ADLs: HeteroCL [FPGA'19], Dahlia [PLDI'20], PyLog [TC'21]
  - Automated DSE for HLS: ScaleHLS [HPCA'22], Merlin [TRETS'22]



Allo achieves (much) higher performance by optimizing data placement with a custom memory hierarchy

# Single-Kernel Evaluation

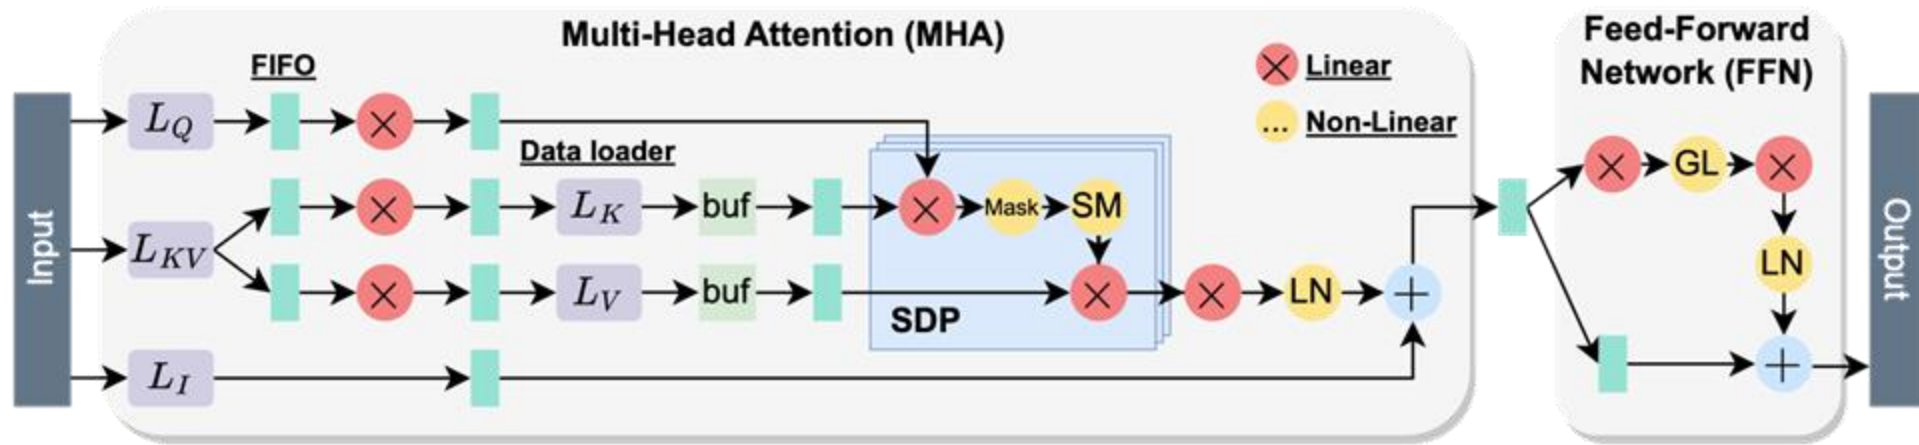
- ▶ Compared to ScaleHLS [HPCA'22], Allo achieves
  - Lower latency with much more effective use of compute resources
  - Higher post place-and-route frequency due to better pipelining

Benchmark	Allo						ScaleHLS				
	Latency (cycles)	II	DSP Usage	PnR Freq. (MHz)	Lines of Allo	Custm. Compile Time (s)	Latency (cycles)	II	DSP Usage	PnR Freq. (MHz)	Compile Time (s)
atax	4.9K (↓ 3.9×)	1	403 (↑ 2.9×)	411	9	1.0	19.4K	4	141	329	36.1
correlation	498.7K (↓ 290.5×)	1	4168 (↑ 38.2×)	362	19	0.8	144.9M	667	109	305	638.8
jacobi-2d	58.8K (↓ 183.1×)	1	3968 (↑ 72.1×)	411	17	0.9	10.8M	28	55	308	47.9
symm	405.7K (↓ 427.4×)	1	1208 (↑ 201.3×)	402	15	1.0	182.4M	13	6	397	3.5
trmm	492.6K (↓ 78.0×)	1	101 (↑ 14.4×)	414	12	0.8	38.4M	4	7	382	1.4

Allo achieves (much) higher performance by optimizing data placement with a custom memory hierarchy

# Multi-Kernel Evaluation: A Complete LLM Accelerator

- ▶ GPT2 model (the only open-source LLM in the GPT family)
  - 355M parameters, 24 hidden layers, 16 heads
  - W4A8 quantization




```
def GPT_layer(inp: float32[L, D], ...)
    -> float32[L, D]:
    # 1. Multi-Head Attention (MHA)
    Q = Linear_layer_qkv(inp, Wq, Bq)
    K = Linear_layer_qkv(inp, Wk, Bk)
    V = Linear_layer_qkv(inp, Wv, Bv)
    attn_sf_outp = Self_attention(Q, K, V)
    # ...
    # 2. Feed Forward Network (FFN)
    # ...
    return ff_n_res_outp
```

Compose all the schedules together

```
s = allo.customize(GPT_layer)
s.compose([s_qkv, ..., s_gelu])
```

# LLM Accelerator Evaluation


- ▶ GPT2: single-batch, low-latency, generative inference settings
  - AMD U280 FPGA (16nm), 250MHz
  - **2.2x speedup in prefill stage** compared to DFX [MICRO'22] (an FPGA-based overlay)
  - **1.9x speedup for long output** sequences and **5.7x more energy-efficient** vs. NVIDIA A100 GPU (7nm)
  - Fewer than 50 lines of schedule code in Allo



	Allo	DFX
Device	U280	U280
Freq.	250MHz	200MHz
Quant.	W4A8	fp16
BRAM	384 (19.0%)	1192 (59.1%)
DSP	1780 (19.73%)	3533 (39.2%)
FF	652K (25.0%)	1107K (42.5%)
LUT	508K (39.0%)	520K (39.9%)

# AMD AI Engine Evaluation

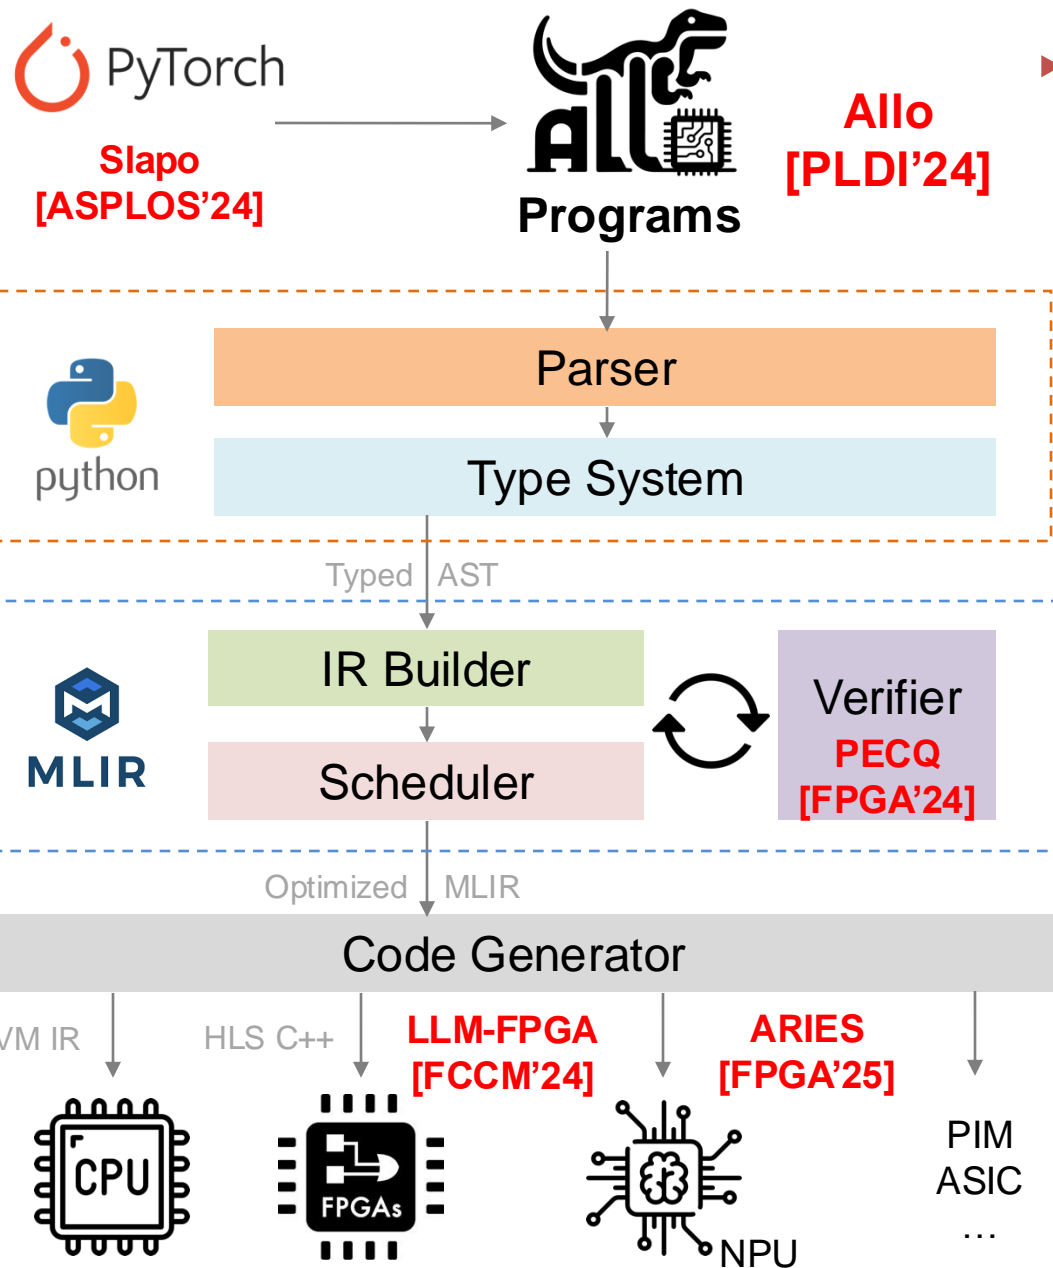
- ▶ CNN on NPU
  - ResNet on AMD Ryzen 7940 SoC
  - 23x speedup compared to Riallto



Designs	Tile	SL	IP	Util (%)	RT (ms)	Speedup
D1 scalar	[16, 16]	1	No	15	57.40	1.24
D2 +vectorized	[16, 16]	4	No	15	16.63	4.29
D3 +vectorized	[16, 16]	8	No	15	9.82	7.27
D4 +conv3×3-2core	[16, 16]	8	No	20	9.18	7.77
D5 +inst-pipeline	[16, 16]	8	Yes	20	8.44	8.45
D6 +opt-tile-size	[32, 16]	8	Yes	20	5.72	13.70
D7 +opt-tile-size	[32, 32]	8	Yes	20	5.21	12.48
D8 +opt-tile-size	[32, 64]	8	Yes	20	4.99	14.30
D9 +more-cores	[32, 64]	8	Yes	40	3.16	22.58
Riallto / RAI-SW[20]	-	-	-	20	71.36	1x



# Summary



## ► Features of Allo ADL

- Pythonic
- Decoupled & verifiable customizations
- Composability



<https://github.com/cornell-zhang/allo>



Cornell University



UNIVERSITY OF ILLINOIS  
URBANA-CHAMPAIGN



BROWN



UNIVERSITY OF TORONTO



UNIVERSITY of VIRGINIA



Microsoft



UNIVERSITY OF ILLINOIS CHICAGO



THE UNIVERSITY OF CHICAGO