


ALAC: A Programming Model for Composable Accelerator Design

Hongzheng Chen*, Niansong Zhang*, Shaojie Xiang,
Zhichen Zeng[†], Mengjia Dai[†], Zhiru Zhang

Cornell University

[†]University of Science and Technology of China

PLDI'24



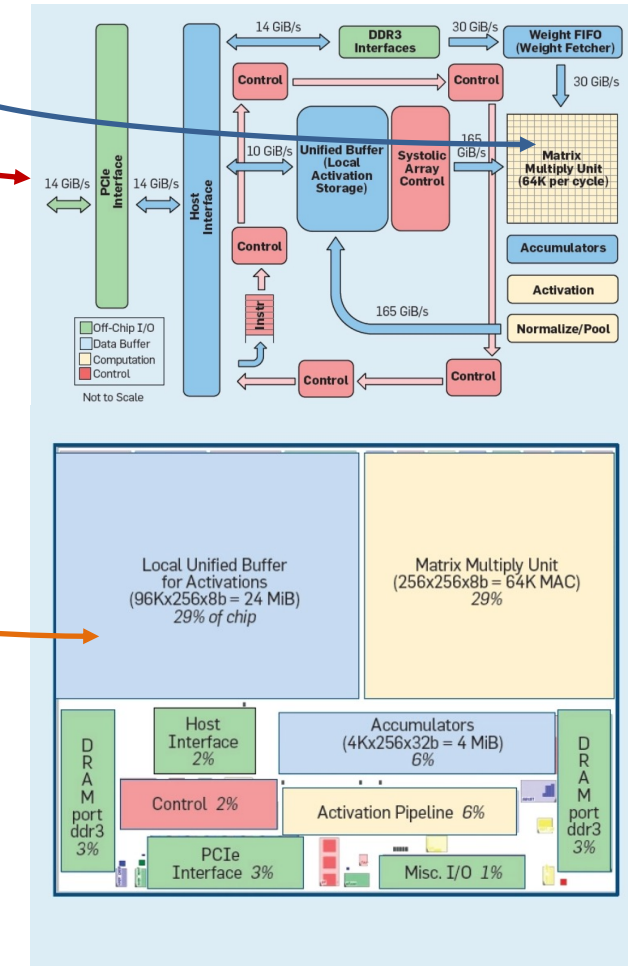

Cornell University

*Equal Contribution



Complexity in Specialized Accelerator Design

- ▶ Accelerator design is different from programming on general processors
 - Custom processing engines (PEs)
 - Custom data communication
 - Custom memory hierarchy



An AI accelerator example

Challenge: Balancing Manual Control & Compiler Optimization

```
void systolic_tile(int8_t A_tile[2][768],
int8_t B_tile[768][2],
int8_t C_tile[2][2]) {
#pragma dataflow
hls::stream<int8_t> A_fifo[2][3], B_fifo[2][3];
#pragma stream variable=A/B_fifo depth=3
#pragma partition variable=A/B/C_tile complete dim=1
for (int k4 = 0; k4 < 768; k4++) {
for (int m = 0; m < 2; m++) {
int8_t v105 = A_tile[m][k4];
A_fifo[m][0].write(v105);
// ... write B_fifo
}
for (int Ti = 0; Ti < 2; ++Ti) {
#pragma HLS unroll
for (int Tj = 0; Tj < 2; ++Tj) {
#pragma HLS unroll
// ... load A/B_fifo
PE_kernel(A_in, A_out, B_in, B_out, C, Ti, Tj);
}
}
}
}
```

Compute
Customization

```
void matmul(int8_t A[512][768], int8_t B[768][768],
int8_t C[512][768]) {
int8_t local_A[2][768], local_B[768][2], local_C[2][2];
for (int mi = 0; mi < 256; mi++) {
for (int ni = 0; ni < 384; ni++) {
// ... load A, B
systolic_tile(local_A, local_B, local_C);
}
}
}
```

Vanilla Matmul (1% theoretical peak perf.)

+ Compute customization (~30% peak)



+ Loop tiling

+ Loop unrolling

+ Loop & function pipelining

Increasing programming effort

Memory



Accelerator

Challenge: Balancing Manual Control & Compiler Optimization

```
void systolic_tile(int8_t A_tile[2][768],
int8_t B_tile[768][2],
int8_t C_tile[2][2]) {
#pragma dataflow
hls::stream<int8_t> A_fifo[2][3], B_fifo[2][3];
#pragma stream variable=A/B_fifo depth=3
#pragma partition variable=A/B/C_tile complete dim=1
for (int k4 = 0; k4 < 768; k4++) {
for (int m = 0; m < 2; m++) {
int8_t v105 = A_tile[m][k4];
A_fifo[m][0].write(v105);
// ... write B_fifo
}
for (int Ti = 0; Ti < 2; ++Ti) {
#pragma HLS unroll
for (int Tj = 0; Tj < 2; ++Tj) {
#pragma HLS unroll
// ... load A/B_fifo
PE_kernel(A_in, A_out, B_in, B_out, C, Ti, Tj);
}
}
}
}
```

**Memory
Customization**

**Compute
Customization**

```
void matmul(int8_t A[512][768], int8_t B[768][768],
int8_t C[512][768]) {
int8_t local_A[2][768], local_B[768][2], local_C[2][2];
for (int mi = 0; mi < 256; mi++) {
for (int ni = 0; ni < 384; ni++) {
// ... load A, B
systolic_tile(local_A, local_B, local_C);
}
}
}
```

Vanilla Matmul (1% theoretical peak perf.)

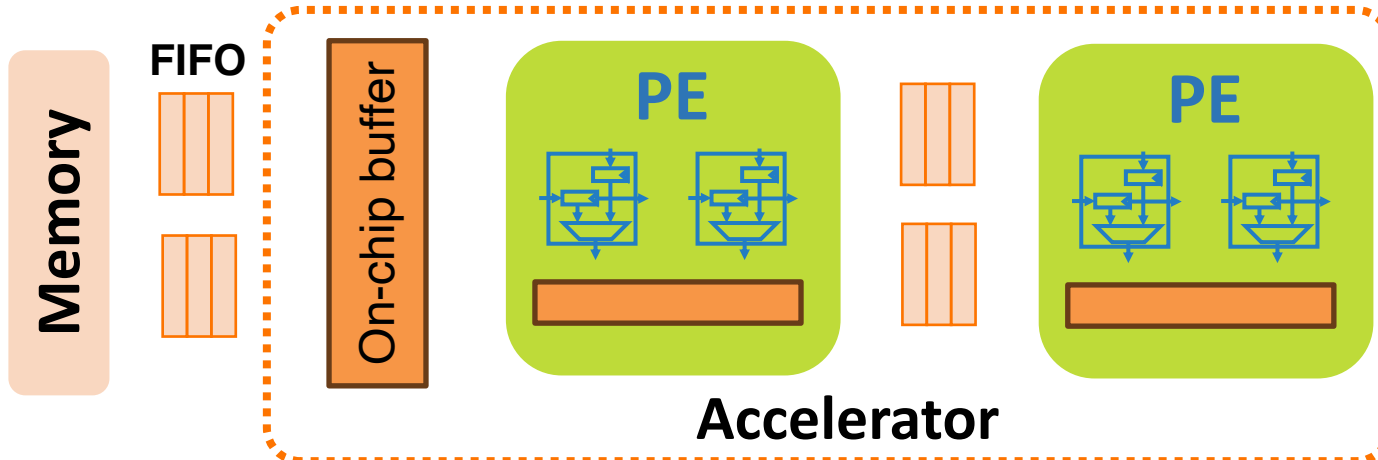
+ Compute customization (~30% peak)

- + Loop tiling
- + Loop unrolling
- + Loop & function pipelining

+ Custom memory hierarchy (~50% peak)

- + Tiling & data reuse buffers
- + Memory banking/partitioning

Increasing programming effort



Challenge: Balancing Manual Control & Compiler Optimization

```
void systolic_tile(int8_t A_tile[2][768],
int8_t B_tile[768][2],
int8_t C_tile[2][2]) {
#pragma dataflow
hls::stream<int8_t> A_fifo[2][3], B_fifo[2][3];
#pragma stream variable=A/B_fifo depth=3
#pragma partition variable=A/B/C_tile complete dim=1
for (int k4 = 0; k4 < 768; k4++) {
for (int m = 0; m < 2; m++) {
int8_t v105 = A_tile[m][k4];
A_fifo[m][0].write(v105);
// ... write B_fifo
}
for (int Ti = 0; Ti < 2; ++Ti) {
#pragma HLS unroll
for (int Tj = 0; Tj < 2; ++Tj) {
#pragma HLS unroll
// ... load A/B_fifo
PE_kernel(A_in, A_out, B_in, B_out, C, Ti, Tj);
}
}
}
}
```

**Communication
Customization**

**Memory
Customization**

**Compute
Customization**

```
void matmul(int8_t A[512][768], int8_t B[768][768],
int8_t C[512][768]) {
int8_t local_A[2][768], local_B[768][2], local_C[2][2];
for (int mi = 0; mi < 256; mi++) {
for (int ni = 0; ni < 384; ni++) {
// ... load A, B
systolic_tile(local_A, local_B, local_C);
}
}
}
```

Vanilla Matmul (1% theoretical peak perf.)

+ Compute customization (~30% peak)

- + Loop tiling
- + Loop unrolling
- + Loop & function pipelining


+ Custom memory hierarchy (~50% peak)

- + Tiling & data reuse buffers
- + Memory banking/partitioning

+ Data movement optimization (~95% peak)

- + Data streaming
- + Data packing (vectorization)
- + Memory coalescing
- + Systolic communication

Increasing programming effort



Challenge: Balancing Manual Control & Compiler Optimization


```
void systolic_tile(int8_t A_tile[2][768],
int8_t B_tile[768][2],
int8_t C_tile[2][2]) {
#pragma dataflow
hls::stream<int8_t> A_fifo[2][3], B_fifo[2][3];
#pragma stream variable=A/B_fifo depth=3
#pragma partition variable=A/B/C_tile complete dim=1
for (int k4 = 0; k4 < 768; k4++) {
for (int m = 0; m < 2; m++) {
int8_t v105 = A_tile[m][k4];
A_fifo[m][0].write(v105);
// ... write B_fifo
}
for (int Ti = 0; Ti < 2; ++Ti) {
#pragma HLS unroll
for (int Tj = 0; Tj < 2; ++Tj) {
#pragma HLS unroll
// ... load A/B_fifo
PE_kernel(A_in, A_out, B_in, B_out, C, Ti, Tj);
}
}
}
}
```

**Communication
Customization**

**Memory
Customization**

**Compute
Customization**

```
void matmul(int8_t A[512][768], int8_t B[768][768],
int8_t C[512][768]) {
int8_t local_A[2][768], local_B[768][2], local_C[2][2];
for (int mi = 0; mi < 256; mi++) {
for (int ni = 0; ni < 384; ni++) {
// ... load A, B
systolic_tile(local_A, local_B, local_C);
}
}
}
```



Vanilla Matmul (1% theoretical peak perf.)

+ Compute customization (~30% peak)

- + Loop tiling **Existing HLS compiler**
- + Loop unrolling **e.g., ScaleHLS [HPCA'22]**
- + Loop & function pipelining

+ Custom memory hierarchy (~50% peak)

- + Tiling & data reuse buffers
- + Memory banking/partitioning


+ Data movement optimization (~95% peak)

- + Data streaming
- + Data packing (vectorization)
- + Memory coalescing
- + Systolic communication

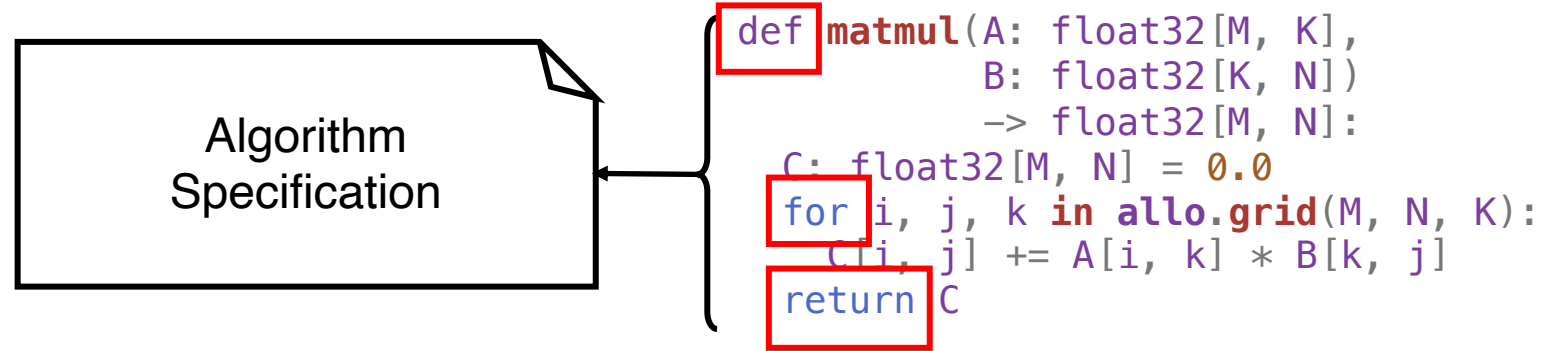
~500 lines of HLS code for a small systolic array
vendor-specific, hard to maintain & reuse

Increasing programming effort

Allo Accelerator Design Language (ADL) and Compiler




Pythonic: No need to learn a new DSL!



Framework	Input	Paradigm
Spatial [PLDI'18]	Scala	Imperative
TVM TE [OSDI'18]	Python	Declarative
Allo	Python	Imperative

Allo Accelerator Design Language (ADL) and Compiler



Pythonic: No need to learn a new DSL!

Algorithm
Specification

```
def matmul(A: float32[M, K],
           B: float32[K, N])
    -> float32[M, N]:
    C: float32[M, N] = 0.0
    for i, j, k in allo.grid(M, N, K):
        C[i, j] += A[i, k] * B[k, j]
    return C
```

Decoupled customization

Compute Cust.

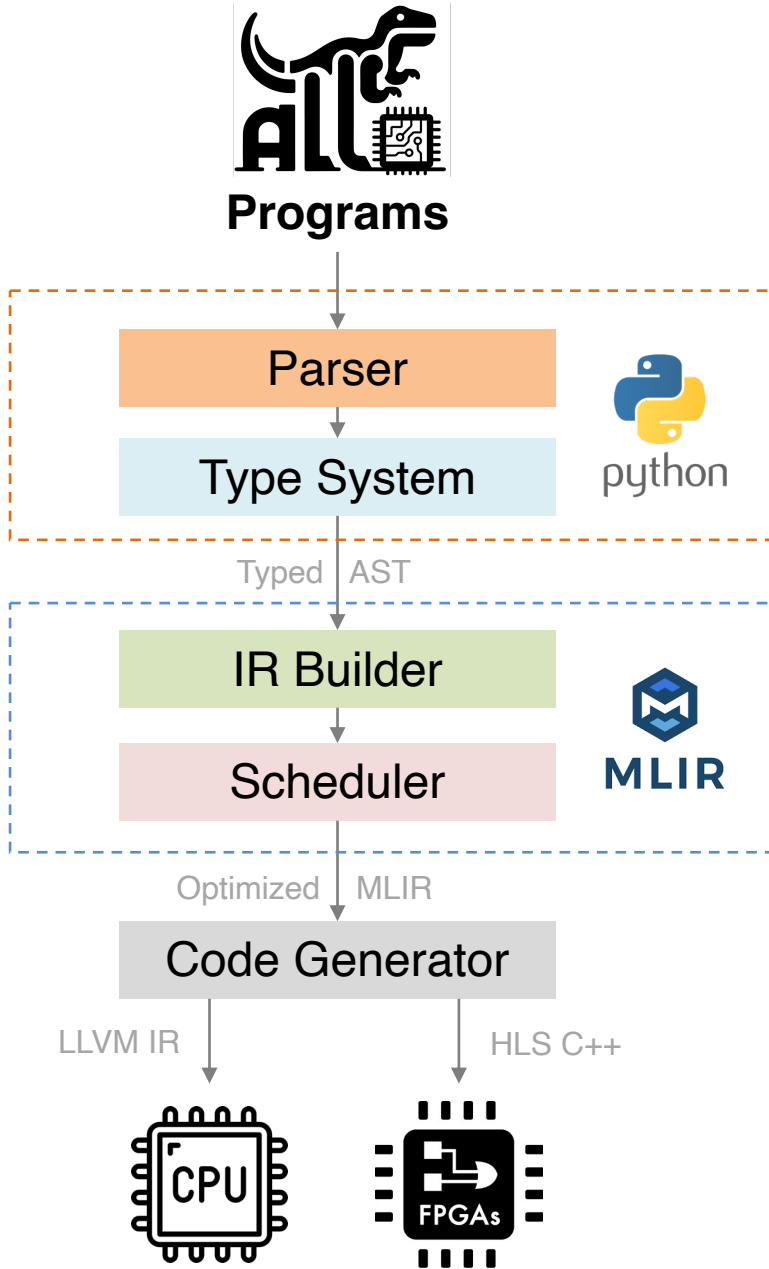
Memory Cust.

Comm. Cust.

```
s = allo.customize(matmul)
s.reorder("k", "j")
s.buffer_at(s.C, axis="i")
s.pipeline("j")
```

Framework	Input	Paradigm	Decoupled Cust.
Spatial [PLDI'18]	Scala	Imperative	✗
TVM TE [OSDI'18]	Python	Declarative	✓
Allo	Python	Imperative	✓

Allo Accelerator Design Language (ADL) and Compiler



Stepwise verifiable rewrites

```

s = allo.customize(gemm)
s.reorder("k", "j")
print(s.module)
  
```

```

s.buffer_at(s.C, axis="i")
print(s.module)
  
```

```

s.pipeline("j")
print(s.module)
  
```



```

module {
  func.func @gemm(%arg0: memref<1024x1024xf32>, %arg1:
    memref<1024x1024xf32>) -> memref<1024x1024xf32> {
    %alloc = memref.alloc() {name = "C"} : memref<1024x1024xf32>
    affine.for %arg2 = 0 to 1024 {
      affine.for %arg3 = 0 to 1024 {
        affine.for %arg4 = 0 to 1024 {
          ...
        } {loop_name = "k"}
      } {loop_name = "j"}
    } {loop_name = "i", op_name = "S_i_j_0"}
    return %alloc : memref<1024x1024xf32>
  }}
  
```

```

module {
  func.func @gemm(%arg0, %arg1) -> memref<1024x1024xf32> {
    %alloc = memref.alloc() {name = "C"} : memref<1024x1024xf32>
    affine.for %arg2 = 0 to 1024 {
      affine.for %arg3 = 0 to 1024 {
        affine.for %arg4 = 0 to 1024 {
          ...
        } {loop_name = "j", pipeline}
      } {loop_name = "k", op_name = "S_k_0", reduction}
    }
    ...
    } {loop_name = "i", op_name = "S_i_j_0"}
    return %alloc : memref<1024x1024xf32>
  }}
  
```

```

module {
  func.func @gemm(%arg0: memref<1024x1024xf32>, %arg1:
    memref<1024x1024xf32>) -> memref<1024x1024xf32> {
    %alloc = memref.alloc() {name = "C"} : memref<1024x1024xf32>
    affine.for %arg2 = 0 to 1024 {
      %alloc_0 = memref.alloc() : memref<1024xf32>
      affine.for %arg3 = 0 to 1024 {
        ...
      } {buffer, loop_name = "j_init", pipeline_ii = 1 : i32}
    } {loop_name = "j"}
    } {loop_name = "k", op_name = "S_k_0", reduction}
    affine.for %arg3 = 0 to 1024 {
      ...
    } {buffer, loop_name = "j_back", pipeline_ii = 1 : i32}
    } {loop_name = "i", op_name = "S_i_j_0"}
    return %alloc : memref<1024x1024xf32>
  }}
  
```

Framework	Input	Paradigm	Decoupled Cust.	Verified Rewrite
Spatial [PLDI'18]	Scala	Imperative	✗	✗
TVM TE [OSDI'18]	Python	Declarative	✓	✗
Allo	Python	Imperative	✓	✓

Transforming GEMM to Systolic Array

→ Algorithm specification


```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N]):  
    for i, j in allo.grid(M, N, "PE"):  
        for k in range(K):  
            C[i, j] += A[i, k] * B[k, j]
```



→ Schedule construction

```
s = allo.customize(matmul)
```

→ Architectural Diagram



* Schedule: A sequence of customization primitives

Transforming GEMM to Systolic Array

→ Algorithm specification

```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N]):  
    for i, j in allo.grid(M, N, "PE"):  
        for k in range(K):  
            C[i, j] += A[i, k] * B[k, j]
```

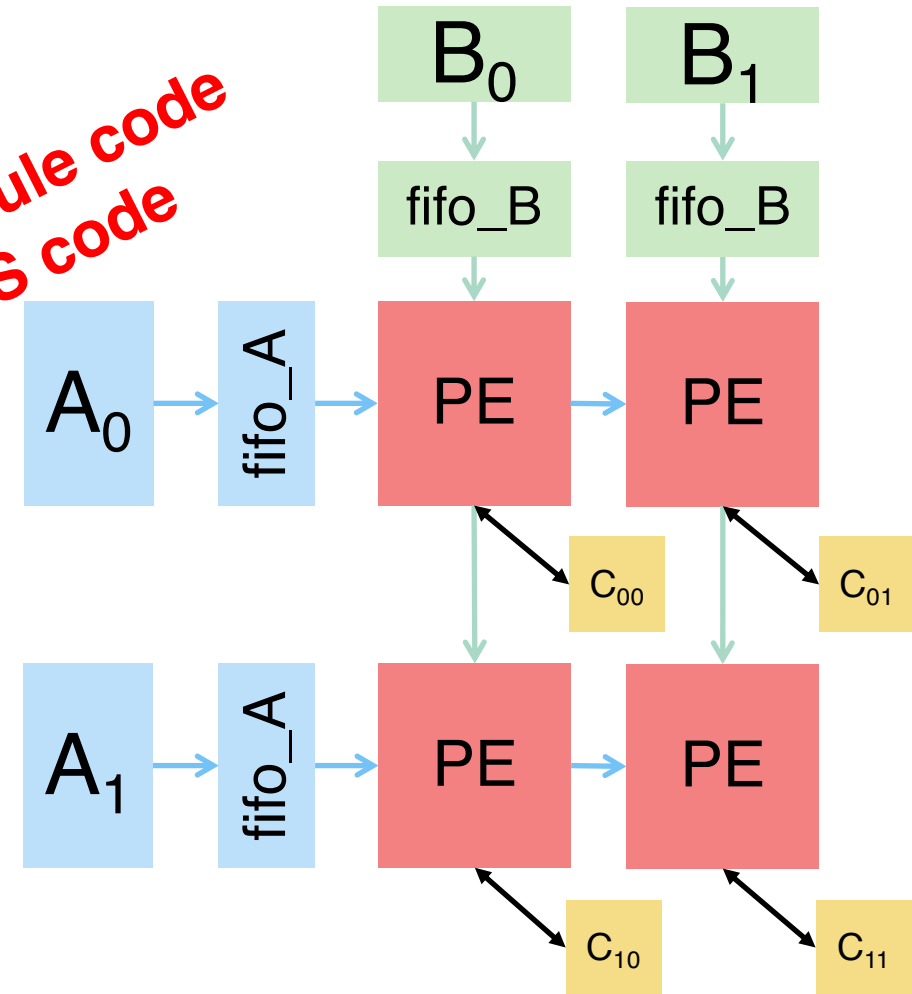


Only 8 lines of schedule code
vs ~500 lines HLS code

→ Schedule construction

```
s = allo.customize(matmul)  
buf_A = s.buffer_at(s.A, "j")  
buf_B = s.buffer_at(s.B, "j")  
pe = s.unfold("PE", axis=[0, 1],  
             factor=[M, N])  
s.partition(s.A, dim=0)  
s.partition(s.B, dim=1)  
s.partition(s.C, dim=[0, 1])  
s.relay(buf_A, pe, axis=0, depth=M + 1)  
s.relay(buf_B, pe, axis=1, depth=N + 1)
```

→ Architectural Diagram



Transforming GEMM to Systolic Array

→ Algorithm specification

```
def matmul(A: int8[M, K],
           B: int8[K, N],
           C: int16[M, N]):
    for i, j in allo.grid(M, N, "PE"):
        for k in range(K):
            C[i, j] += A[i, k] * B[k, j]
```



→ Schedule construction


```
s = allo.customize(matmul)
buf_A = s.buffer_at(s.A, "j")
buf_B = s.buffer_at(s.B, "j")
pe = s.unfold("PE", axis=[0, 1],
              factor=[M, N])

s.partition(s.A, dim=0)
s.partition(s.B, dim=1)
s.partition(s.C, dim=[0, 1])
s.relay(buf_A, pe, axis=0, depth=M + 1)
s.relay(buf_B, pe, axis=1, depth=N + 1)
```

→ CPU Simulation

```
mod = s.build(target="llvm")

A = np.random.randint(-8, 8, size=(M, K)) \
    .astype(np.int8)
B = np.random.randint(-8, 8, size=(K, N)) \
    .astype(np.int8)
C = np.zeros((M, N), dtype=np.int16)
mod(A, B, C)
np.testing.assert_allclose(C, A @ B, atol=1e-3)
```



Transforming GEMM to Systolic Array

→ Algorithm specification

```
def matmul(A: int8[M, K],
           B: int8[K, N],
           C: int16[M, N]):
    for i, j in allo.grid(M, N, "PE"):
        for k in range(K):
            C[i, j] += A[i, k] * B[k, j]
```





→ Schedule construction

```
s = allo.customize(matmul)
buf_A = s.buffer_at(s.A, "j")
buf_B = s.buffer_at(s.B, "j")
pe = s.unfold("PE", axis=[0, 1],
              factor=[M, N])
s.partition(s.A, dim=0)
s.partition(s.B, dim=1)
s.partition(s.C, dim=[0, 1])
s.relay(buf_A, pe, axis=0, depth=M + 1)
s.relay(buf_B, pe, axis=1, depth=N + 1)
```

→ Formal Verification [FPGA'24 Best Paper*]

```
mod = s.verify()
```



- A **formal equivalence checker** of source-to-source HLS transformations via symbolic execution
 - Support statically interpretable control-flow (SICF)
- Verification in time/space linear w.r.t. OPs executed
 - ~500K Ops/sec in verification throughput
 - A complex **64x64 systolic array verified in 16 minutes**

Transforming GEMM to Systolic Array


→ Algorithm specification

```
def matmul(A: int8[M, K],  
           B: int8[K, N],  
           C: int16[M, N]):  
    for i, j in allo.grid(M, N, "PE"):  
        for k in range(K):  
            C[i, j] += A[i, k] * B[k, j]
```



→ Bitstream generation

```
mod = s.build(target="vitis_hls",  
              mode="hw",  
              project="systolic.prj")  
  
# Automatically invoke the HLS toolchain  
mod(A, B, C)
```



First FFN layer in BERT-base (512, 768)x(768, 3072) w/ 16x16 SA


	Latency (ms)	BRAM	DSP	FF	LUT
Allo	15.73	0 (0%)	128 (1%)	79969 (3%)	244439 (18%)
AutoSA	15.71	514 (12%)	256 (2%)	100138 (3%)	244032 (18%)

Same level of performance but much lower resource usage

Transforming GEMM to Systolic Array

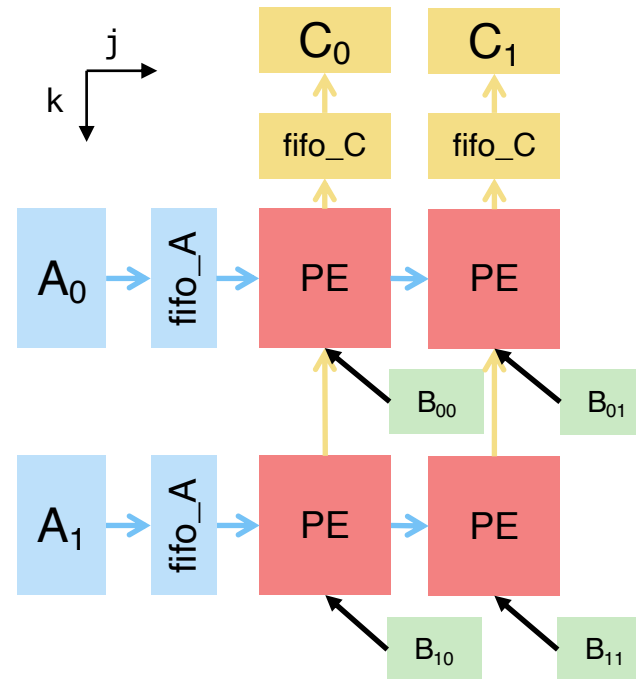
→ Spatial loop: i, j

```
s = allo.customize(matmul)
buf_A = s.buffer_at(s.A, "j")
buf_B = s.buffer_at(s.B, "j")
pe = s.unfold("PE", axis=[0, 1],
              factor=[M, N])
s.partition(s.A, dim=0)
s.partition(s.B, dim=1)
s.partition(s.C, dim=[0, 1])
s.relay(buf_A, pe, axis=0, depth=M + 1)
s.relay(buf_B, pe, axis=1, depth=N + 1)
```




→ Spatial loop: k, j

```
s = allo.customize(matmul)
s.reorder("k", "j", "i")
buf_A = s.buffer_at(s.A, "j")
buf_C = s.buffer_at(s.C, "j")
pe = s.unfold("PE", axis=[0, 1],
              factor=[K, N])
s.partition(s.A, dim=0)
s.partition(s.B, dim=[0, 1])
s.partition(s.C, dim=1)
s.relay(buf_A, pe, axis=0)
s.relay(buf_C, pe, axis=1)
```




→ Spatial loop: j

```
s = allo.customize(matmul)
buf_A = s.buffer_at(s.A, "j")
buf_B = s.buffer_at(s.B, "j")
pe = s.unfold("PE", axis=1,
              factor=M)
s.partition(s.B, dim=1)
s.partition(s.C, dim=1)
s.relay(buf_A, pe, axis=0)
s.relay(buf_B, pe, axis=1)
```




**Realize different dataflows
with minimal schedule code**

How to Compose Optimized Kernels into Complete Accelerator?



**Optimized
Kernels**



**High-performance
Accelerator**

Composable Schedules

→ Given different optimized kernel implementations

```
def matmul(A: int8[M, K],  
          B: int8[K, N],  
          C: int16[M, N])
```




→ Schedule composition

```
# Previous customizations for matmul  
s_matmul = allo.customize(matmul)  
# ...
```

```
s_top = allo.customize(top)  
s_top.compose(s_matmul)
```

→ Algorithm specification (Hierarchical)

```
def top(X: int8[M, K], W_A: int8[K, N],  
        W_B: int8[N, K], Y: int8[M, K]):  
    Y: int8[M, N] = 0  
    Z: int8[M, K] = 0  
    matmul(X, W_A, Y)  
    matmul(Y, W_B, Z)  
    return Z
```



Composable Schedules

→ Given optimized kernel implementations

```
def matmul(A: int8[M, K],  
           B: int8[K, N],  
           C: int16[M, N])
```



Mismatched
interface

`s2.partition(s.A, dim=1)`



`s1.partition(s.C, dim=0)`



Matched
interface!




`s2.partition(s.A, dim=0)`



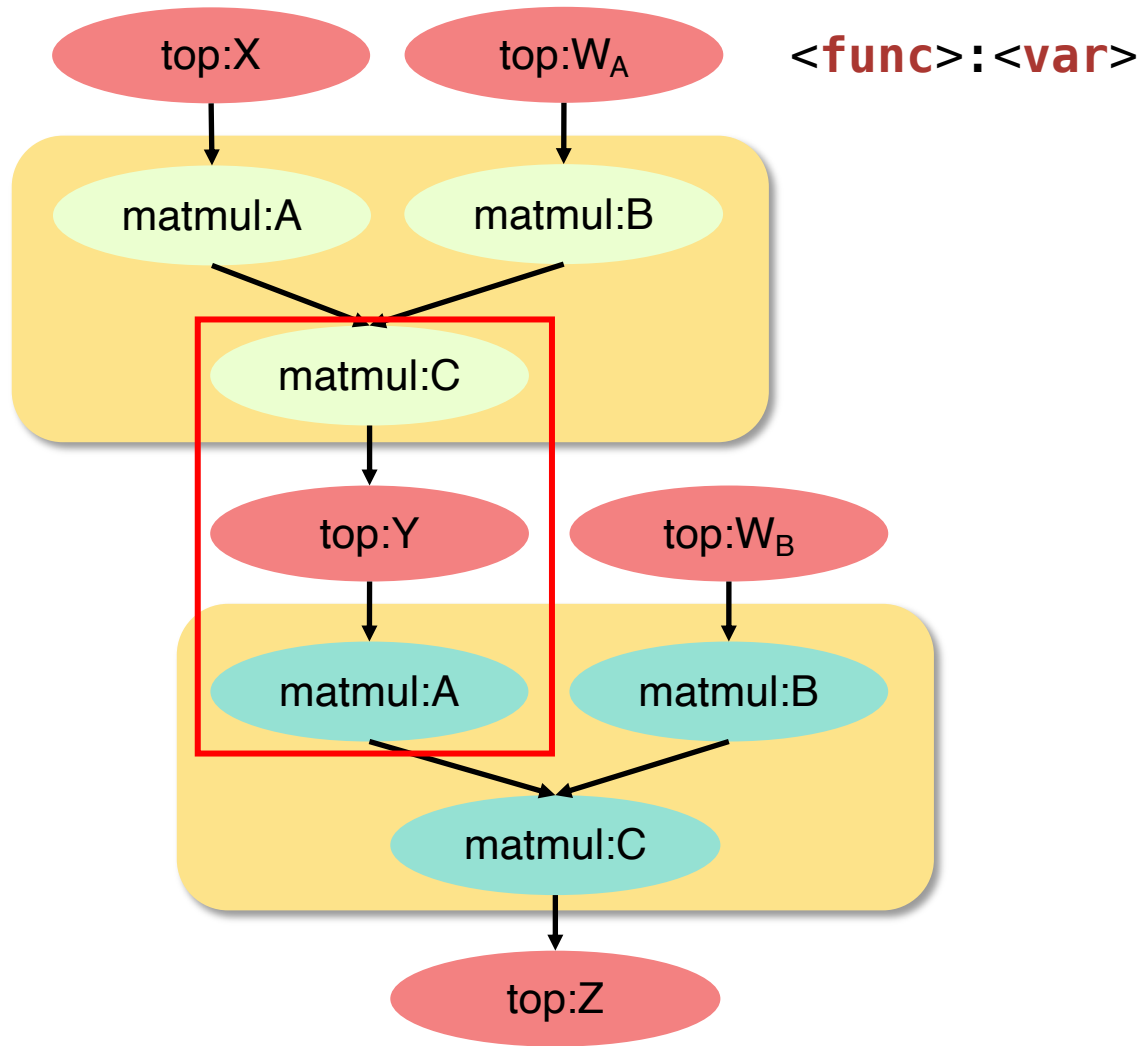
→ Algorithm specification (Hierarchical)

```
def top(X: int8[M, K], W_A: int8[K, N],  
        W_B: int8[N, K], Y: int8[M, K]):  
    Y: int8[M, N] = 0  
    Z: int8[M, K] = 0  
    matmul(X, W_A, Y)  
    matmul(Y, W_B, Z)  
    return Z
```



Need to ensure
interface consistency

Composable Schedules



Hierarchical dataflow graph

→ Algorithm specification (Hierarchical)

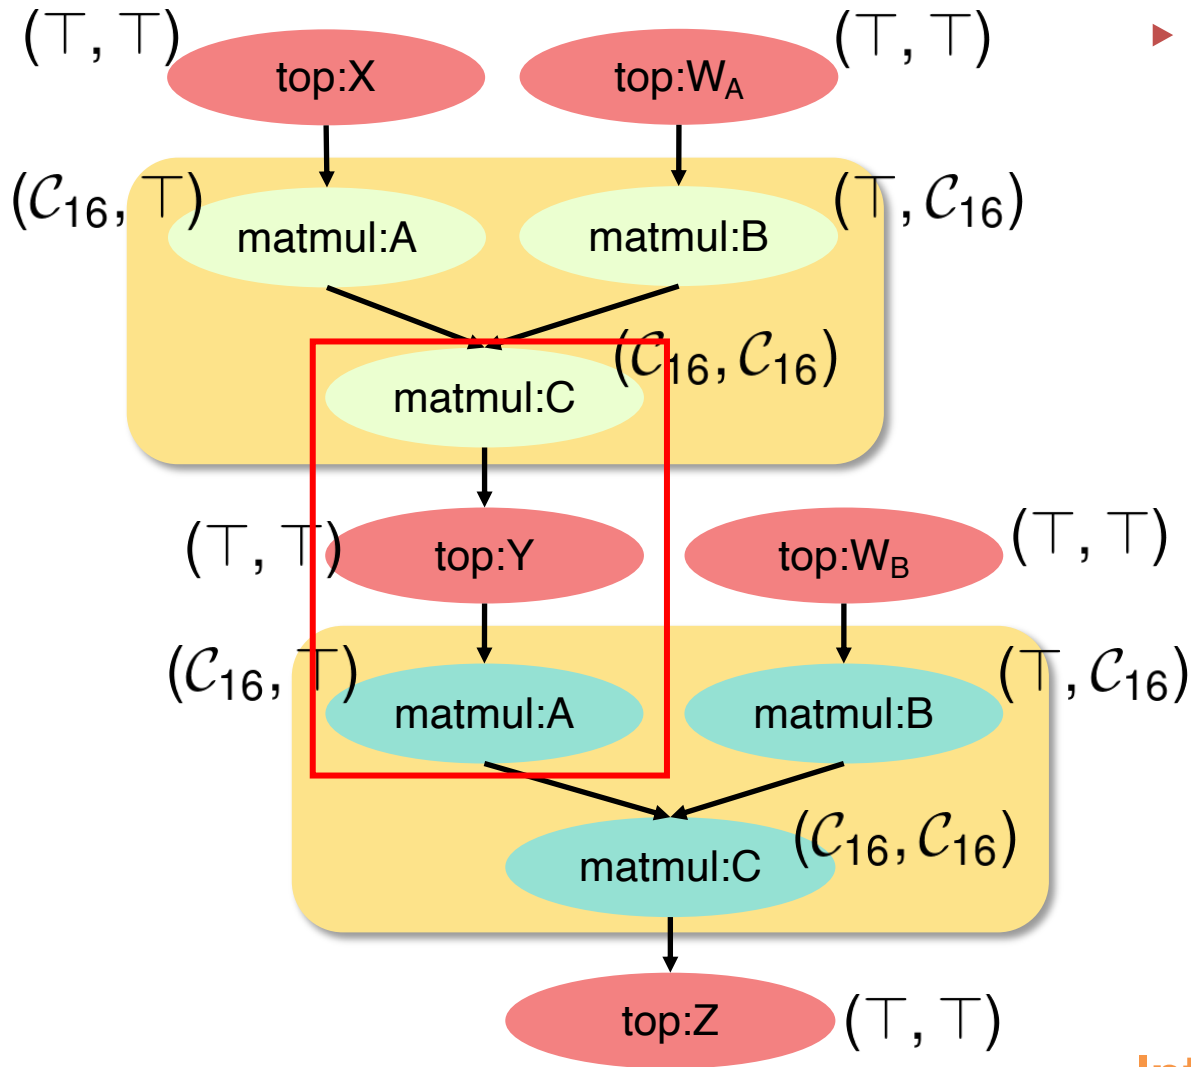
```
def top(X: int8[M, K], W_A: int8[K, N],  
        W_B: int8[N, K], Y: int8[M, K]):  
    Y: int8[M, N] = 0  
    Z: int8[M, K] = 0  
    matmul(X, W_A, Y) ← Caller definition  
    matmul(Y, W_B, Z)  
    return Z
```

→ Callee definition

```
def matmul(A: int8[M, K],  
           B: int8[K, N],  
           C: int16[M, N])
```

Need to ensure interface consistency


Composable Schedules



Hierarchical dataflow graph

- **Key idea: Model data layout as a type** to ensure the kernel interfaces are consistent

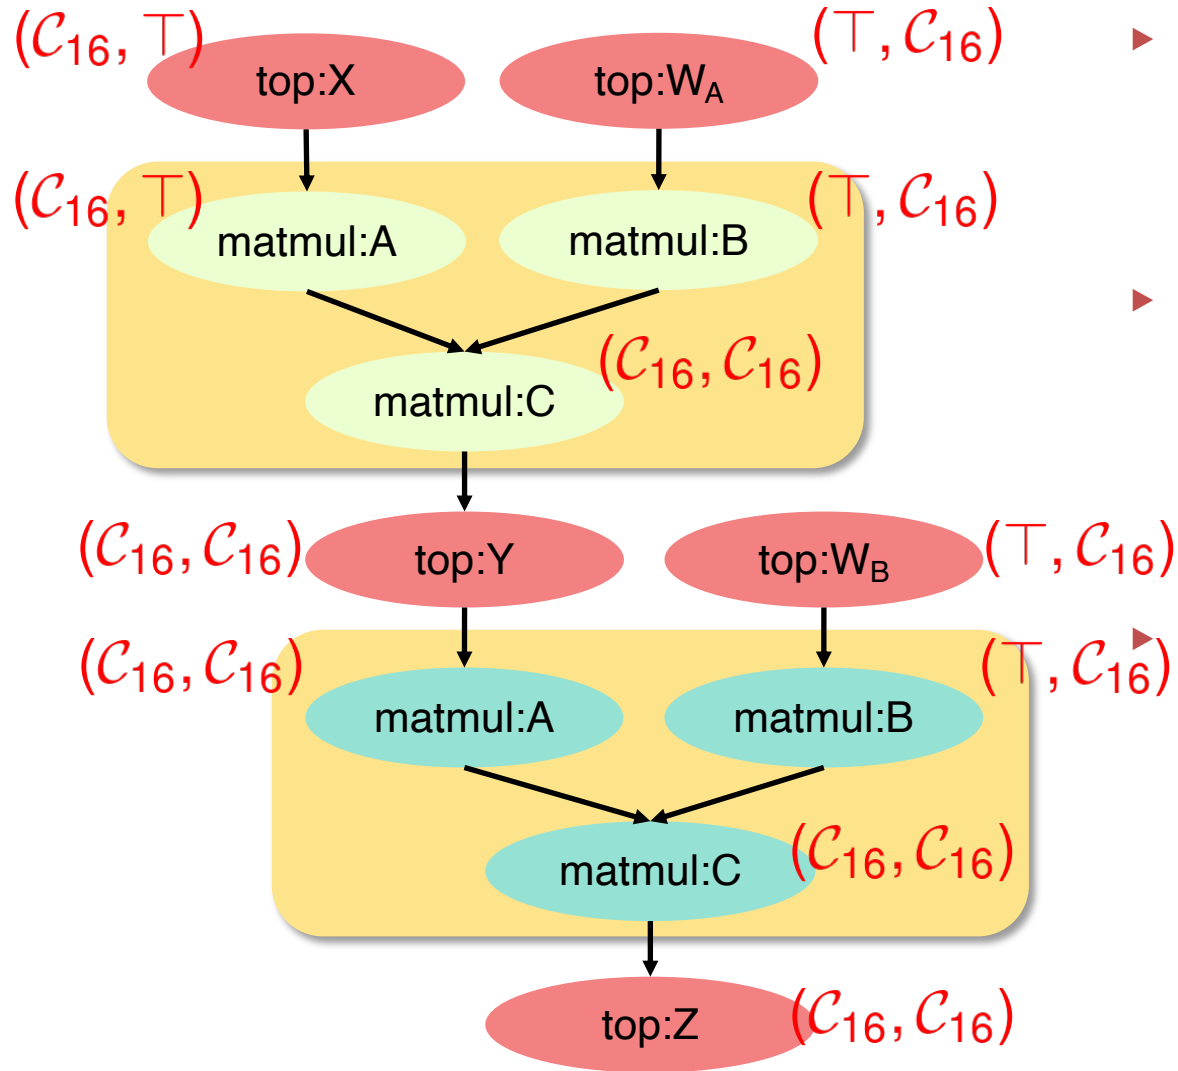
Subtyping relation forms a **lattice**!



Intuition:

We can supply more read/write parallelism, but not less!

Composable Schedules




Hierarchical dataflow graph

- ▶ **Key idea: Model data layout as a type** to ensure the kernel interfaces are consistent
- ▶ Lattice-based subtyping relation permits **linear-time layout inference** through Worklist algorithm
- ▶ Proof in supplementary material

Single-Kernel Evaluation

- ▶ Benchmarks from PolyBench; Target hardware: AMD U280 FPGA
- ▶ Normalized against AMD VitisHLS-auto (pragmas automatically inserted)
- ▶ Other baselines
 - ADLs: HeteroCL [FPGA'19], Dahlia [PLDI'20], PyLog [TC'21]
 - Automated DSE for HLS: ScaleHLS [HPCA'22], Merlin [TRETS'22]



Allo achieves (much) higher performance by optimizing data placement with a custom memory hierarchy

Single-Kernel Evaluation

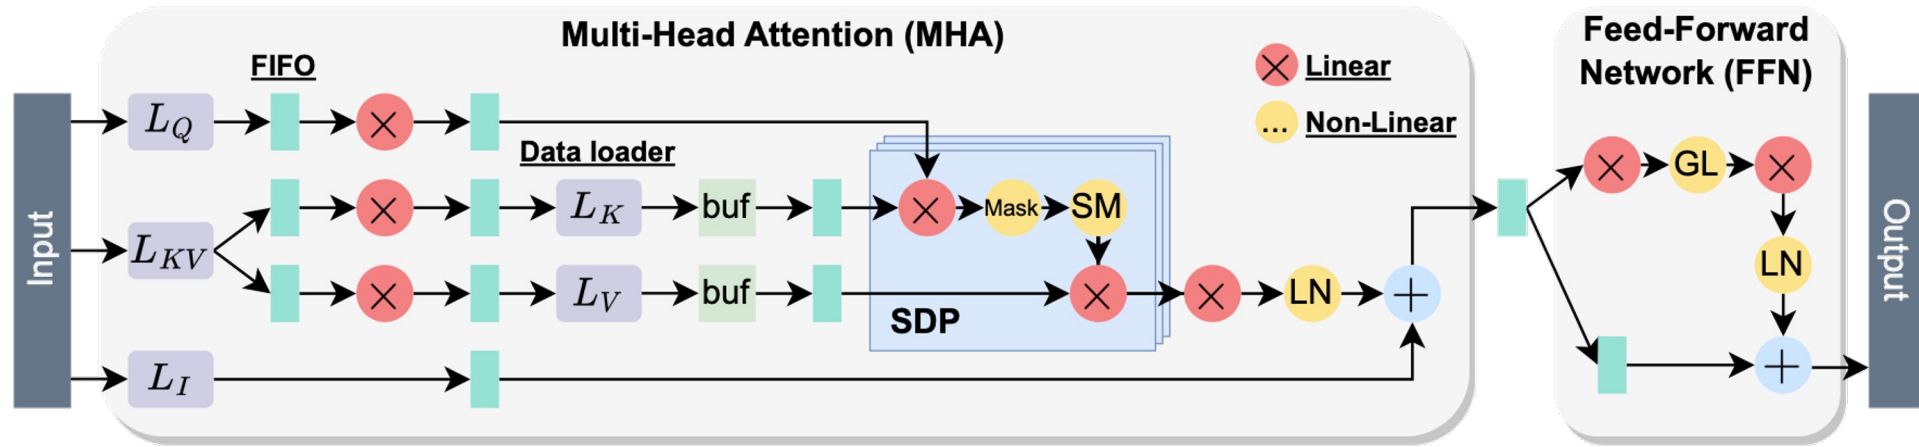
- ▶ Compared to ScaleHLS [HPCA'22], Allo achieves
 - Lower latency with much more effective use of compute resources
 - Higher post place-and-route frequency due to better pipelining

Benchmark	Allo						ScaleHLS				
	Latency (cycles)	II	DSP Usage	PnR Freq. (MHz)	Lines of Allo Custm.	Compile Time (s)	Latency (cycles)	II	DSP Usage	PnR Freq. (MHz)	Compile Time (s)
atax	4.9K (↓ 3.9×	1	403 (↑ 2.9×	411	9	1.0	19.4K	4	141	329	36.1
correlation	498.7K (↓ 290.5×	1	4168 (↑ 38.2×	362	19	0.8	144.9M	667	109	305	638.8
jacobi-2d	58.8K (↓ 183.1×	1	3968 (↑ 72.1×	411	17	0.9	10.8M	28	55	308	47.9
symm	405.7K (↓ 427.4×	1	1208 (↑ 201.3×	402	15	1.0	182.4M	13	6	397	3.5
trmm	492.6K (↓ 78.0×	1	101 (↑ 14.4×	414	12	0.8	38.4M	4	7	382	1.4

Allo achieves (much) higher performance by optimizing data placement with a custom memory hierarchy

Multi-Kernel Evaluation: A Complete LLM Accelerator

- ▶ GPT2 model (the only open-source LLM in the GPT family)
 - 355M parameters, 24 hidden layers, 16 heads
 - W4A8 quantization




```
def GPT_layer(inp: float32[L, D], ...)
    -> float32[L, D]:
    # 1. Multi-Head Attention (MHA)
    Q = Linear_layer_qkv(inp, Wq, Bq)
    K = Linear_layer_qkv(inp, Wk, Bk)
    V = Linear_layer_qkv(inp, Wv, Bv)
    attn_sf_outp = Self_attention(Q, K, V)
    # ...
    # 2. Feed Forward Network (FFN)
    # ...
    return ffn_res_outp
```

Compose all the schedules together

```
s = allo.customize(GPT_layer)
s.compose([s_qkv, ..., s_gelu])
```


LLM Accelerator Evaluation

- ▶ GPT2: single-batch, low-latency, generative inference settings
 - AMD U280 FPGA (16nm), 250MHz
 - **2.2x speedup in prefill stage** compared to DFX [MICRO'22] (an FPGA-based overlay)
 - **1.9x speedup for long output** sequences and **5.7x more energy-efficient** vs. NVIDIA A100 GPU (7nm)
 - Fewer than 50 lines of schedule code in Allo



	Allo	DFX
Device	U280	U280
Freq.	250MHz	200MHz
Quant.	W4A8	fp16
BRAM	384 (19.0%)	1192 (59.1%)
DSP	1780 (19.73%)	3533 (39.2%)
FF	652K (25.0%)	1107K (42.5%)
LUT	508K (39.0%)	520K (39.9%)

High-Level PyTorch Frontend

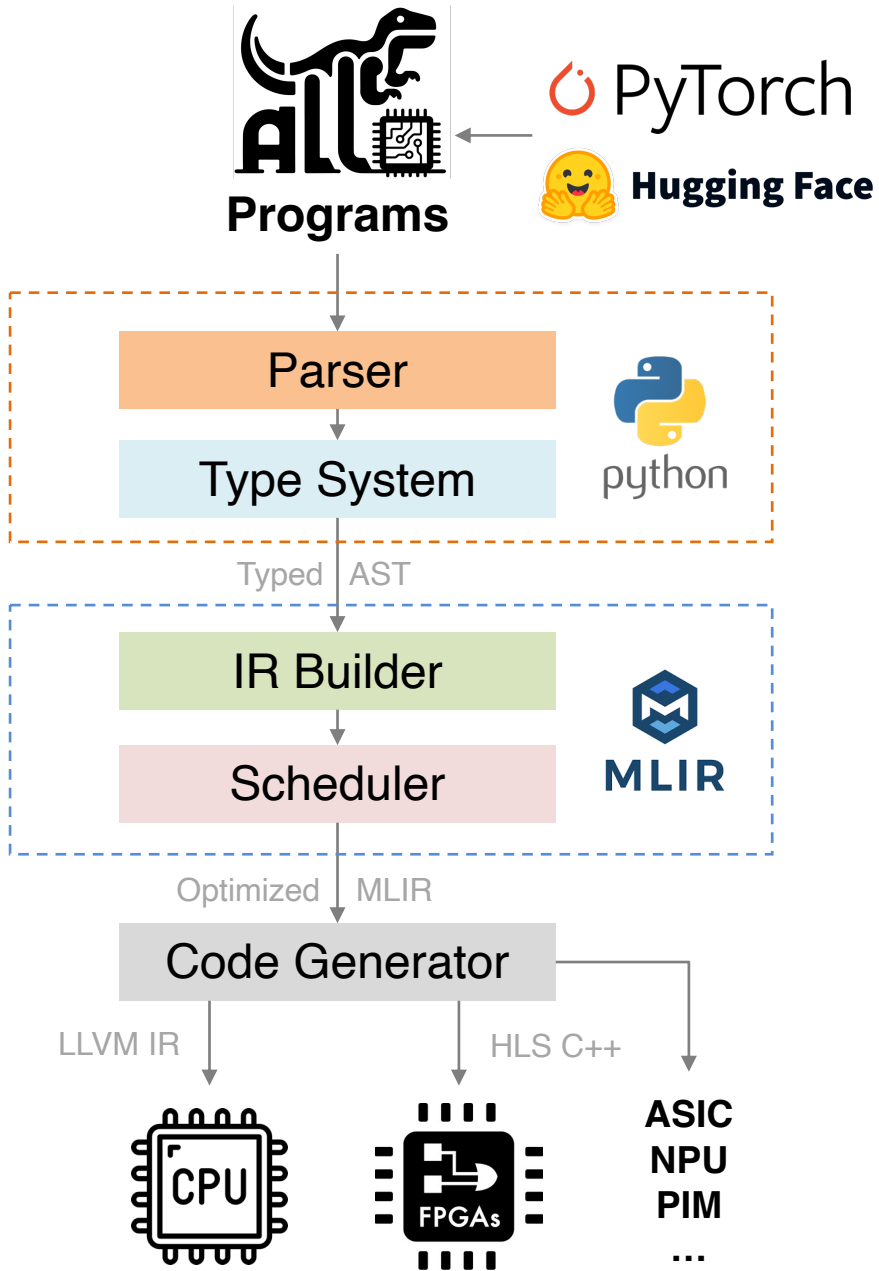
- ▶ Predefined schedules for commonly used NN operators
- ▶ Can directly import model from PyTorch and build optimized xcel design
 - Through TorchDynamo and torch.fx

```
import torch
import allo
import numpy as np
from transformers import AutoConfig
from transformers.models.gpt2.modeling_gpt2 import GPT2Model

bs, seq, hs = 1, 512, 1024

example_inputs = [torch.rand(bs, seq, hs)]
config = AutoConfig.from_pretrained("gpt2")
module = GPT2Model(config).eval()
mlir_mod = allo.frontend.from_pytorch(
    module,
    example_inputs=example_inputs,
)
```

Summary



- Features of Allo ADL
 - Pythonic
 - Decoupled & verifiable customizations
 - Composability



<https://github.com/cornell-zhang/allo>

- Ongoing work
 - ASIC & other programmable accelerators
 - NPUs (e.g., AMD AIE)
 - Processing-in-memory (PIM)
 - Automatic generation of high-performance schedules
 - Sparsity support
 - Integration with UniSparse [OOPSLA'24]

