

Large Language Model Acceleration with Allo


Hongzheng Chen

PI: Zhiru Zhang

Cornell University

UIUC-HACC


04/10/2024



Cornell University



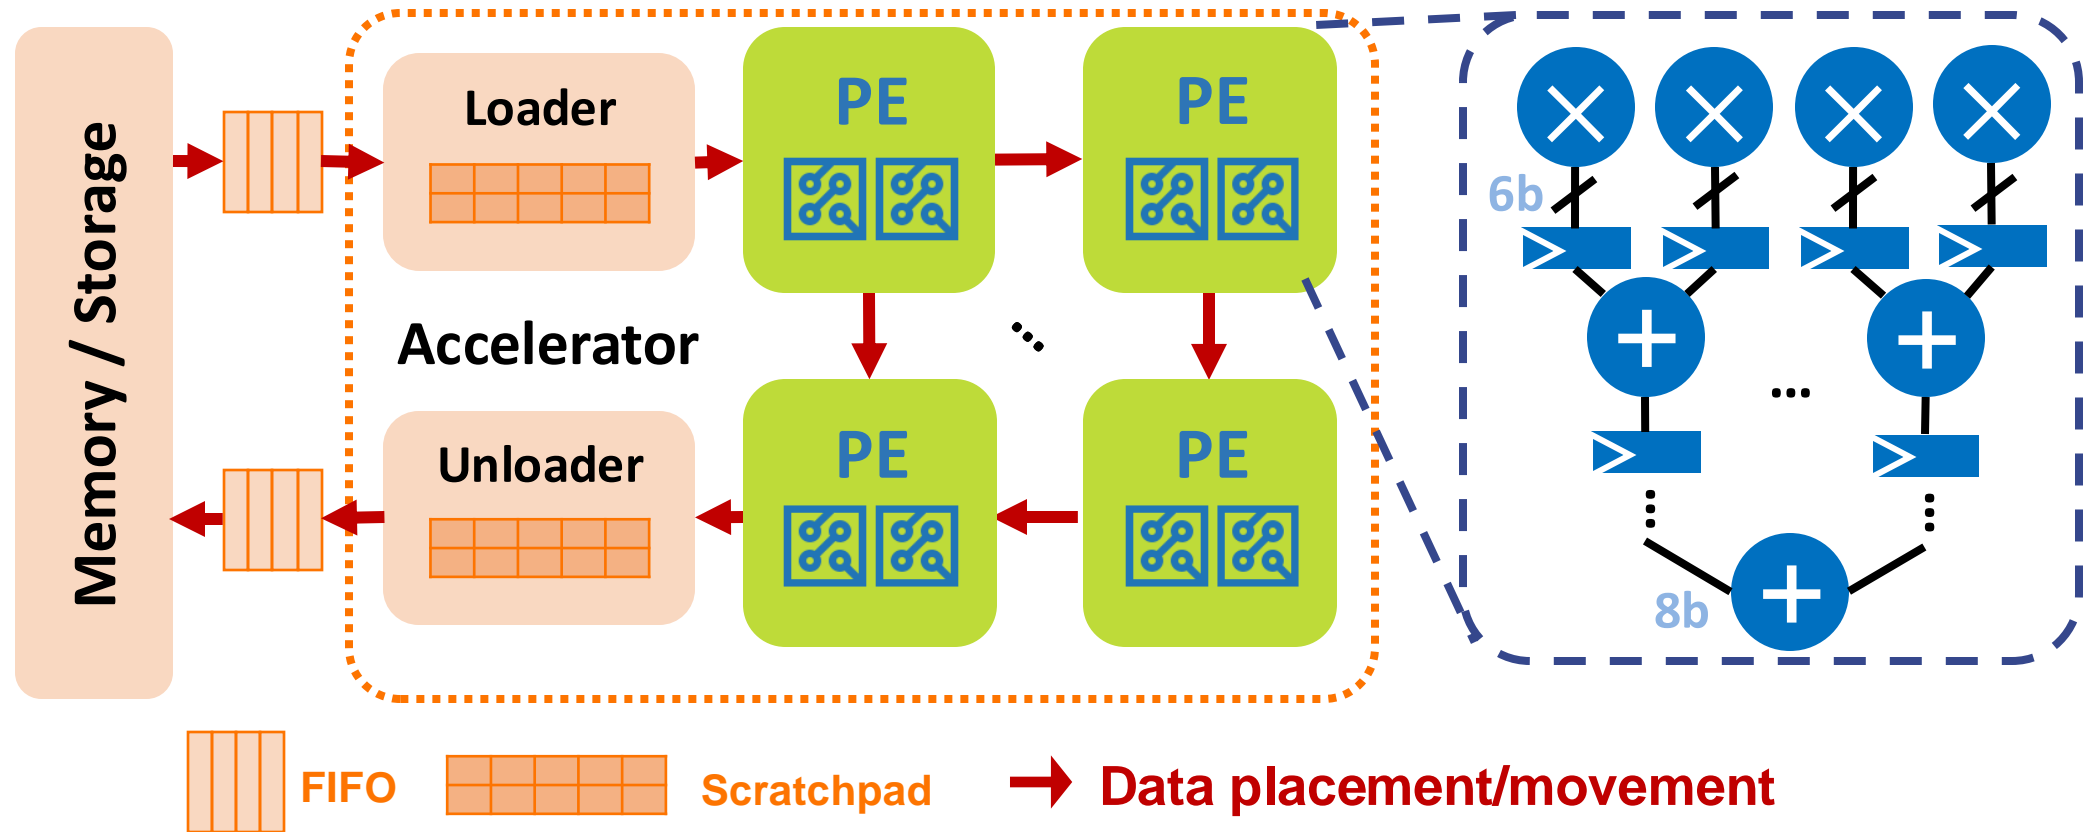
The Era of Large Language Models (LLMs)



- The growth in computation demands outpaces the increase in compute power offered by current hardware
- Special-purpose hardware accelerator for large deep learning models (e.g., Google TPU, AWS Inferentia)

Specialized Accelerator Design

- ▶ Accelerator design is different from programming on general processors
 - Custom processing engines (PEs)
 - Custom non-standard data type
 - Custom memory hierarchy
 - Custom data communication



Challenge 1: Balancing Manual Control & Compiler Optimization

(a) Manual optimization

- 😊 Fine-grained optimizations, high performance
- ☐ Rewrite the application code, hard to maintain

```
void conv1(...) {  
  #pragma HLS array_partition variable=Filter dim=0  
  hls::LineBuffer<3, N, ap_fixed<8,4> > buf;  
  hls::Window<3, 3, ap_fixed<8,4> > window;  
  for(int y = 0; y < N; y++) {  
    for(int xo = 0; xo < N/M; xo++) {  
      #pragma HLS pipeline II=1  
      for(int xi = 0; xi < M; xi++) {
```

Custom compute
(Loop tiling)

```
        int x = xo*M + xi;  
        ap_fixed<8,4> acc = 0;  
        ap_fixed<8,4> in = Input[y][x];  
        buf.shift_up(x);  
        buf.insert_top(in, x);  
        window.shift_left();  
        for(int r = 0; r < 2; r++)  
            window.insert(buf.getval(r,x),  
                          i, 2);  
        window.insert(in, 2, 2);
```

Custom data type
(Quantization)

Custom data
placement
(Reuse buffers)

```
        if (y >= 2 && x >= 2) {  
            for(int r = 0; r < 3; r++) {  
                for(int c = 0; c < 3; c++) {  
                    acc += window.getval(r,c) * Filter[r][c];  
                }  
            }  
            Out[y-2][x-2] = acc;  
        }  
    }  
}
```

(b) Compiler optimization

- 😊 Fully automatic, least manual effort
- ☐ No control on memory/communication, not general, hard to debug

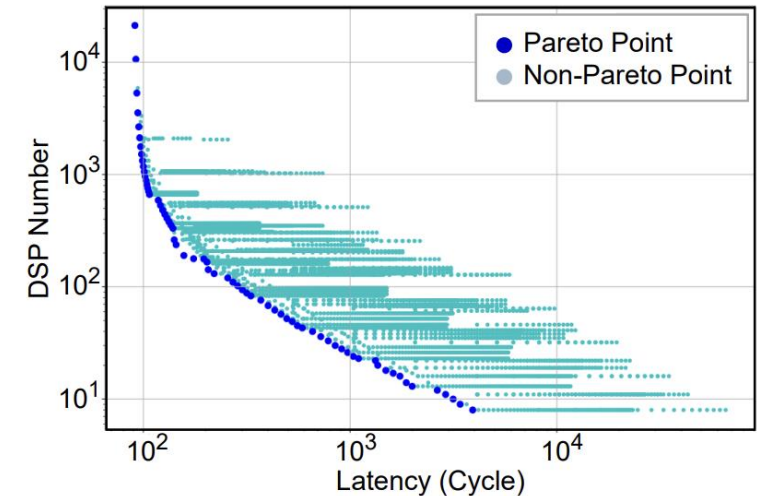


Fig source: ScaleHLS [HPCA'22]


- DSE only searches for hyperparams (e.g., tiling factors) but not program transformations (e.g., buffering)
- No transformation correctness guarantees

Challenge 2: Bridging the Gap from Single-Kernel to Multi-Kernel Design

- Existing accelerator design languages (ADLs) only consider opt. inside a kernel
 - e.g., HeteroCL[FPGA'19], Spatial[PLDI'18], Dahlia[PLDI'20]

```
27 // Specify the accelerator design
28 Accel {
29   // Produce C in M x N tiles
30   Foreach (A.rows by M, B.cols by N) { (ii,jj) =>
31     val tileC = SRAM[Half] (M, N)
32
33     // Combine intermediates across common dimension
34     MemReduce (tileC) (A.cols by P) { kk =>
35       // Allocate on-chip scratchpads
36       val tileA = SRAM[Half] (M, P)
37       val tileB = SRAM[Half] (P, N)
38       val accum = SRAM[Half] (M, N)
39
40       // Load tiles of A and B from DRAM
41       tileA load A(ii::ii+M, kk::kk+P) // M x P
42       tileB load B(kk::kk+P, jj::jj+N) // P x N
43
44       // Combine intermediates across a chunk of P
45       MemReduce (accum) (P by 1 par PAR_K) { k =>
46         val partC = SRAM[Half] (M, N)
47         Foreach (M by 1, N by 1 par PAR_J) { (i,j) =>
48           partC(i,j) = tileA(i,k) * tileB(k,j)
49         }
50         partC
51       // Combine intermediates with element-wise add
52       }{(a,b) => a + b }
53       }{(a,b) => a + b }
54
55       // Store the tile of C to DRAM
56       C(ii::ii+M, jj::jj+N) store tileC
57     }
58 }
```

(a) Code snippet in Spatial



Host-accelerator

```
void blur(DTYPE* input0, ..., DTYPE* input6,
DTYPE* output0, ..., DTYPE* output6) {
  #pragma HLS interface port=input0 bundle=g0 burst=32
  #pragma HLS interface port=input1 bundle=g1 burst=32
```

```
  stream<DTYPE> fifo_in[8], fifo_out[8];
  input_io_schedule(fifo_in, input0, ..., input6);
```

```
  #pragma HLS dataflow
  #pragma HLS stream var=fifo_inter[0] depth=32
  #pragma HLS stream var=fifo_inter[1] depth=32
  conv1(fifo_in, fifo_inter);
  conv2(fifo_inter, fifo_out);
```

Inter-kernel

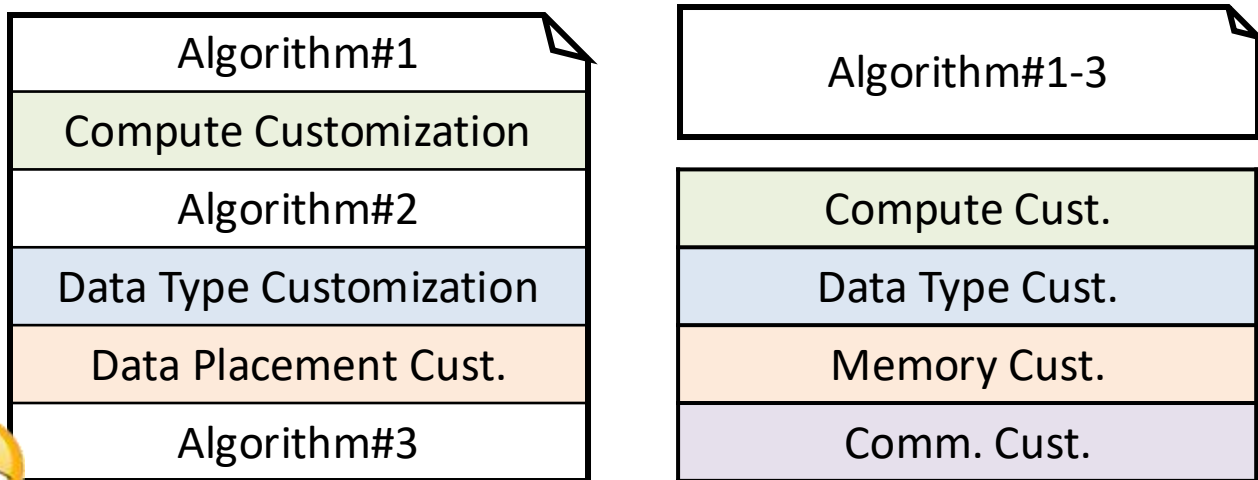
```
  output_io_schedule(fifo_out, output0, ..., output6);
}
```

(b) Code snippet in C++ HLS

Allo: A Programming Model for Composable Accelerator Design [PLDI'24]

Challenge 1: Manual vs compiler opt.

⚙️ Progressive hardware customization



Entangled algorithm specs and customization schemes


Fully decoupled customization schemes (i.e., schedule)

📄 Reusable parameterized kernel templates


```
1 def systolic[TyA, TyB, TyC, Mt: index, Nt: index, K: index]
2   (A: TyA[Mt, K], B: TyB[K, Nt], C: TyC[Mt, Nt])
3
4 def tiled_systolic[TyA, TyB, TyC, M: index, N: index, K: index]
5   (A: TyA[M, K], B: TyB[K, N], C: TyC[M, N]):
6   local_A: TyA[8, K]; local_B: TyB[K, 8]; local_C: TyC[8, 8]
7   for mi, ni in allo.grid(M // 8, N // 8, name="outer_tile"):
8     # ... load_A_tile, load_B_tile
9     systolic[TyA, TyB, TyC, 8, 8, K](local_A, local_B, local_C)
10    # ... store_C_tile
```

Challenge 2: Single -> Multi-kernel


🔗 Composable schedules



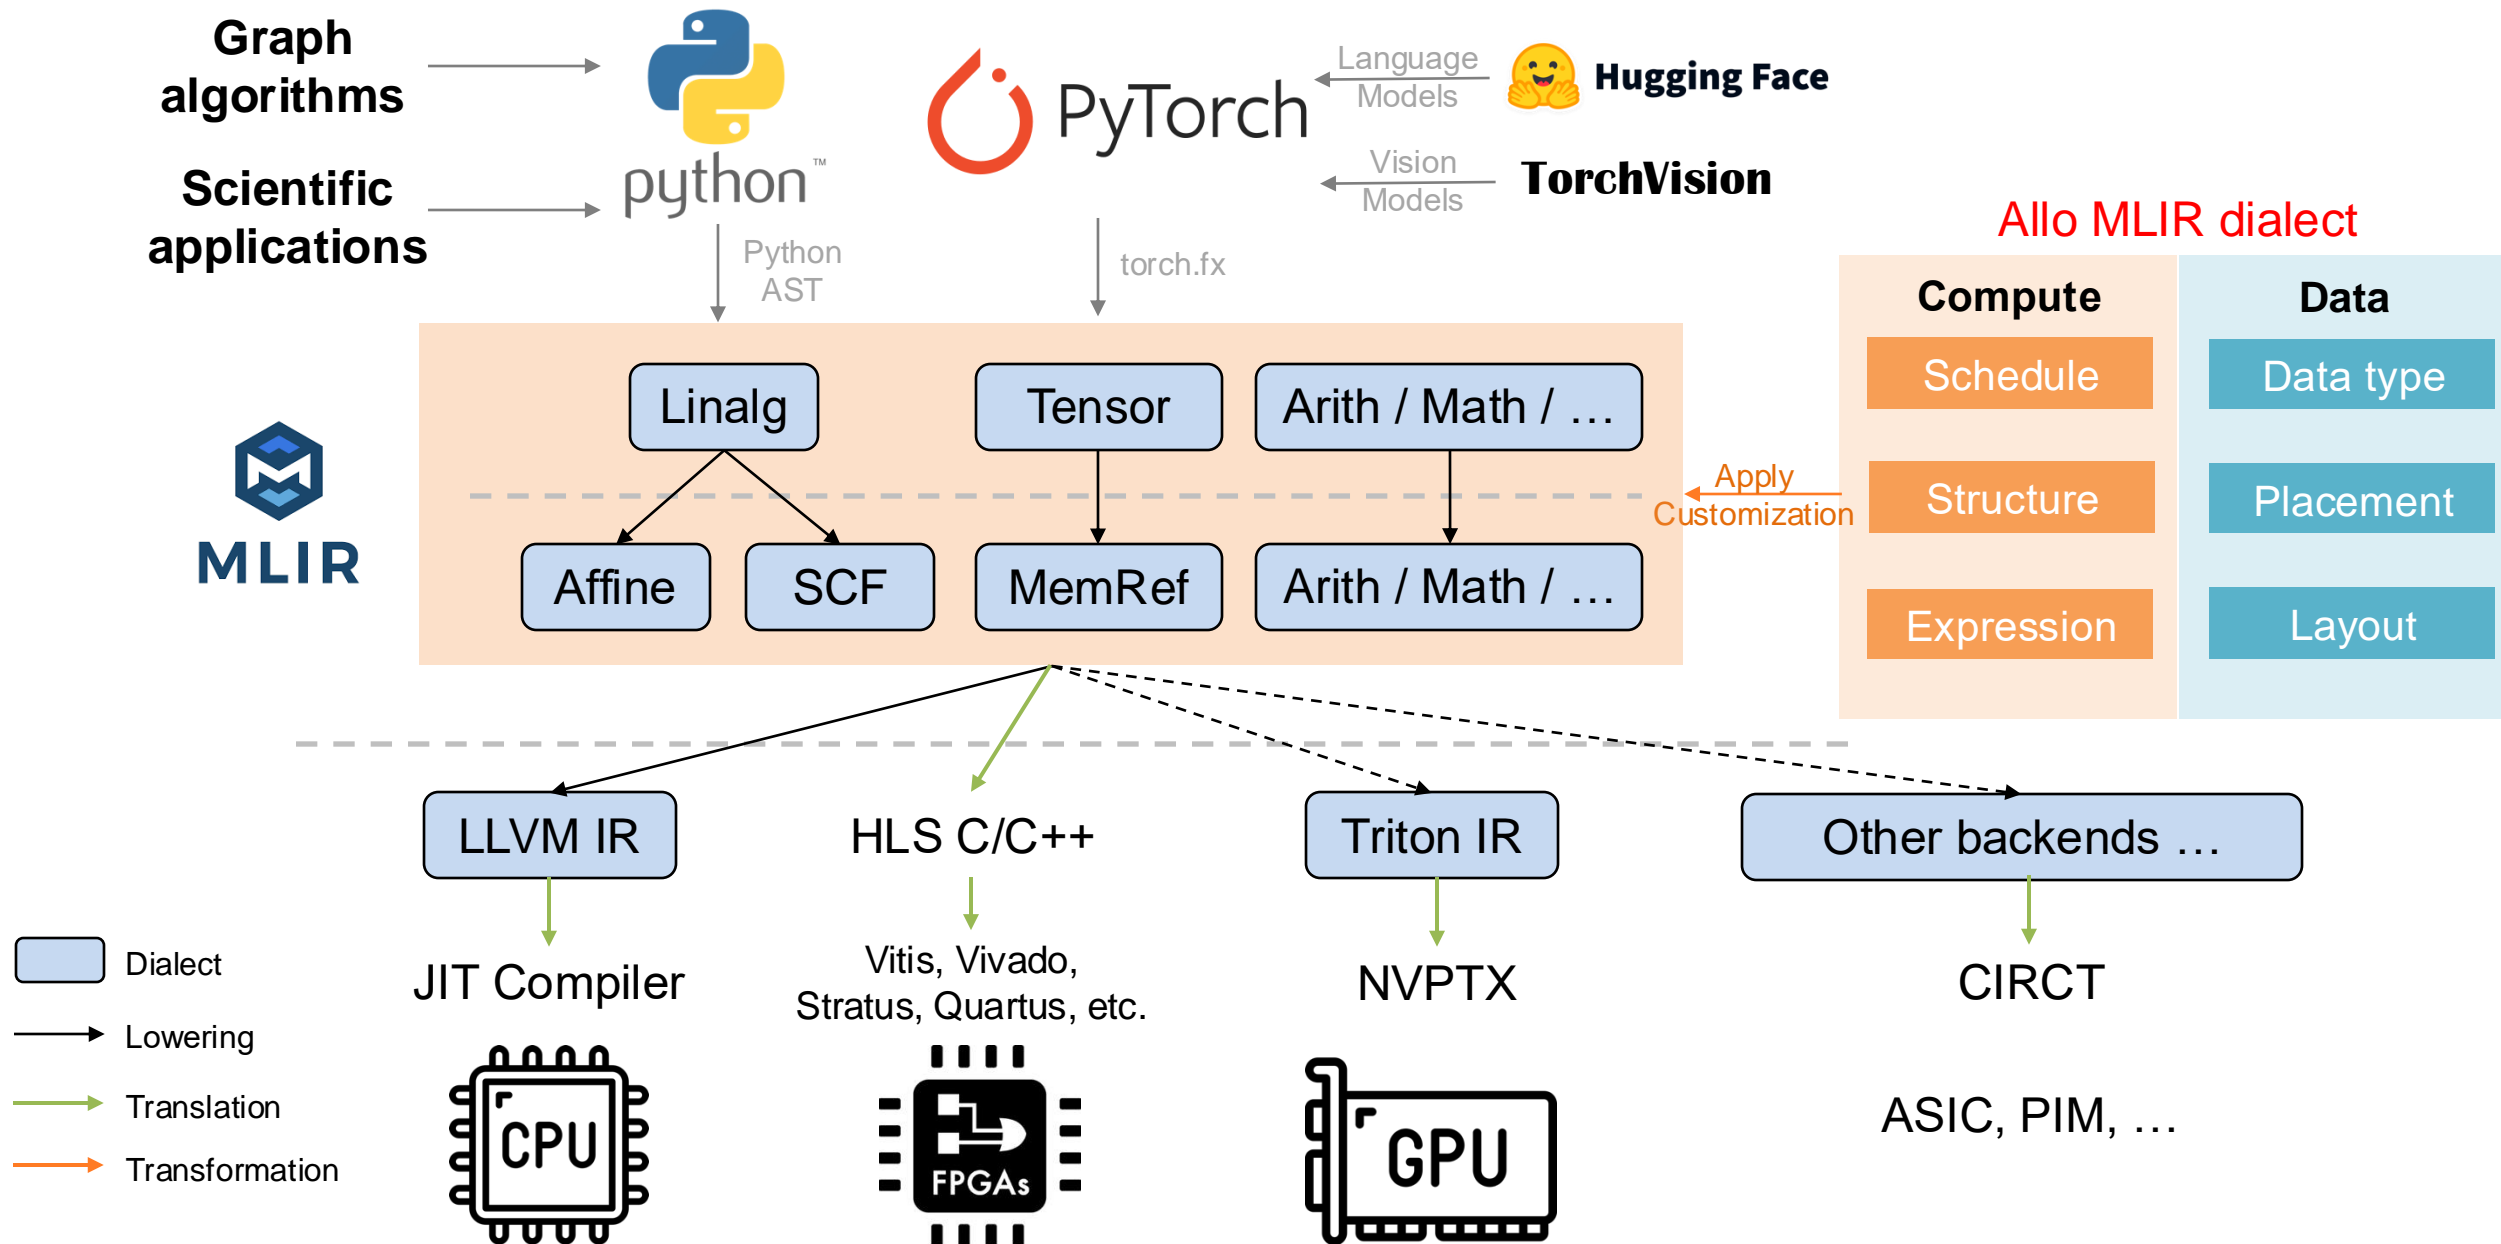
Type-safe composition



⚖️ Holistic dataflow optimizations



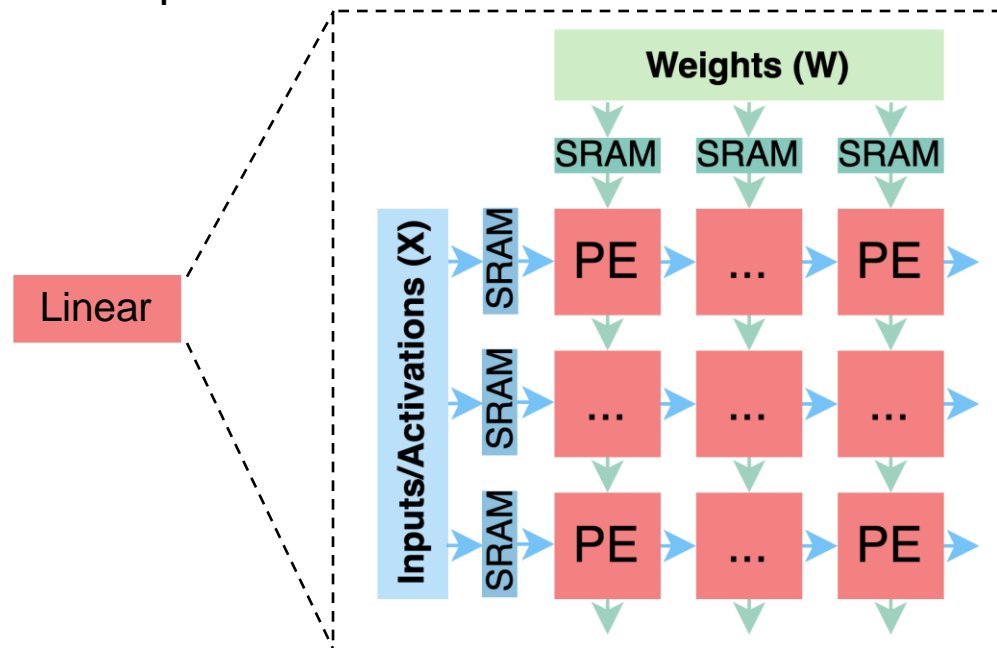
Overview of Allo ADL and Compilation Stack



Goal: Design a High-Performance LLM Accelerator

Step 1: Construct building blocks

- ▶ Performance
- ▶ Correctness
- ▶ Reusability
- Linear operators



- Non-linear operators

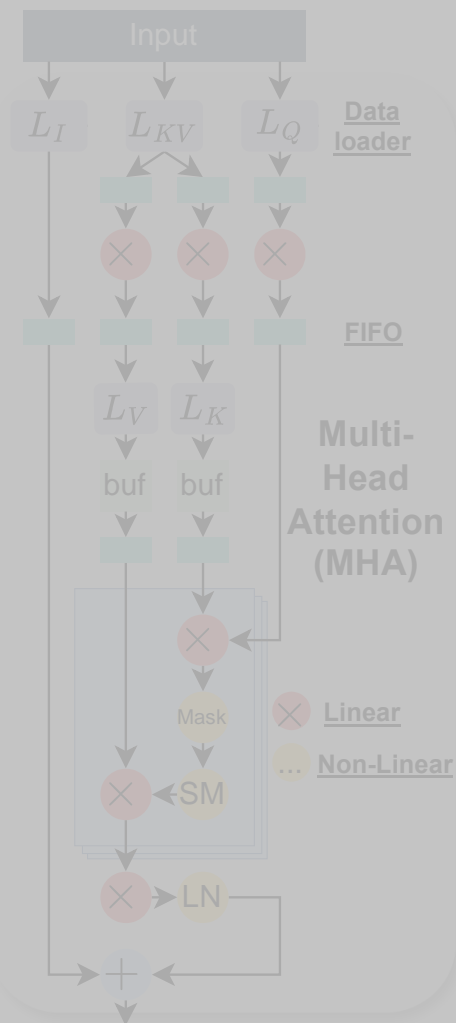
Softmax

LayerNorm

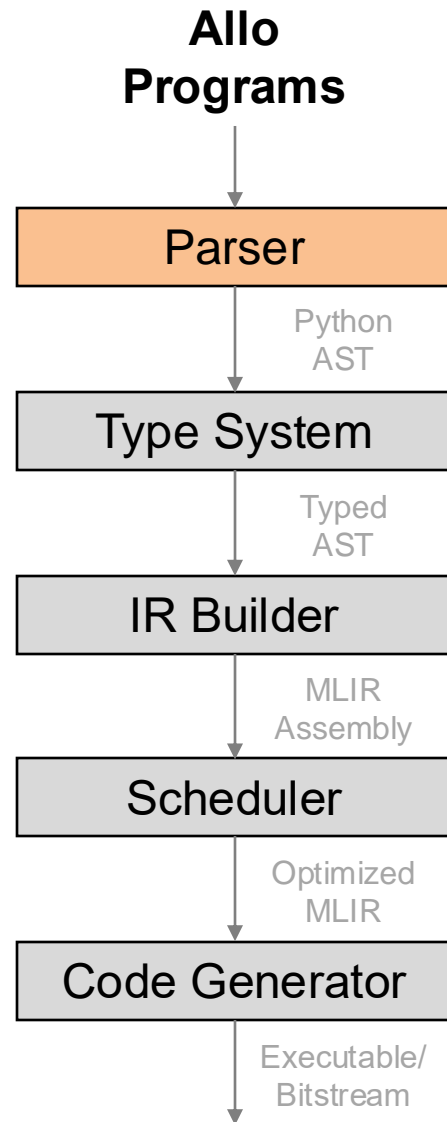
GELU

Step 2: Ensemble into a complete design

- ▶ Connect different operators bottom-up



Allo ADL Compilation Stack



Pythonic: No need to learn a new DSL!

- Free-form imperative programming
- Python native keywords (e.g., for, if, else)

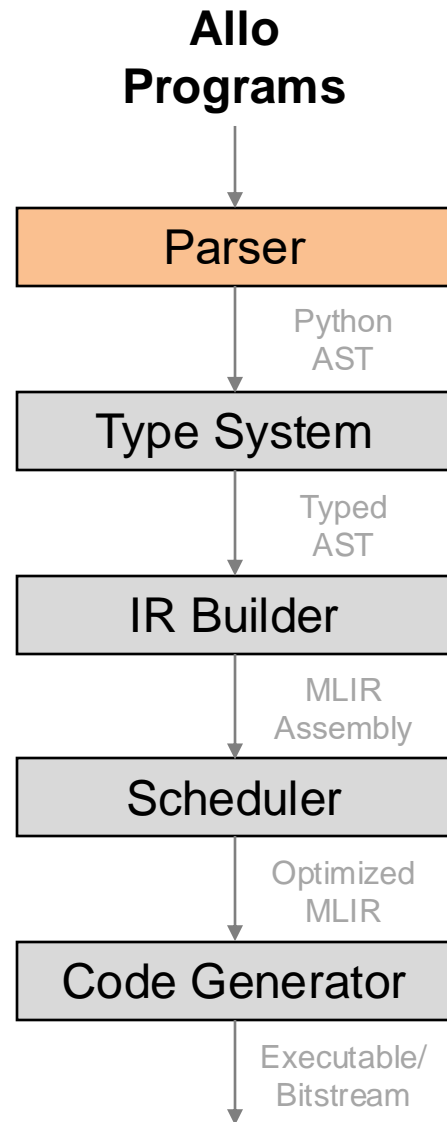


```
def gemm(A: int8[M, K],  
        B: int8[K, N]) -> int8[M, N]:  
    C: int8[M, N] = 0  
    for i, j, k in allo.grid(M, N, K):  
        C[i, j] += A[i, k] * B[k, j]  
    return C
```

Red boxes highlight the function definition 'def gemm', the 'allo.grid' call, the array access 'A[i, k]', and the 'return C' statement. Red arrows point from the 'Pythonic' header to these specific parts of the code.

```
s = allo.customize(gemm)
```

Allo ADL Compilation Stack



Pythonic: No need to learn a new DSL!

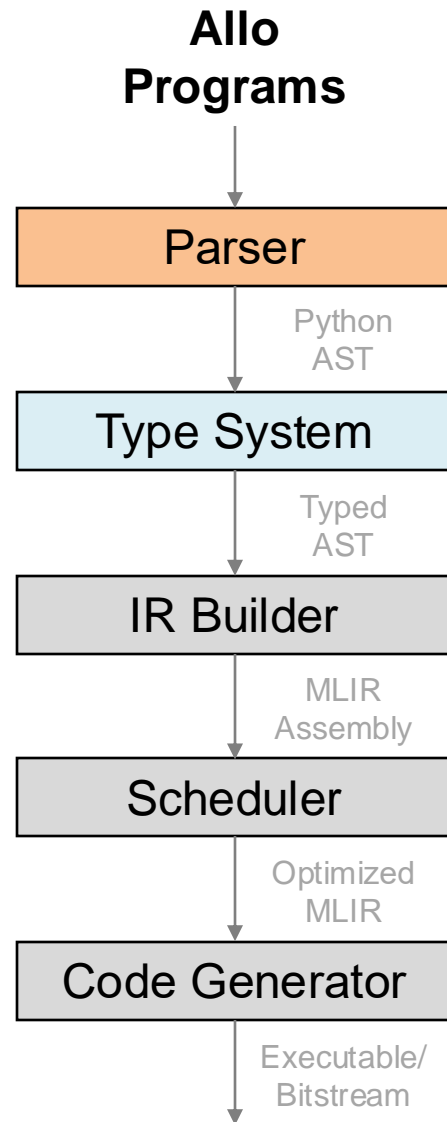
- Free-form imperative programming
- Python native keywords (e.g., for, if, else)
- Explicit type annotation

```
def gemm(A: int8[M, K],  
         B: int8[K, N]) -> int8[M, N]:  
    C: int8[M, N] = 0  
    for i, j, k in allo.grid(M, N, K):  
        C[i, j] += A[i, k] * B[k, j]  
    return C
```

```
s = allo.customize(gemm)
```

```
T ::= eleTy | eleTy [ shape ]  
shape ::= dim | dim , shape  
dim ::= [0-9]+ | constvar | expr  
expr ::= dim binop dim  
binop ::= + | - | * | /  
eleTy ::= intTy | floatTy | fixedTy  
intTy ::= int[0-9]+ | uint[0-9]+  
floatTy ::= float16 | float32 | float64  
fixedTy ::= fixed(width, frac)  
           | ufixed(width, frac)  
constvar ::= [a-zA-Z_][a-zA-Z0-9_]*  
width ::= [0-9]+  
frac ::= [0-9]+
```

Allo ADL Compilation Stack



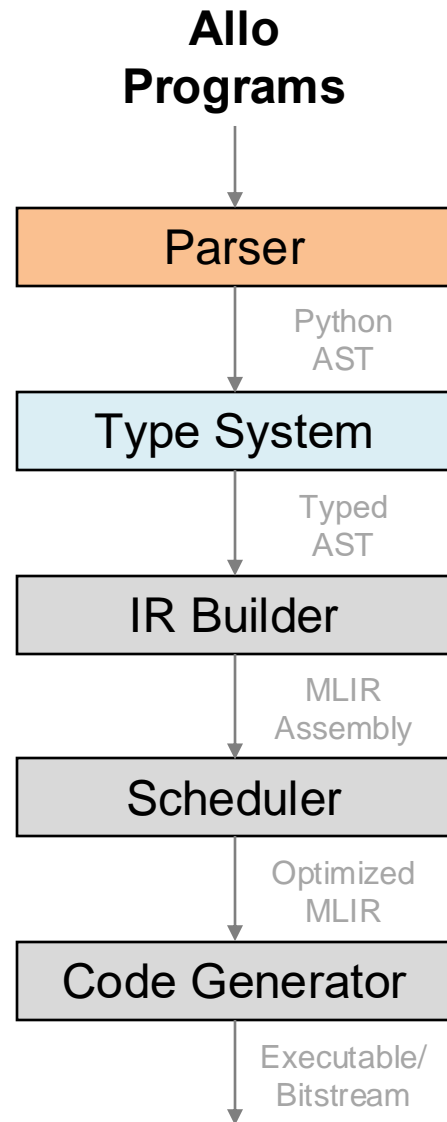
Type System

- Type checking: Inferred type different from annotated
- Type conversion: Type inference and implicit casting
- Shape propagation: Array broadcasting

```
M, N, K = 32, 32, 32
```

```
def gemm(A: int8[M, K], B: int8[K, N]) -> int8[M, N]:  
    C: int8[M, N] = 0  
    for i, j in allo.grid(M, N):  
        v: int14 = 0  
        for k in range(K):  
            v += A[i, k] * B[k, j]  
        C[i, j] = v  
    return C
```

Allo ADL Compilation Stack



Type System


- Type checking: Inferred type different from annotated
- Type conversion: Type **inference** and implicit **casting**
- Shape propagation: Array **broadcasting**

Type inling & constant folding

```
def gemm(A: int8[32, 32], B: int8[32, 32])  
    -> int8[32, 32]:  
    C: int8[32, 32] = 0: int8 -> int8[32, 32]  
    for i: index, j: index in allo.grid(M, N):  
        v: int14 = 0: int14  
        for k: index in range(K):  
            v: int14 += (A[i, k]: int8  
                        * B[k, j]: int8) -> int16 -> int14  
        C[i, j]: int8 = v: int14 -> int8  
    return C: int8[32, 32]
```

The code snippet shows a function `gemm` that takes two 32x32 integer arrays `A` and `B` and returns a 32x32 integer array `C`. The function uses a nested loop structure to compute the dot product of rows and columns. Annotations include type inference (e.g., `int14` for the accumulator `v`) and implicit casting (e.g., `int8` to `int16` to `int14`).

Allo ADL Compilation Stack



Type System

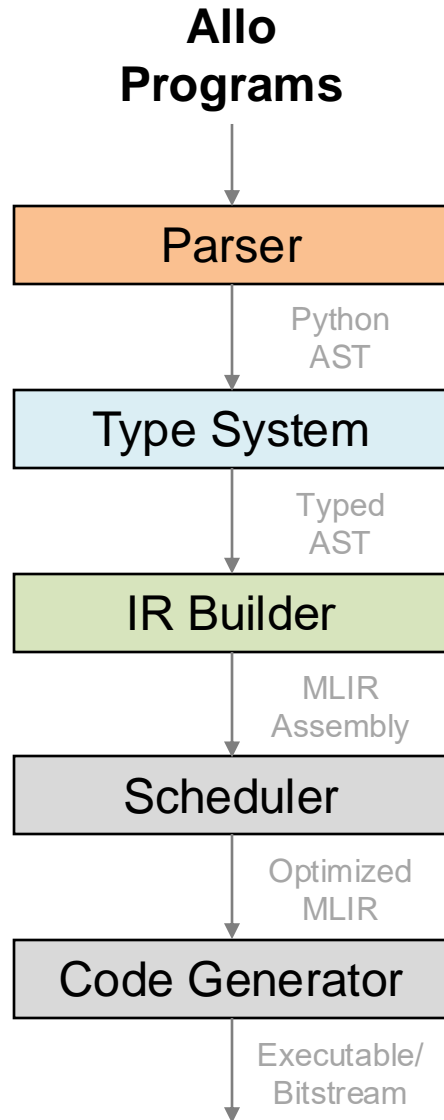
- Type checking: Inferred type different from annotated
- Type conversion: Type **inference** and implicit **casting**
- Shape propagation: Array **broadcasting**
- Quantization: Automatic fixed point to int conversion

```
T_IN, T_OUT = Fixed(7, 1), Fixed(15, 4)
```

```
def gemm(A: T_IN[M, K], B: T_IN[K, N]) -> T_OUT[M, N]:  
    C: T_OUT[M, N] = 0  
    for i, j, k in allo.grid(M, N, K, name="C"):  
        C[i, j] += A[i, k] * B[k, j]  
    return C
```

```
s = allo.customize(gemm)
```

Allo ADL Compilation Stack



IR Builder


- MLIR builtin dialects
- One-to-one mapping, strictly follow program structure

```
def gemm(A: int8[M, K],  
         B: int8[K, N])  
    -> int8[M, N]:  
    C: int8[M, N] = 0  
    for i, j, k in allo.grid(M, N, K):  
        C[i, j] += A[i, k] * B[k, j]  
    return C
```

```
s = allo.customize(gemm)  
print(s.module)
```

```
module {  
  func.func @gemm(%arg0: memref<32x32xi8>,  
                  %arg1: memref<32x32xi8>)  
    -> memref<32x32xi8> {  
    %0 = memref.alloc() {name = "C"} : memref<32x32xi8>  
    %c0_i8 = arith.constant 0 : i8  
    linalg.fill ins(%c0_i8 : i8) outs(%0 : memref<32x32xi8>)  
    affine.for %arg2 = 0 to 32 {  
      affine.for %arg3 = 0 to 32 {  
        affine.for %arg4 = 0 to 32 {  
          ...  
        } {loop_name = "k"}  
      } {loop_name = "j"}  
    } {loop_name = "i", op_name = "S_i_j_k_0"}  
    return %0 : memref<32x32xi8>  
  }  
}
```

Allo ADL Compilation Stack



Scheduler


- Decoupled customizations
- Real-time transformation

```
def gemm(A: int8[M, K],  
         B: int8[K, N])  
    -> int8[M, N]:  
    C: int8[M, N] = 0  
    for i, j, k in allo.grid(M, N, K):  
        C[i, j] += A[i, k] * B[k, j]  
    return C
```

```
s = allo.customize(gemm)  
i_out, i_in = s.split("i", 8)  
print(s.module)
```

```
#map = affine_map<(d0, d1) -> (d0 + d1 * 8)>  
module {  
    func.func @gemm(%arg0: memref<32x32xi8>,  
                   %arg1: memref<32x32xi8>)  
        -> memref<32x32xi8> {  
        %0 = memref.alloc() {name = "C"} : memref<32x32xi8>  
        %c0_i8 = arith.constant 0 : i8  
        linalg.fill ins(%c0_i8 : i8) outs(%0 : memref<32x32xi8>)  
        affine.for %arg2 = 0 to 4 {  
            affine.for %arg3 = 0 to 8 {  
                affine.for %arg4 = 0 to 32 {  
                    affine.for %arg5 = 0 to 32 {  
                        ...  
                    } {loop_name = "k"}  
                } {loop_name = "j"}  
            } {loop_name = "i.inner"}  
        } {loop_name = "i.outer", op_name = "S_i_j_k_0"}  
        return %0 : memref<32x32xi8>  
    }  
}
```

Allo ADL Compilation Stack



Scheduler

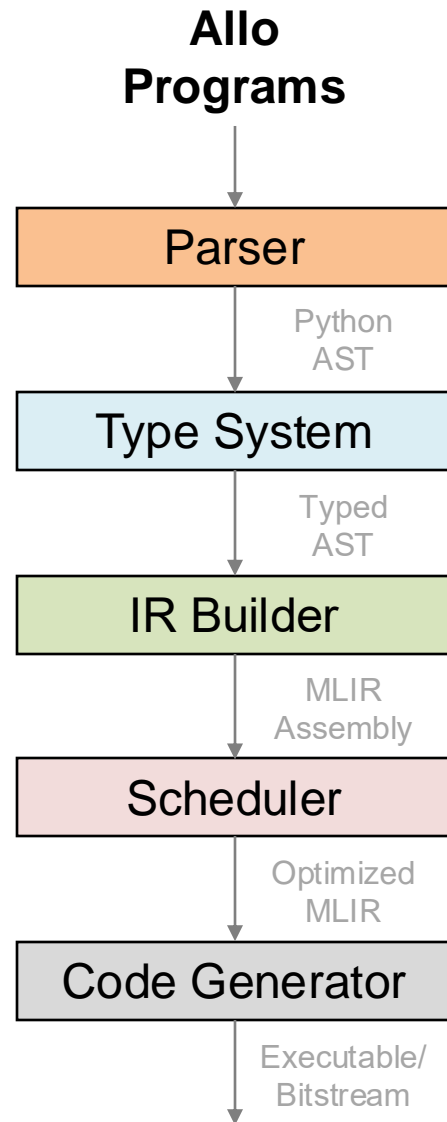
- Decoupled customizations
- Real-time transformation

```
def gemm(A: int8[M, K],  
         B: int8[K, N])  
    -> int8[M, N]:  
    C: int8[M, N] = 0  
    for i, j, k in allo.grid(M, N, K):  
        C[i, j] += A[i, k] * B[k, j]  
    return C
```

```
s = allo.customize(gemm)  
i_out, i_in = s.split("i", 8)  
print(s.module)  
j_out, j_in = s.split("j", 8)  
print(s.module)
```

```
#map = affine_map<(d0, d1) -> (d0 + d1 * 8)>  
module {  
    func func @gemm(%arg0: memref<32x32xi8>,  
                   %arg1: memref<32x32xi8>)  
        -> memref<32x32xi8> {  
        %0 = memref.alloc() {name = "C"} : memref<32x32xi8>  
        %c0_i8 = arith.constant 0 : i8  
        linalg.fill ins(%c0_i8 : i8) outs(%0 : memref<32x32xi8>)  
        affine.for %arg2 = 0 to 4 {  
            affine.for %arg3 = 0 to 8 {  
                affine.for %arg4 = 0 to 4 {  
                    affine.for %arg5 = 0 to 8 {  
                        affine.for %arg6 = 0 to 32 {  
                            ...  
                        } {loop_name = "k"}  
                    } {loop_name = "j.inner"}  
                } {loop_name = "j.outer"}  
            } {loop_name = "i.inner"}  
        } {loop_name = "i.outer"}  
        return %0 : memref<32x32xi8>}}
```


Allo ADL Compilation Stack



Scheduler

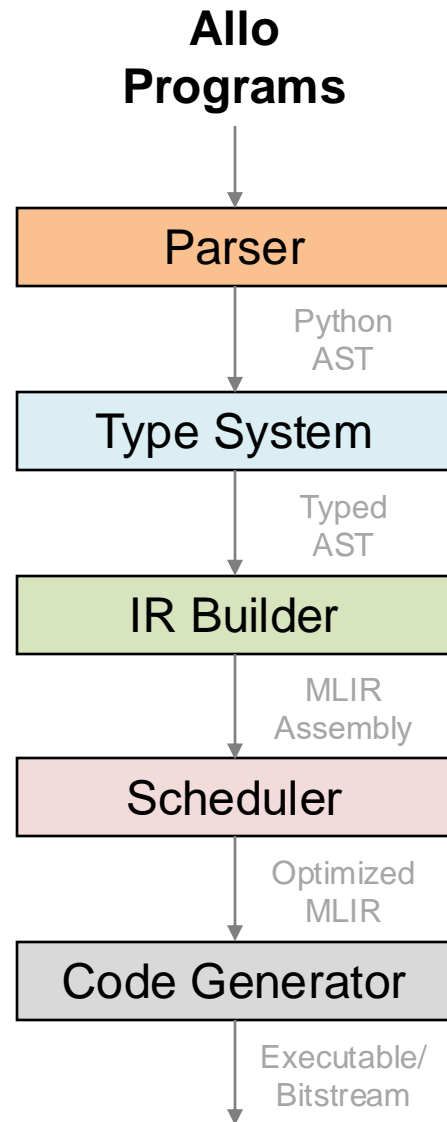
- Decoupled customizations
- Real-time transformation

```
def gemm(A: int8[M, K],  
         B: int8[K, N])  
    -> int8[M, N]:  
    C: int8[M, N] = 0  
    for i, j, k in allo.grid(M, N, K):  
        C[i, j] += A[i, k] * B[k, j]  
    return C
```

```
s = allo.customize(gemm)  
i_out, i_in = s.split("i", 8)  
print(s.module)  
j_out, j_in = s.split("j", 8)  
print(s.module)  
s.reorder(i_out, j_out, i_in, j_in)  
print(s.module)
```

```
#map = affine_map<(d0, d1) -> (d0 + d1 * 8)>  
module {  
    func.func @gemm(%arg0: memref<32x32xi8>,  
                  %arg1: memref<32x32xi8>)  
        -> memref<32x32xi8> {  
        %0 = memref.alloc() {name = "C"} : memref<32x32xi8>  
        %c0_i8 = arith.constant 0 : i8  
        linalg.fill ins(%c0_i8 : i8) outs(%0 : memref<32x32xi8>)  
        affine.for %arg2 = 0 to 4 {  
            affine.for %arg3 = 0 to 8 {  
                affine.for %arg4 = 0 to 4 {  
                    affine.for %arg5 = 0 to 8 {  
                        affine.for %arg6 = 0 to 32 {  
                            ...  
                        } {loop_name = "k"}  
                    } {loop_name = "j.inner"}  
                } {loop_name = "i.inner"}  
            } {loop_name = "j.outer"}  
        } {loop_name = "i.outer"}  
        return %0 : memref<32x32xi8>}}
```

Allo ADL Compilation Stack



Codegen

- Default: CPU backend for simulation testing

```
def gemm(A: int8[M, K],  
         B: int8[K, N])  
    -> int8[M, N]:  
    C: int8[M, N] = 0  
    for i, j, k in allo.grid(M, N, K):  
        C[i, j] += A[i, k] * B[k, j]  
    return C
```

```
s = allo.customize(gemm)  
# ... customizations (omitted)  
X = np.random.randint(-8, 8, size=(M, K)).astype(np.int8)  
A = np.random.randint(-8, 8, size=(K, N)).astype(np.int8)  
np_outs = gemm(X, A)
```

```
f = s.build(target="llvm")  
outs = f(X, A)  
np.testing.assert_allclose(outs, np_outs, atol=1e-3)
```

Allo ADL Example - Systolic Array


- ▶ Transform a vanilla GEMM implementation into a high-performance systolic array
- ▶ Only 8 lines of schedule code is needed!

```

1  # Algorithm specification
2  def gemm(A: int8[M, K], B: int8[K, N],
3          C: int16[M, N]):
4      for i, j in allo.grid(M, N, "PE"):
5          for k in range(K):
6              C[i, j] += A[i, k] * B[k, j]
7
8  # Schedule construction
9  s = allo.customize(gemm)
10 buf_A = s.buffer_at(s.A, "j")
11 buf_B = s.buffer_at(s.B, "j")
12 pe = s.unfold("PE", axis=[0, 1])
13 s.partition(s.C, dim=[0, 1])
14 s.partition(s.A, dim=0)
15 s.partition(s.B, dim=1)
16 s.relay(buf_A, pe, axis=1, depth=M + 1)
17 s.relay(buf_B, pe, axis=0, depth=N + 1)

```

p_0
 # p_1
 # p_2
 # p_3
 # p_4
 # p_5
 # p_6
 # p_7



First FFN layer in BERT-base (512, 768)x(768, 3072) w/ 16x16 SA

	Latency (ms)	BRAM	DSP	FF	LUT
Ours (w/o DSP packing)	15.73	0 (0%)	256 (2%)	88284 (3%)	168190 (12%)
Ours (w/ DSP packing)	15.73	0 (0%)	128 (1%)	79969 (3%)	244439 (18%)
AutoSA [75]	15.71	514 (12%)	256 (2%)	100138 (3%)	244032 (18%)

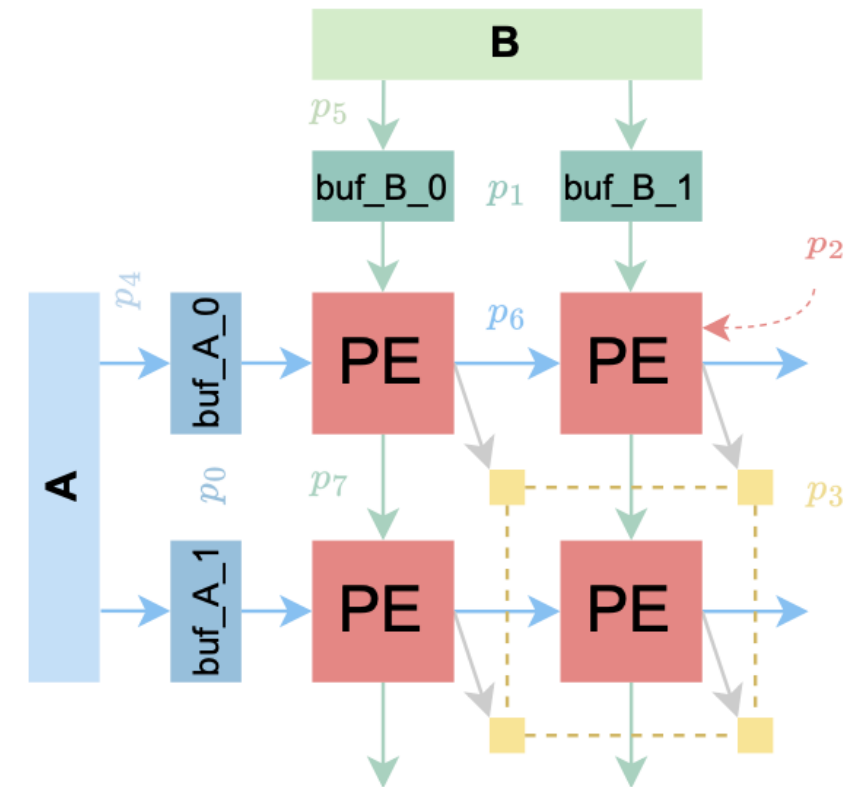
These schedules can indeed generate high-performance designs

Verifiable Schedule [FPGA'24 Best Paper]

- ▶ `s.verify(orig_sch, new_sch)`
 - Integrated CDAG verifier
 - Able to verify statically interpretable control-flow (SICF) programs

```
1  # Algorithm specification
2  def gemm(A: int8[M, K], B: int8[K, N],
3          C: int16[M, N]):
4      for i, j in allo.grid(M, N, "PE"):
5          for k in range(K):
6              C[i, j] += A[i, k] * B[k, j]
7
8  # Schedule construction
9  s = allo.customize(gemm)
10 buf_A = s.buffer_at(s.A, "j")
11 buf_B = s.buffer_at(s.B, "j")
12 pe = s.unfold("PE", axis=[0, 1])
13 s.partition(s.C, dim=[0, 1])
14 s.partition(s.A, dim=0)
15 s.partition(s.B, dim=1)
16 s.relay(buf_A, pe, axis=1, depth=M + 1)
17 s.relay(buf_B, pe, axis=0, depth=N + 1)
```

p_0
p_1
p_2
p_3
p_4
p_5
p_6
p_7




Parameterized Kernel Template

- Reusability: Parameterized templates
 - Parameter types: New feature in Python 3.12
 - Works like C++ template
 - Parameterizable shapes and types
 - Build kernel libraries in Python

```
1  def systolic[TyA, TyB, TyC, Mt: index, Nt: index, K: index]
2      (A: TyA[Mt, K], B: TyB[K, Nt], C: TyC[Mt, Nt])
3
4  def tiled_systolic[TyA, TyB, TyC, M: index, N: index, K: index]
5      (A: TyA[M, K], B: TyB[K, N], C: TyC[M, N]):
6      local_A: TyA[8, K]; local_B: TyB[K, 8]; local_C: TyC[8, 8]
7      for mi, ni in allo.grid(M // 8, N // 8, name="outer_tile"):
8          # ... load_A_tile, load_B_tile
9          systolic[TyA, TyB, TyC, 8, 8, K](local_A, local_B, local_C)
10         # ... store_C_tile
```

Single-Kernel Evaluation

- ▶ Normalized against VitisHLS auto baseline (no pragma inserted)
- ▶ Automated DSE: ScaleHLS, Merlin
- ▶ Human designed customizations: Allo, HeteroCL, PyLog, Dahlia



Single-Kernel Evaluation

- ▶ Compared to ScaleHLS DSE results, Allo achieves:
 - Lower II, II=1
 - More efficient use of computation resources
 - Higher PnR frequency due to better pipelined designs


Benchmark	Allo						ScaleHLS				
	Latency (cycles)	II	DSP Usage	PnR Freq. (MHz)	Lines of Allo Custm.	Compile Time (s)	Latency (cycles)	II	DSP Usage	PnR Freq. (MHz)	Compile Time (s)
atax	4.9K (↓ 3.9×	1	403 (↑ 2.9×	411	9	1.0	19.4K	4	141	329	36.1
correlation	498.7K (↓ 290.5×	1	4168 (↑ 38.2×	362	19	0.8	144.9M	667	109	305	638.8
jacobi-2d	58.8K (↓ 183.1×	1	3968 (↑ 72.1×	411	17	0.9	10.8M	28	55	308	47.9
symm	405.7K (↓ 427.4×	1	1208 (↑ 201.3×	402	15	1.0	182.4M	13	6	397	3.5
trmm	492.6K (↓ 78.0×	1	101 (↑ 14.4×	414	12	0.8	38.4M	4	7	382	1.4

Memory & Communication customization can introduce
larger design space leading to better performance

Goal: Design a High-Performance LLM Accelerator

Step 1: Construct building blocks

- ▶ Performance
- ▶ Correctness
- ▶ Reusability
- Linear operators



- Non-linear operators

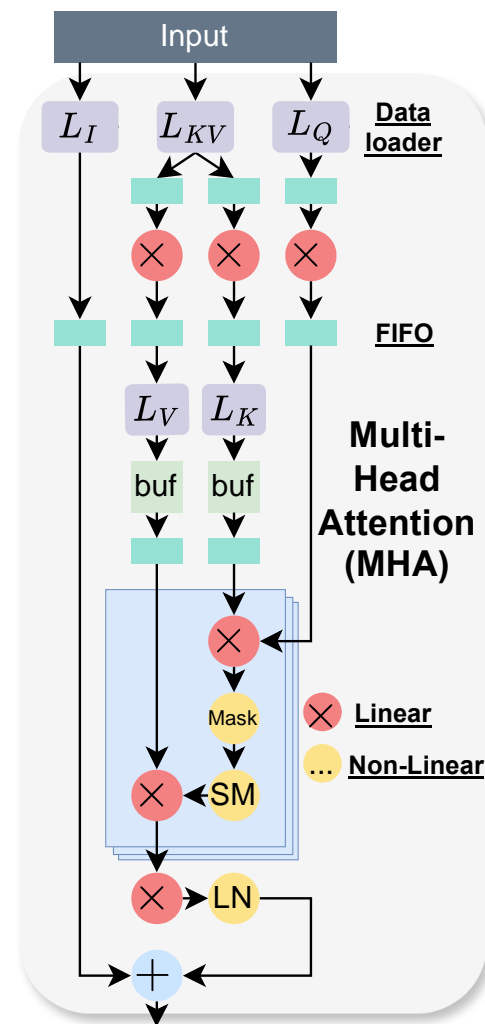
Softmax

LayerNorm

GELU

Step 2: Ensemble into a complete design

- ▶ Connect different operators bottom-up




We have the kernel library,
but how to compose them
to form a Transformer layer?

\times Linear
... Non-Linear

Allo ADL Example - Transformer Kernel (Alg. Specification)

- ▶ Allo ADL
 - High-level tensor-based operations
 - NumPy-like interface



Each kernel has its own schedule,
so how to glue them together?


```
def attention(hidden_states: float32[2, 512, 768])
    -> float32[2, 512, 768]:
    q_proj_weight: float32[768, 768] = global_q_proj_weight
    q_proj_bias: float32[768] = global_q_proj_bias
    k_proj_weight: float32[768, 768] = global_k_proj_weight
    k_proj_bias: float32[768] = global_k_proj_bias
    v_proj_weight: float32[768, 768] = global_v_proj_weight
    v_proj_bias: float32[768] = global_v_proj_bias
    q_proj = allo.linear(hidden_states, q_proj_weight, q_proj_bias)
    reshape = allo.reshape(q_proj, (2, 512, 12, 64))
    permute = allo.transpose(reshape, (0, 2, 1, 3))
    k_proj = allo.linear(hidden_states, k_proj_weight, k_proj_bias)
    reshape_1 = allo.reshape(k_proj, (2, 512, 12, 64))
    permute_1 = allo.transpose(reshape_1, (0, 2, 1, 3))
    v_proj = allo.linear(hidden_states, v_proj_weight, v_proj_bias)
    reshape_2 = allo.reshape(v_proj, (2, 512, 12, 64))
    permute_2 = allo.transpose(reshape_2, (0, 2, 1, 3))
    transpose = allo.transpose(permute_1, (-2, -1))
    matmul = allo.matmul(permute, transpose)
    truediv = matmul / 8.0
    softmax = allo.softmax(truediv)
    matmul_1 = allo.matmul(softmax, permute_2)
    permute_3 = allo.transpose(matmul_1, (0, 2, 1, 3))
    reshape_3 = allo.reshape(permute_3, (2, 512, 768))
    return reshape_3

s = allo.customize(attention)
mod = s.build()
example_inputs = [np.random.randn(2, 512, 768).astype(np.float32)]
out = mod(*example_inputs)
```

Composable Schedules

- ▶ `s.compose(<new_schedule>, <id>)`
 - `<id>`: Distinguish between different function calls
 - Can also compose external HLS IP modules

```
1  def top(X: int8[32, 64]) -> int8[64, 32]:
2      Z: int8[32, 64] = 0
3      Y: int8[64, 32] = 0
4      W_A: int8[64, 64] = W_A_cst
5      W_B: int8[64, 64] = W_B_cst
6      tiled_systolic[int8, int8, int16, \
7          32, 64, 64, "FFN1"](X, W_A, Z)
8      tiled_systolic[int8, int8, int16, \
9          64, 32, 64, "FFN2"](Z, W_B, Y)
10     return Y
11
12     s_top = allo.customize(top)
13     s_top.compose(s) # `s` is an optimized schedule
14     s_top.relay(s_top.Z, "tiled_systolic_FF2")
```




Composable Schedules: Implementation

► Hierarchical Dataflow Graph

- Traditional DFG are flattened
- Missing function boundaries info, less optimization opportunity

```
1 def add(A: T[M, N]) -> T[M, N]:  
2   B = A + 1  
3   return B  
4 def mul(A: T[M, K], B: T[K, N]) -> T[M, N]:  
5   # ... Calculate matrix multiply of A and B  
6 def top(A: T[M, K], B: T[K, N]) -> T[M, N]:  
7   C = mul(A, B)  
8   D = add(C)  
9   return D
```



How to resolve conflicts?

e.g., mul:C is partitioned in mul, but mul wants to be composed into top, where top:C is NOT partitioned

► Schedule Replay

- Replays customizations on design modules after composition

Algorithm 1: Composing multiple schedules with schedule replay

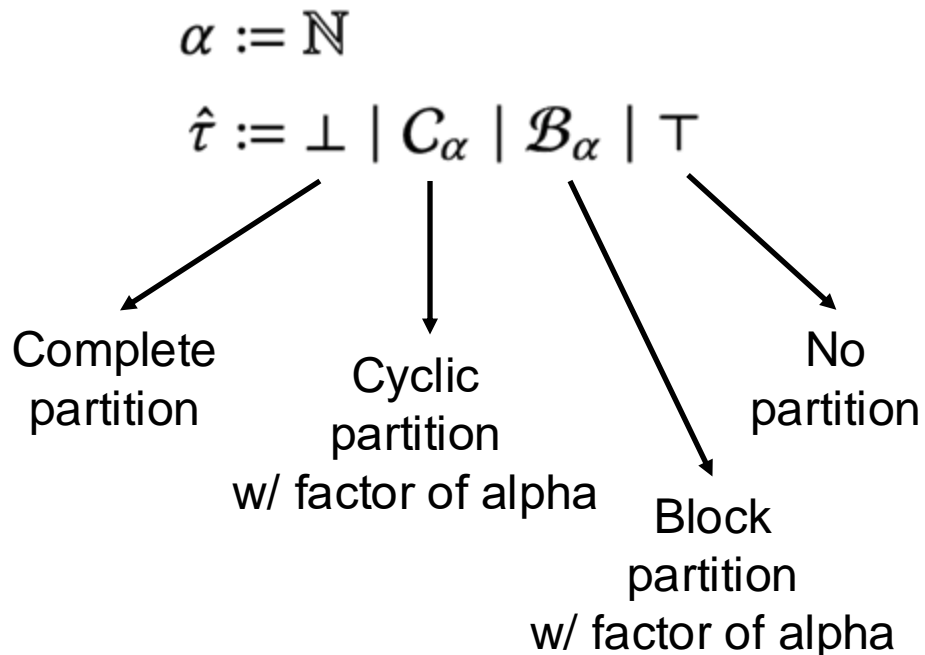
Data: Two schedules S_P and S_Q for programs P and Q

Result: Composition of the schedules $S_{out} = S_Q \circ S_P$ and the output program P' after applying S_{out}

```
1 Initialize  $S_{out} = S_P$ ;  
2 foreach primitive  $p_i \in S_Q$  do  
3   Update the arguments of  $p_i$  to refer to the functions and arguments in program  $P$ ;  
4   if  $p_i$  conflicts with primitives in  $S_{out}$  then  
5     Composition fails, raise an error;  
6   Append  $p_i$  to  $S_{out}$ ;  
7 Apply each primitive in  $S_{out}$  to the program  $P$  to obtain  $P'$ 
```

Memory Layout Composition

- ▶ Goal: Ensure the layouts of function arguments are consistent
- ▶ Key idea: Model data layout as a type
- ▶ Data layout propagation -> type inference
 - N-D array layout (composite type): $\tau := (\hat{\tau}_1, \dots, \hat{\tau}_N)$
 - Base type for each dimension:



Subtyping relation: $X <: Y$

The code expecting a memory with partition type Y is also compatible with a memory with partition type X

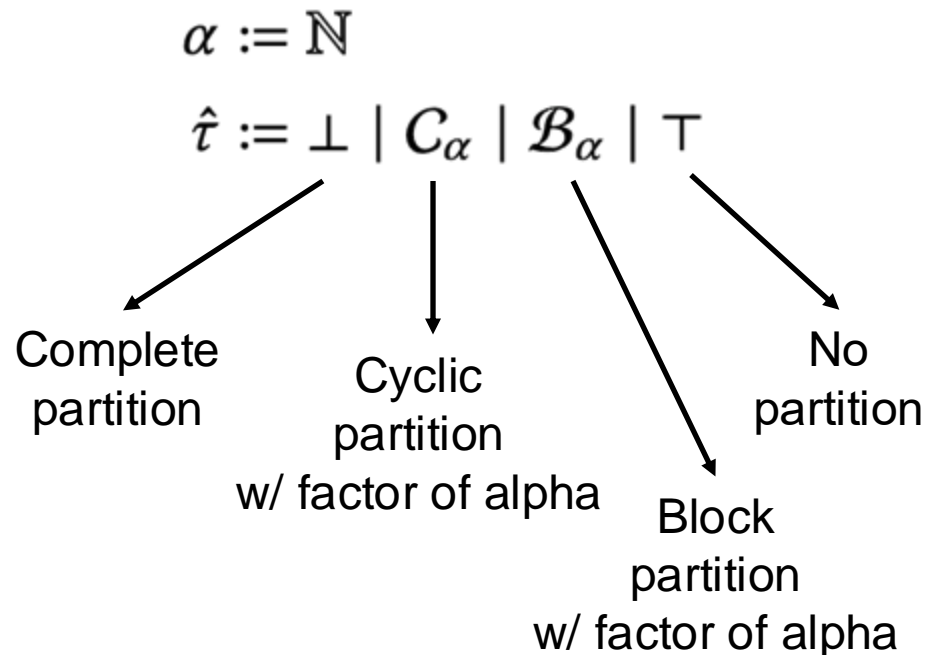
e.g., $\perp <: C_2$ since complete partitioning already partition the array into cyclic with a factor of 2

Intuition:


You can supply more read/write parallelism, but not less!

Memory Layout Composition

- ▶ Goal: Ensure the layouts of function arguments are consistent
- ▶ Key idea: Model data layout as a type
- ▶ Data layout propagation -> type inference
 - N-D array layout (composite type): $\tau := (\hat{\tau}_1, \dots, \hat{\tau}_N)$
 - Base type for each dimension:



Subtyping relation construct a **lattice**!



An example lattice of partition types for a 1D array of shape (8,)

Memory Layout Composition

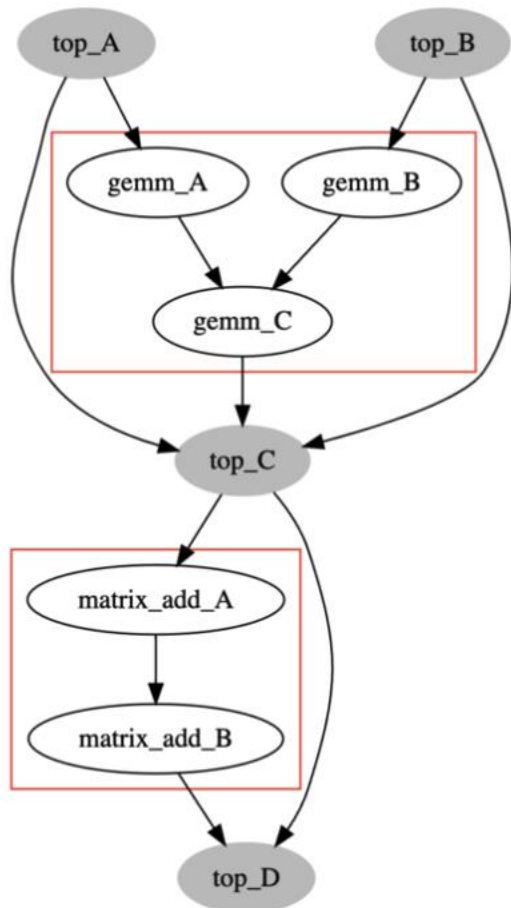
- ▶ Goal: Ensure the layouts of function arguments are consistent
- ▶ **Key idea: Model data layout as a type**
- ▶ Data layout propagation -> type inference
 - We check if layout is correctly composed with typing rules

$\frac{}{\perp <: C_\alpha}$ <p>S-BOTTOM-C</p>	$\frac{}{\perp <: \mathcal{B}_\alpha}$ <p>S-BOTTOM-B</p>	$\frac{\alpha_2 \equiv 0(\text{mod } \alpha_1)}{C_{\alpha_2} <: C_{\alpha_1}}$ <p>S-CYCLIC</p>	$\frac{\alpha_2 \equiv 0(\text{mod } \alpha_1)}{\mathcal{B}_{\alpha_2} <: \mathcal{B}_{\alpha_1}}$ <p>S-BLOCK</p>	$\frac{}{C_\alpha <: \top}$ <p>S-TOP-C</p>	$\frac{}{\mathcal{B}_\alpha <: \top}$ <p>S-TOP-B</p>
$\frac{\exists i \in \{1, \dots, N\} : \hat{\tau}_i <: \hat{\tau}'_i}{(\hat{\tau}_1, \dots, \hat{\tau}_N) <: (\hat{\tau}'_1, \dots, \hat{\tau}'_N)}$ <p>S-ARRAY</p>		$\frac{\Gamma \vdash f : \tau_1 \rightarrow \tau_2 \quad \Gamma \vdash e : \tau_3 \quad \tau_3 <: \tau_1}{\Gamma \vdash fe : \tau_2}$ <p>FUNCAPP</p>			

Memory Layout Composition

- ▶ Unification algorithm in functional programming does not work for type systems with subtypes
- ▶ Lattice has good property! We can directly use dataflow analysis for type inference (Worklist algorithm)

Guaranteed to terminate in linear time by the *fixed-point theorem* if (1) the structure is a finite lattice and (2) the transfer function is monotonic



Algorithm 2: Partition type inference (Memory layout propagation)

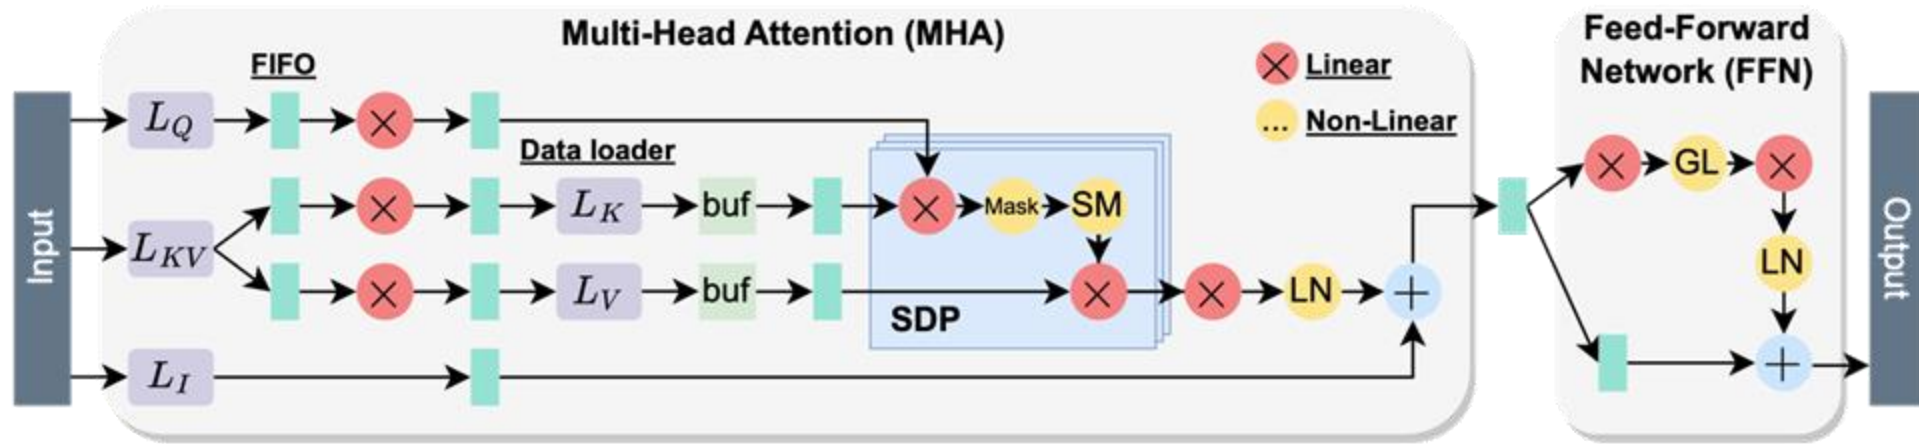
Data: The partition type $(t_1^{(0)}, \dots, t_M^{(0)})$ of the nodes (n_1, \dots, n_M) in the hierarchical dataflow graph, and a `.partition()` primitive on node n_{in} that transforms type t_{in} to t'_{in}

Result: Result partition type $(t_1^{(out)}, \dots, t_M^{(out)})$

```
1 Initialize Worklist  $\leftarrow \{(n_{in}, t'_{in})\}$ ;
2 while Worklist is not empty do
3   Pick an item of dataflow node and target type  $(n, t')$  from Worklist;
4   Update type  $t_n^{(next)} \leftarrow t' \sqcap t_n^{(curr)}$ ;
5   if  $t_n^{(next)} \neq t_n^{(curr)}$  then
6     foreach predecessors and successors  $\tilde{n}$  of  $n$  do
7       if  $\tilde{n}$  and  $n$  are in different functions then
8         Add  $(\tilde{n}, t_n^{(next)})$  to Worklist;
```

A Complete LLM Accelerator

- ▶ GPT2 model (the only open-source LLM in the GPT family)
 - 355M parameters, 24 hidden layers, 16 heads
 - W4A8 quantization



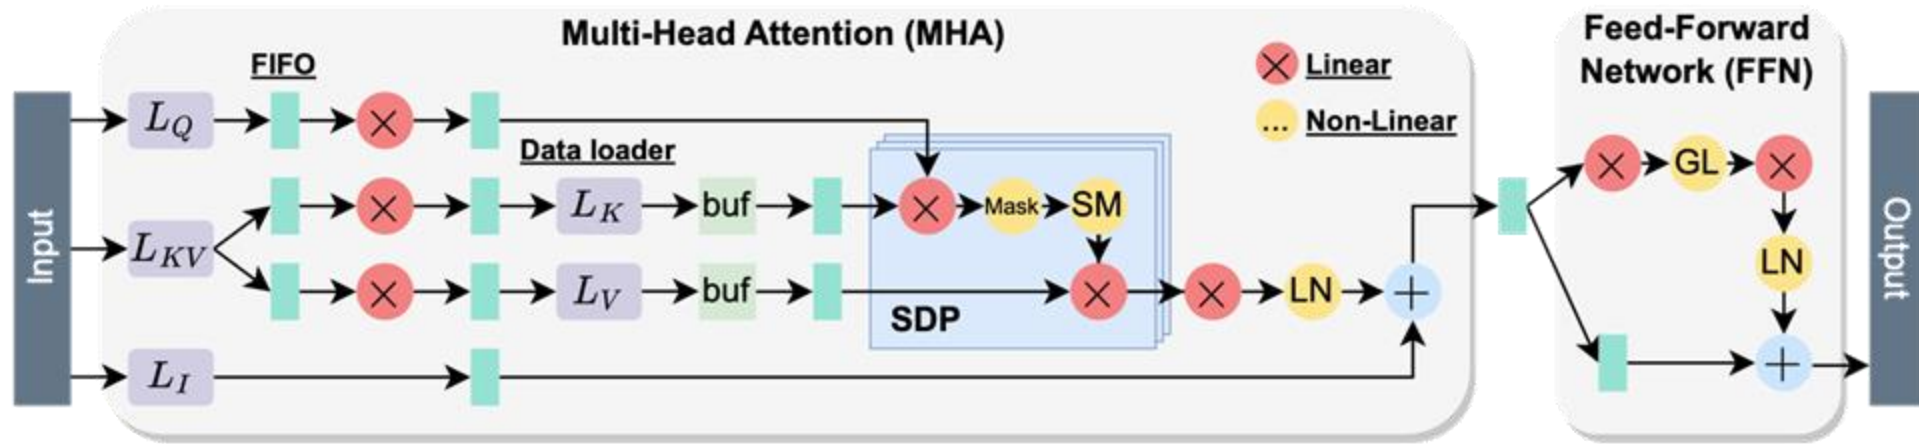
```
def GPT_layer(inp: float32[inp_num, inp_len], ...)
    -> float32[inp_num, inp_len]:
    # 1. Multi-Head Attention (MHA)
    Q = Linear_layer_qkv(inp, Wq, Bq)
    K = Linear_layer_qkv(inp, Wk, Bk)
    V = Linear_layer_qkv(inp, Wv, Bv)
    attn_sf_outp = Self_attention(Q, K, V)
    # ...
    # 2. Feed Forward Network (FFN)
    # ...
    return ffn_res_outp
```

Compose all the schedules together

```
s = allo.customize(GPT_layer)
s.compose(s_qkv)
s.compose(s_sfa)
s.compose(s_ds0)
s.compose(s_res)
s.compose(s_ln)
s.compose(s_ds1)
s.compose(s_ds2)
s.compose(s_gelu)
```


A Complete LLM Accelerator

- ▶ GPT2 model (the only open-source LLM in the GPT family)
 - 355M parameters, 24 hidden layers, 16 heads
 - W4A8 quantization



Is it the end?

No! We need to determine how many resources are allocated to each operator!

Compose all the schedules together

```
s = allo.customize(GPT_layer)
s.compose(s_qkv)
s.compose(s_sfa)
s.compose(s_ds0)
s.compose(s_res)
s.compose(s_ln)
s.compose(s_ds1)
s.compose(s_ds2)
s.compose(s_gelu)
```

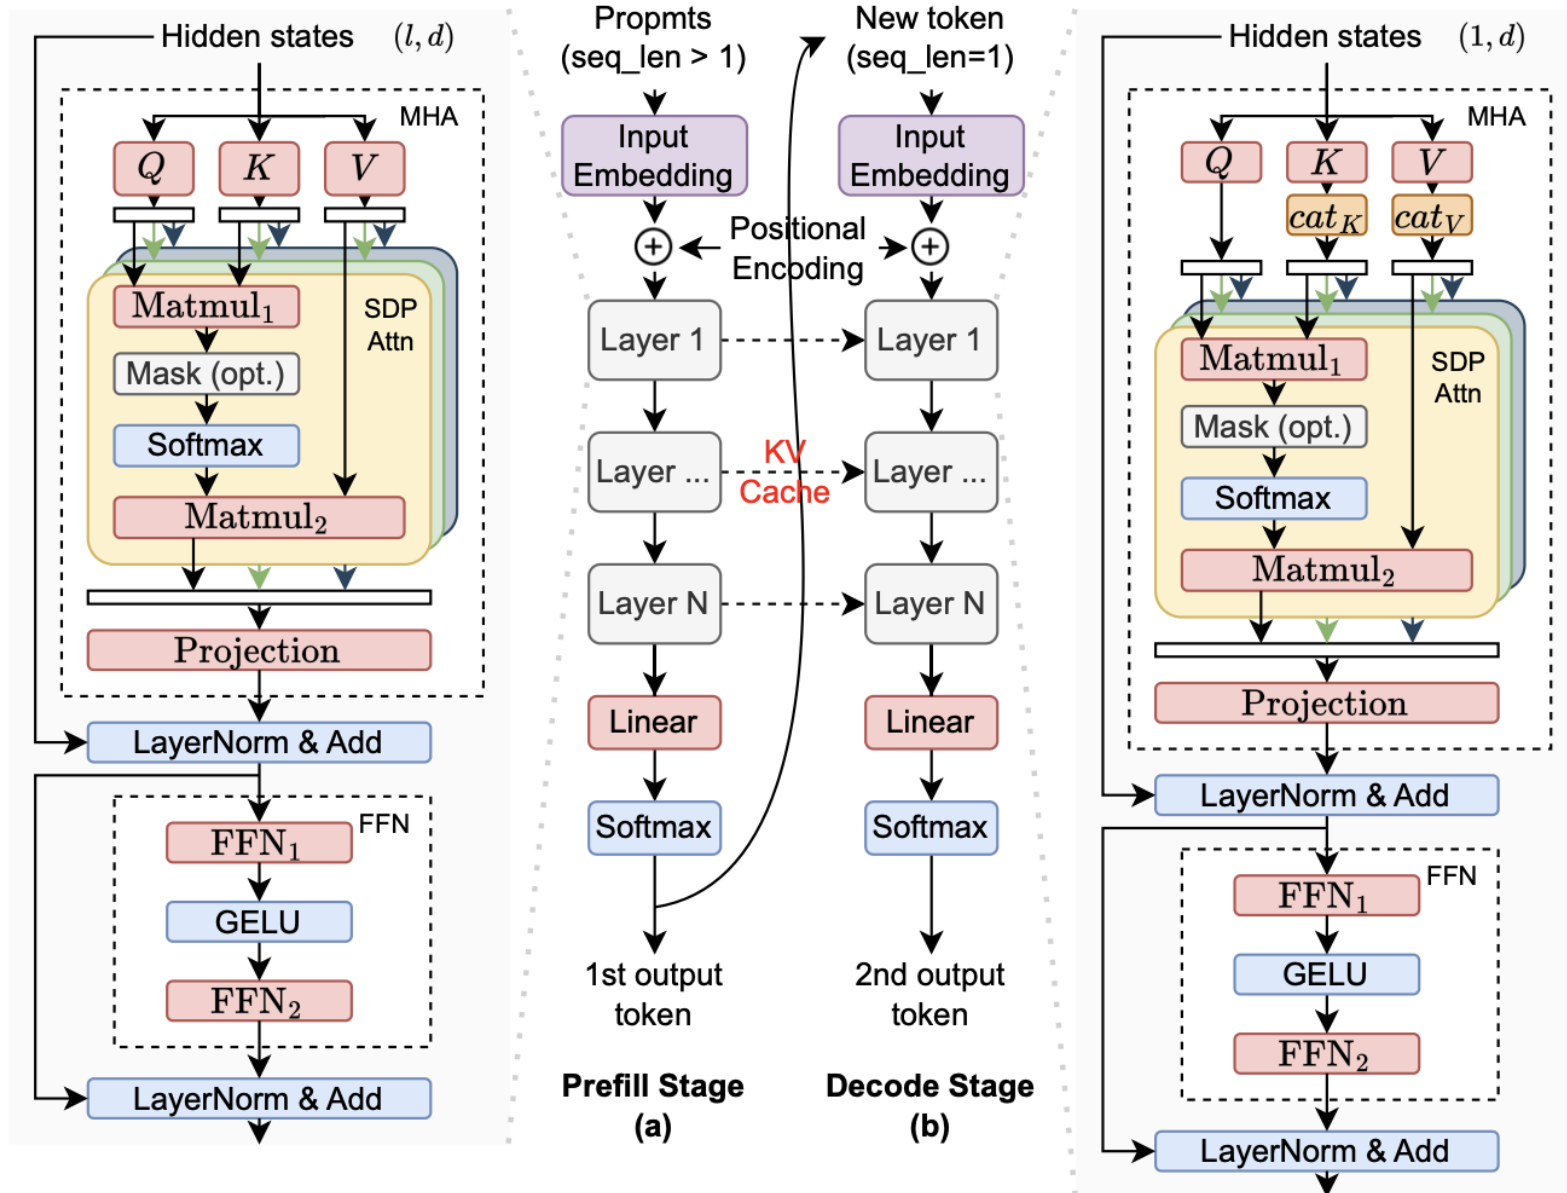
Two-Stage Generative Inference for LLM

- **Stage 1: Prefill**

- Take in user prompts and generate the 1st token
- $\text{seq_len} > 1$
- GEMM

- **Stage 2: Decode**

- Take in previous generated token and generate new tokens one at a time in an auto-regressive way
- $\text{seq_len} = 1$
- GEMV



Analytical Model for LLMs [FCCM/TRETS'24]

Linear Layer	Abbreviations	Input Matrices	Prefill	Decode
Q/K/V linear	q, k, v	XW_Q, XW_K, XW_V	$3ld^2$	$3d^2$
Matmul ₁	a_1	QK^T	l^2d	$(l+1)d$
Matmul ₂	a_2	$X_{sm}V$	l^2d	$(l+1)d$
Projection	p	$X_{sdp}W_{Proj}$	ld^2	d^2
FFN ₁	f_1	$X_{mha}W_{FFN_1}$	ldd_{FFN}	dd_{FFN}
FFN ₂	f_2	$X_{act}W_{FFN_2}$	ldd_{FFN}	dd_{FFN}

- **Compute resource:** M is compute power in MACs/cycle and C is the # of layers per FPGA

$$\sum M_i C < M_{tot}, i \in \{q, k, v, a_1, a_2, p, f_1, f_2\}$$

- **Memory capacity:** S is buffer size

$$S_{param} C < DRAM_{tot},$$

$$\sum S_i C < SRAM_{tot}, i \in \{tile, KV, FIFO\}$$

- **Memory port:** s is tensor size and b is bitwidth

$$R_i = \left\lceil \frac{s_i b_{BRAM}}{M_i / r_i \times S_{BRAM}} \right\rceil \times \frac{M_i / r_i}{k}$$

$$\sum_i C R_i + 2C(R_{a_1} + R_{a_2}) < SRAM_{tot}, i \in \{q, k, v, p, f_1, f_2\}$$

- **Memory bandwidth:** B is bandwidth

$$\sum_i C B_i < B_{tot}, i \in \{q, k, v, p, f_1, f_2\}$$




Figure : Pipeline diagram. Different colors stand for different input samples. Different blocks stand for different linear operators which also constitute the pipeline stages. h is the number of attention heads.

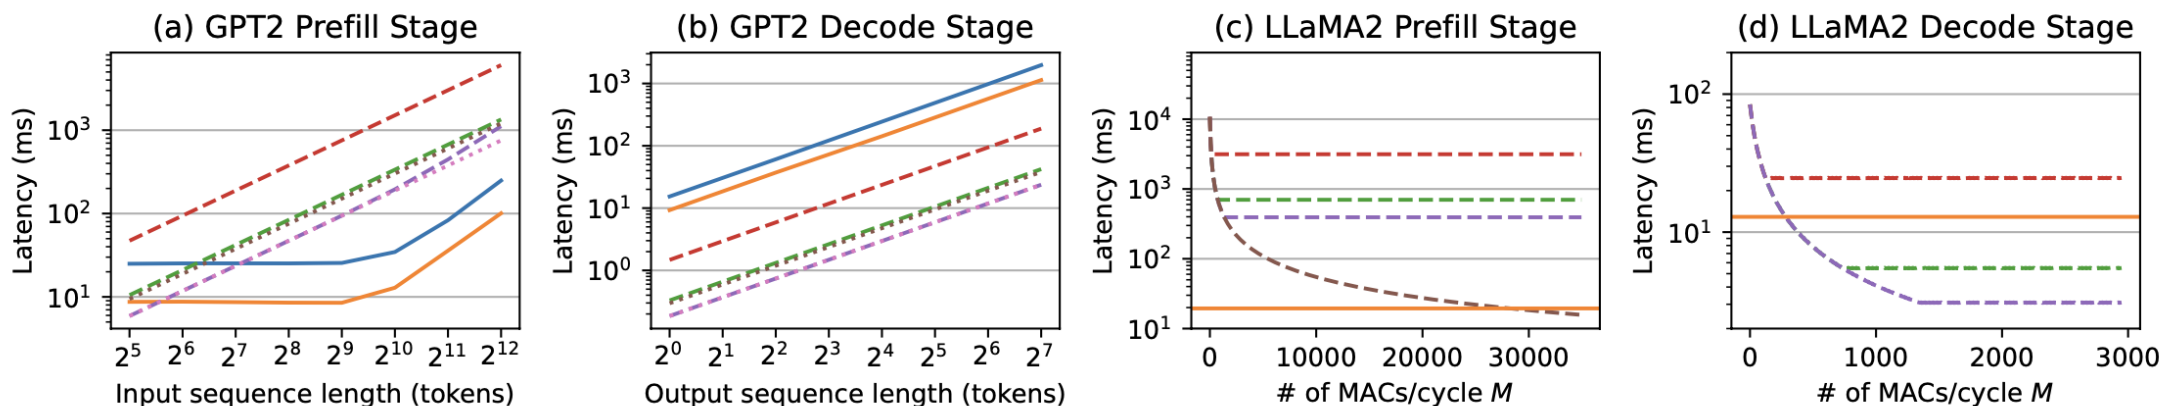
$$T_{prefill} = \frac{1}{freq} \frac{N}{C} \left(\frac{ld^2}{M_k} + C_{max} \left(\frac{ld^2}{M_k}, \frac{l^2d}{M_{a_1}}, \frac{ldd_{FFN}}{M_{f_1}}, T_{mem} \right) \right)$$

$$T_{decode} = \frac{1}{freq} \frac{N}{C} \left(\frac{d^2}{M_k} + C_{max} \left(\frac{d^2}{M_k}, \frac{(l_{max}+1)d}{M_{a_1}}, \frac{dd_{FFN}}{M_{f_1}}, T_{mem} \right) \right)$$

Analytical Model for LLMs

	AMD Xilinx FPGA			Intel FPGA		Nvidia GPU	
	Alveo U280 [83]	Versal VCK5000 [86]	Versal VHK158 [90]	Stratix 10 NX 2100 [37]	Agilex 7 AGM039 [27]	GeForce RTX 2080 Ti	Tesla A100
Process Node	TSMC 16nm	TSMC 7nm	TSMC 7nm	Intel 14nm	Intel 7nm	TSMC 12nm	TSMC 7nm
Release Date	2018	2022	2023	2020	2022	2018	2021
Thermal Design Power	225W	225W	180W	225W	225W	250W	300W
Peak Throughput	24.5 INT8 TOPS	145 INT8 TOPS	56 INT8 TOPS	143 INT8 TOPS	88.6 INT8 TOPS	14.2 TFLOPS	312 TFLOPS
Specialized Blocks	-	400× AI Engine	-	3960× AI Tensor Block	-	544× Tensor Cores	432× Tensor Cores
DSP/CUDA Cores	9024	1968	7392	-	12300	4352	6912
BRAM18K/M20K	4032	967	5063	6847	18960	-	-
URAM/eSRAM	960	463	1301	2	-	-	-
On-chip Memory Capacity	41MB	24MB	63.62MB	30MB	46.25MB	5.5MB	40MB
Off-chip Memory Capacity	8GB HBM2 & 32GB DDR	16GB DDR	32GB HBM2e & 32GB DDR	16GB HBM2	32GB HBM2e	11GB DDR	80GB HBM2e
On-chip Memory Bandwidth	460GB/s & 38GB/s	102.4GB/s	819.2GB/s & 102.4GB/s	512GB/s	820GB/s	616GB/s	1935GB/s


— RTX 2080Ti
 — Tesla A100
 - - - Alveo U280 (est.)
 - - - Versal VCK5000 (est.)
 - - - Versal VHK158 (est.)
 Stratix 10 (est.)
 Agilex 7 (est.)



Existing FPGAs are inferior in the compute-intensive prefill stage but can outperform GPUs in the memory-intensive decode stage.

LLM Accelerator Evaluation

- ▶ GPT2: single-batch, low-latency settings, adjust input/output token numbers
 - Use Allo to implement a design point in the analytical model (M=256)
 - 2.2x speedup in prefill stage compared to DFX (an overlay FPGA-based xcel)
 - 1.7x speedup for long output sequences and 5.4x more energy-efficient vs A100
 - < 50 lines of schedule code



	Allo	DFX
Device	U280	U280
Freq.	250MHz	200MHz
Quant.	W4A8	fp16
BRAM	384 (19.0%)	1192 (59.1%)
DSP	1780 (19.73%)	3533 (39.2%)
FF	652K (25.0%)	1107K (42.5%)
LUT	508K (39.0%)	520K (39.9%)

Composing Operators into Complete Design

- ▶ Predefined schedules for commonly used NN operators
- ▶ Can directly import model from PyTorch and build optimized xcel design
 - Through TorchDynamo and torch.fx

```
import torch
import allo
import numpy as np
from transformers import AutoConfig
from transformers.models.gpt2.modeling_gpt2 import GPT2Model

bs, seq, hs = 1, 512, 1024

example_inputs = [torch.rand(bs, seq, hs)]
config = AutoConfig.from_pretrained("gpt2")
module = GPT2Model(config).eval()
mlir_mod = allo.frontend.from_pytorch(
    module,
    example_inputs=example_inputs,
)
```


Summary

- ▶ Features of Allo ADL
 - Pythonic
 - Decoupled customizations
 - Composability
- ▶ Single-kernel
 - High-performance
 - verifiable
 - reusable
- ▶ Multi-kernel
 - First time to leverage an ADL to design a large-scale LLM accelerator
- ▶ Future work
 - Autoscheduling
 - Build system

Accpeted to PLDI'24

<https://arxiv.org/abs/2404.04815>

Allo: A Programming Model for Composable Accelerator Design

HONGZHENG CHEN*, Cornell University, USA

NIANSONG ZHANG*, Cornell University, USA

SHAOJIE XIANG, Cornell University, USA

ZHICHEN ZENG†, University of Science and Technology of China, China

MENGJIA DAI†, University of Science and Technology of China, China

ZHIRU ZHANG, Cornell University, USA



Available



Functional



Reusable

Accpeted to FCCM'24 Journal Track (TRETS)

<https://arxiv.org/abs/2312.15159>

Understanding the Potential of FPGA-Based Spatial Acceleration for Large Language Model Inference

HONGZHENG CHEN, Cornell University, USA

JIAHAO ZHANG*, Tsinghua University, China

YIXIAO DU, SHAOJIE XIANG, and ZICHAO YUE, Cornell University, USA

NIANSONG ZHANG, YAOHUI CAI, and ZHIRU ZHANG, Cornell University, USA



<https://github.com/cornell-zhang/allo>

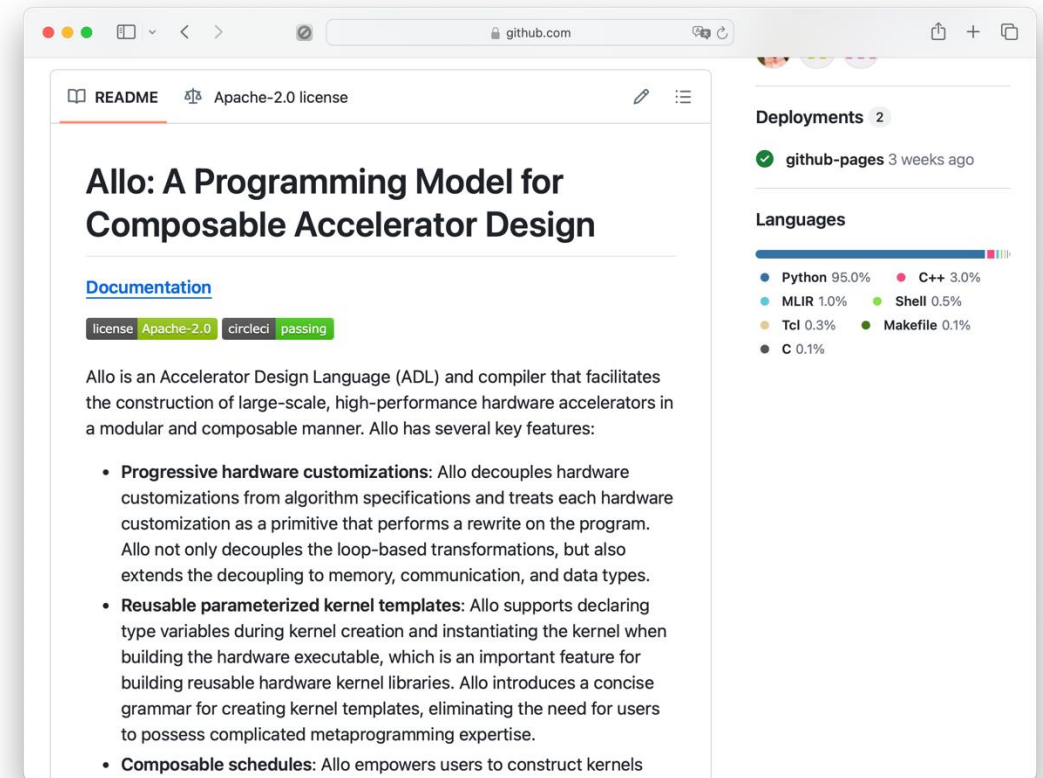
Acknowledgements

Contributors & Collaborators

- **Cornell:** Shaojie Xiang, Andrew Butt, Matthew Hofmann, Jie Liu, Zhongyuan Zhao
- **UIUC:** Hanchen Ye
- **UCLA:** Licheng Guo, Jason Lau, Yuze Chi, Jason Cong
- **UIC:** Debjit Pal
- **CSU:** Louis-Noël Pouchet, Emily Tucker
- **AWS:** Yi-Hsiang Lai
- **Tsinghua:** Jiahao Zhang, Wenbo Zhu
- **Peking:** Haiyu Wang
- **USTC:** Zhichen Zeng, Mengjia Dai
- **Intel:** Jeremy Casas, Pasquale Cocchini, Zhenkun Yang, Jin Yang, Hongbo Rong



<https://github.com/cornell-zhang/allo>



Sponsors



Semiconductor
Research
Corporation



CENTER FOR
EVOLVABLE
COMPUTING

